

G520SC

OPERATING SYSTEMS AND

CONCURRENCY

Processes and Threads

Dr Jason Atkin

Office Hours and Labs

- For lab questions please ask in the lab
 - If we need more time I can get lab helpers to help at other times too
- For coursework issues, we will have extra lab sessions with the lab helpers
- For course questions or other issues please see me either:
 - After the Tuesday G52CPP lecture (5pm outside LT3)
 - Or in office hours (C83), 11-12noon Wednesday
 - Or in the labs on Friday (either G52OSC or G52CPP lab)

This Lecture

- Linux : Creating a file, compiling
 - First three labs cover windows, but we will see Linux versions in lab 4, with walkthrough
- Linux create process and thread
 - fork() and pthreads
- Windows create process and thread
 - CreateProcess() and CreateThread()
 - Handles and WaitFor...
- Shared data vs separate
- Basic windows program

Processes and Threads

- An application consists of one or more processes.
- A *process*, in the simplest terms, is an executing program.
- One or more threads run in the context of the process.
- A *thread* is the basic unit to which the operating system allocates processor time.
- A thread can execute any part of the process code, including parts currently being executed by another thread.
- Source: <http://msdn.microsoft.com/en-gb/library/windows/desktop/ms684841%28v=vs.85%29.aspx>

Creating a process

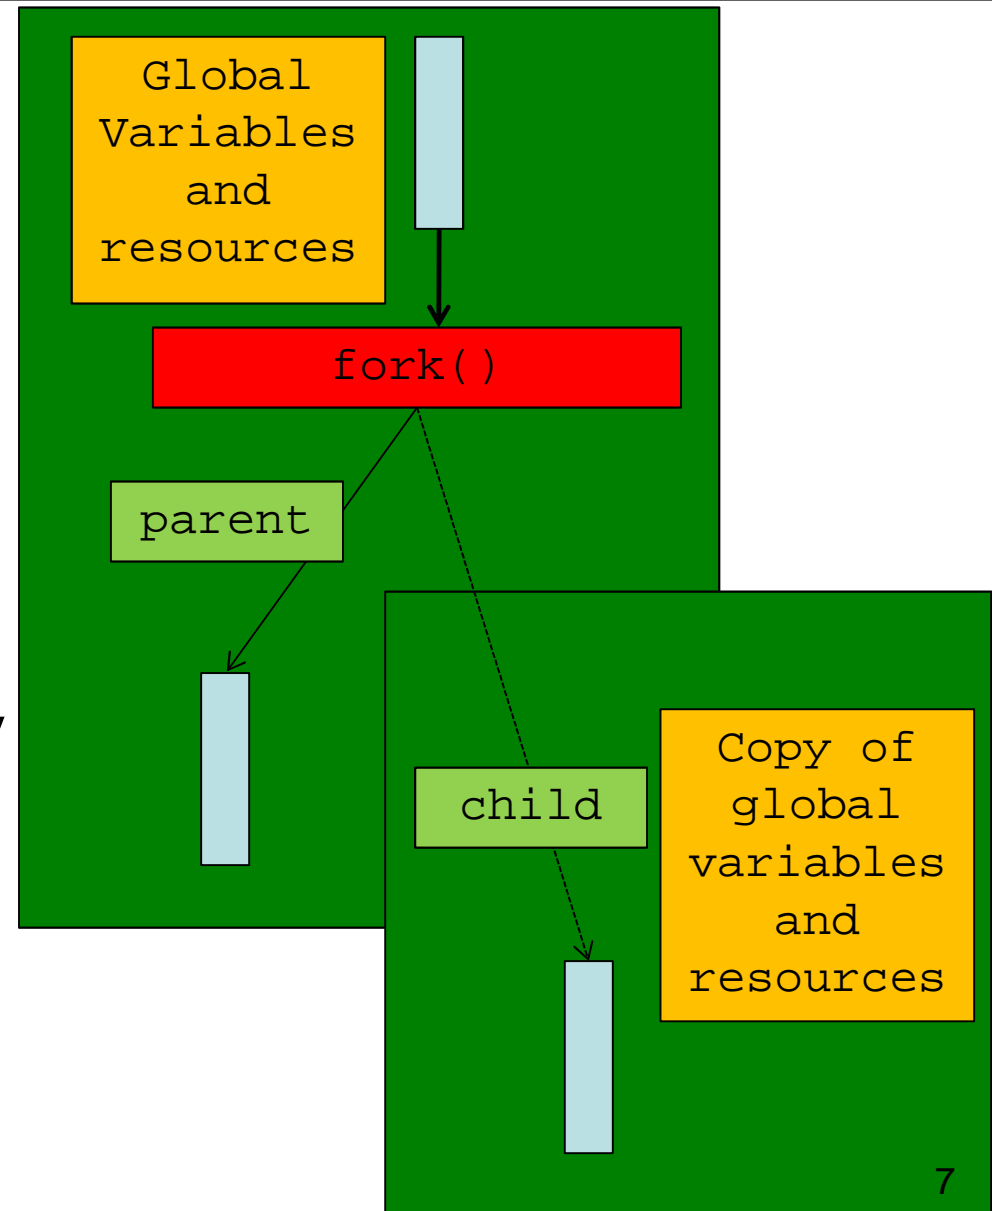
- Processes have their own state
- Instance of a running program
 - Will have an associated program/executable
- Create a process:
 - Tell the operating system to run that program
- Special way in Unix/Linux:
 - `fork()` – copy the current process

Overview: fork()

- Decide to start a new (copy) process
 - Call fork()
 - Check the return code
- Both processes will continue from the next operation in the program
 - The parent will be told the ID of the new child process as the return value
 - The child will get a 0 as the return value
 - This is the way you know which is which!

fork() – create new process

- Each process has its own resources
- You *can* share memory between both processes
 - but it is a pain
 - have to do it manually
 - often map a file (e.g. file ramdisk) into both processes
 - Next week



Example program with fork()

```
#include <stdio.h> // printf
#include <stdlib.h> // system
#include <unistd.h> // sleep
int main()
{
    int iProcID = fork();
    if ( iProcID == 0 )
    {
        printf( "In child. Waiting 3 seconds.\n" );
        sleep( 3 );
        printf( "Child ending now.\n" );
    }
    else
    {
        printf( "In parent. Child proc id is %d.\n",iProcID );
        sleep( 1 );
        printf( "In parent: process list:\n" );
        system( "ps -f" );
        sleep( 3 );
        printf( "Parent ending now.\n" );
    }
    return 0;
}
```


Creating a process

- You get a completely new process environment
- A copy of the old one
 - Often actually a copy-on-write
- From program point of view, a copy of:
 - Global variables
 - Allocated memory
 - System resources
 - etc

Global variable example (1)

```
#include <stdio.h>
#include <stdlib.h> // system
#include <unistd.h> // sleep
```

```
int global_variable = 1;
```

```
int main()
{
```

```
    int iProcID = fork();
```

```
    global_variable = 2;
```

```
    if ( iProcID == 0 )
    {
        global_variable = 3;
        printf(
            "In child. Waiting 3 seconds. Global variable is: %d\n",
            global_variable );
        sleep( 3 );
        printf( "Child ending now. Global variable is: %d\n",
            global_variable );
    }
```

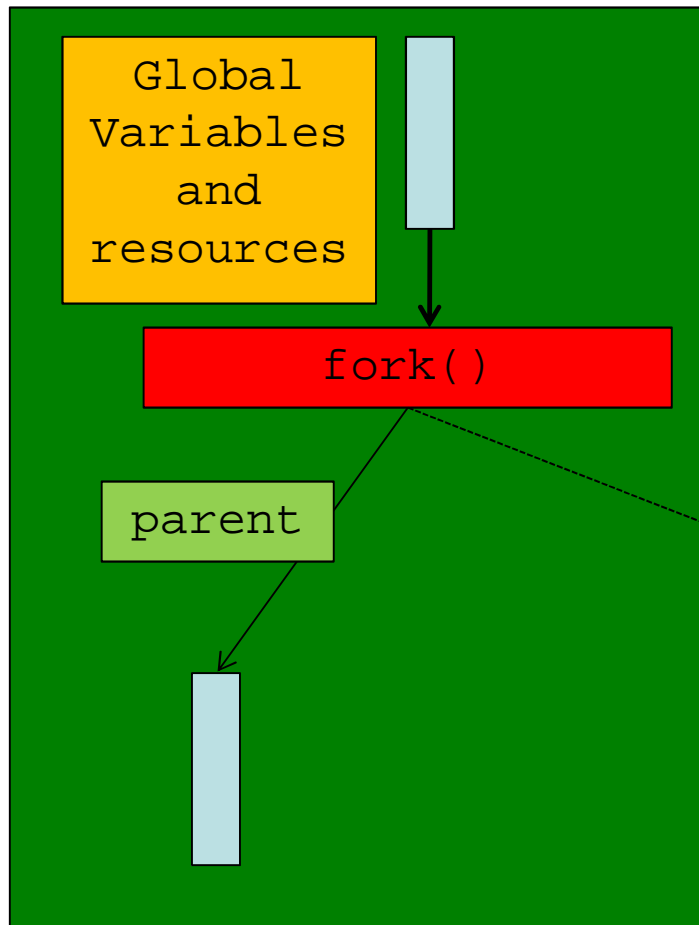
Global variable example (2)

else

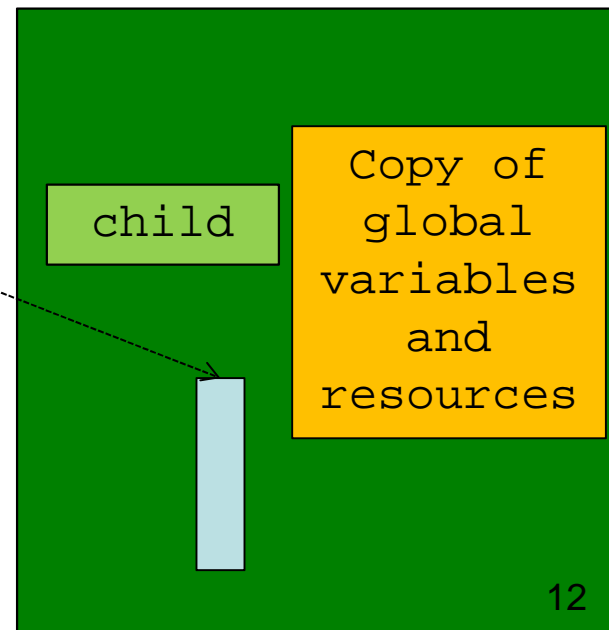
```
{  
    global_variable = 4;  
  
    printf( "In parent. Child proc id is %d. Wait 1 second.  
Global variable is: %d\n",  
            iProcID, global_variable );  
    sleep( 1 );  
    printf( "In parent: process list:\n" );  
    system( "ps -f" );  
    sleep( 3 );  
    printf( "Parent ending now. Global variable is: %d\n",  
            global_variable );  
}  
return 0;  
}
```

Two processes

Parent process



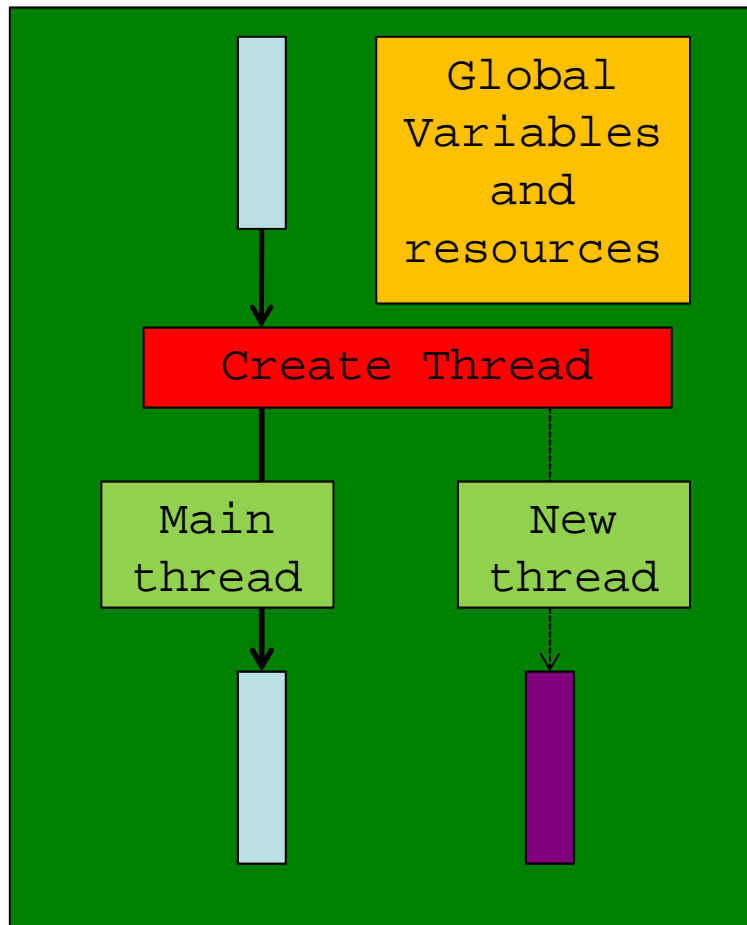
Child process



Threads

Threads – within a single process


Single process




- Another execution point within the same process
- **Same** copy of global variables and resources
- Starts from a specified (new) function

Fork() : TWO processes

Parent process




Child process



CreateThread : ONE process

Single process



The thread function and pre-amble

```
// Note: compile with 'gcc <filename> -lpthread'

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUM_THREADS      10

void* PrintHello(void *threadid) /*Single parameter*/
{
    long tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    sleep( 5 );
    printf( "Goodbye World! From thread #%ld!\n",tid );
    pthread_exit( NULL );
}
```

The main() function

```
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t;
    for( t=0; t<NUM_THREADS; t++ )
    {
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello,
                           (void *)t); /* Pass thread number in */
        if (rc)
        {
            printf("ERROR; return code was %d\n", rc);
            exit(-1);
        }
    }
    /* Exit the main thread - proc ends when the others end */
    pthread_exit(NULL);
    /* Could also call pthread_join to wait for threads.*/
}
```

An example of dividing
the workload
using multiple concurrent
operations
in Windows

Call the function a number of times

```
#define WIN32_LEAN_AND_MEAN
#include <Windows.h>
#include <stdio.h>
#include <stdlib.h>
```

```
#define NUM_THREADS 2
```

```
volatile DWORD dwTotal = 0;
```

```
DWORD WINAPI thread_function(
    LPVOID lpParam )
{
    for ( int i = 0;
          i < 1000000; i++ )
    {
        dwTotal++;
    }
    return 0;
}
```

```
int main()
{
    int iTN = 0;
    dwTotal = 0;
    for ( iTN = 0;
          iTN < NUM_THREADS;
          ++iTN )
    {
        thread_function( ... );
    }
    printf("Total %d\n",dwTotal);

    printf( "Press RETURN" );
    while ( getchar() != '\n' )
        ;
    return 0;
}
```

Windows types

- To try to promote portability, windows has a lot of custom types
 - Prevents sizes changing between platforms
- There are standard naming conventions for these
- LP : long pointer – basically means pointer to
- WORD : 2 byte value, unsigned
- DWORD : double word, 4 byte value, unsigned
- SIZE_T : common type for storing sizes, e.g. number of bytes
- LPVOID : Long pointer to a type that you don't care about – usually then cast to appropriate type

Reminders: windows things

```
volatile DWORD dwTotal = 0;
```

- Use volatile (C specifier) when the variable may be accessed from multiple threads/processes
 - Much more about this later

- DWORD specifies a size of 4 bytes, and is a name for the 'unsigned long' type here

```
DWORD WINAPI thread_function( LPVOID  
lParam )
```

- WINAPI is a flag to tell the compiler the format to lay out the function
- LPVOID means void* (long pointer to void)

The main part which calls function

```
int main()
{
    int iTN = 0;
    dwTotal = 0;
    for ( iTN = 0;
          iTN < NUM_THREADS;
          ++iTN )
    {
        thread_function( ... );
    }
    printf("Total %d\n",dwTotal);

    printf( "Press RETURN" );
    while ( getchar() != '\n' )
        ;
    return 0;
}
```

- Loops **NUM_THREADS** times, calling the **thread_function** manually each time
- We could tell it to do these simultaneously


Thread creation

- You create one new thread at a time
- There is an overhead in creating threads
 - Function `CreateThread()` takes time
 - Operating system has to allocate the resources for the thread
- You have no idea what order the threads will execute, where any interleaving happens, etc
 - Parts of one may be executed, then parts of another, etc. You don't even know which will finish first
- Your original thread can do some of the work itself (avoiding one creation)
- You may need to wait for things to finish

Parallelising the code (3 times)



- We could make these simultaneous
- By creating multiple threads
 - Create them
 - Run one function call in each
 - Wait for them to exit



Original program, ready for threads

```
#define WIN32_LEAN_AND_MEAN
#include <Windows.h>
#include <stdio.h>
#include <stdlib.h>
```

```
#define NUM_THREADS 2
```

```
volatile DWORD dwTotal = 0;
```

```
DWORD WINAPI thread_function(
    LPVOID lpParam )
{
    for ( int i = 0;
          i < 1000000; i++ )
    {
        dwTotal++;
    }
    return 0;
}
```

```
int main()
{
    int iTN = 0;
    dwTotal = 0;
    for ( iTN = 0;
          iTN < NUM_THREADS;
          ++iTN )
    {
        thread_function( ... );
    }
    printf("Total %d\n",dwTotal);

    printf( "Press RETURN" );
    while ( getchar() != '\n' )
        ;
    return 0;
}
```

This part remains unchanged

```
#define NUM_THREADS 2

// Windows name for an unsigned long - 4 bytes
volatile DWORD dwTotal = 0;

// Function called by every thread
DWORD WINAPI thread_function( LPVOID lParam )
{
    for ( int i = 0; i < 1000000; i++ )
    {
        dwTotal++;
    }
    return 0;
}
```

Threaded version of main()

```
HANDLE arrdwThreadHandles[NUM_THREADS];

for ( iTN = 0; iTN < NUM_THREADS - 1; ++iTN )
    arrdwThreadHandles[iTN] = CreateThread(
        NULL, /* No security change */
        0, /* Default stack size */
        thread_function, /* Name of function to call */
        (LPVOID)iTN, /* parameter you can give */
        0, /* Extra flags */
        NULL /* You can get the thread id if you wish */
    );

/* Do the last one in the current thread */
thread_function( (LPVOID)(NUM_THREADS - 1) );

WaitForMultipleObjects( NUM_THREADS - 1,
    arrdwThreadHandles, /* Array of handles */
    TRUE, 10000 ); /* Wait for all, for up to 10 secs */
```

CreateThread

- <http://msdn.microsoft.com/en-us/library/windows/desktop/ms682453%28v=vs.85%29.aspx>

```
HANDLE WINAPI CreateThread(  
    LPSECURITY_ATTRIBUTES    lpThreadAttributes,  
    SIZE_T                   dwStackSize,  
    LPTHREAD_START_ROUTINE   lpStartAddress,  
    LPVOID                    lpParameter,  
    DWORD                     dwCreationFlags,  
    LPDWORD                   lpThreadId );
```

- When you see LPSECURITY_ATTRIBUTES, just pass NULL to say to give it the permissions of the current process
- Stack size says how much memory it needs for functions and local variables (give it a lot if you use recursion)
- LPTHREAD_START_ROUTINE asks for the function to run in the thread. It must be labelled "WINAPI" since you are not calling it

Windows Handles

- A `Handle` gives you something to use to 'grip' a windows object that the system owns
- Processes, threads, windows, events, mutexes etc all have handles that you can use
- You can't do anything with these apart from use them to refer to the object
- `CreateThread` returns a `Handle` for the thread that was created
- We store all of these thread handles in an array in case we need them later

```
arrdwThreadHandles[iTN] =  
    CreateThread( ... );
```

Waiting for handles

- A thread handle can be used to determine whether a thread has finished or not
- You can wait for handles
 - Waiting for an event to go off
 - For a resource to be free
 - Or for a thread to finish
- Waiting for multiple handles to all be ready can be better than waiting for them one at a time

WaitFor...

- Wait for one or all of multiple handles to be 'ready'

```
WaitForMultipleObjects(  
    NUMBER_THREADS - 1, /*Count*/  
    arrdwThreadHandles, /*Array*/  
    TRUE,                /*All vs one?*/  
    10000 );             /*Timeout ms*/
```

```
DWORD WINAPI WaitForMultipleObjects(  
    DWORD nCount,  
    const HANDLE *lpHandles,  
    BOOL bWaitAll,  
    DWORD dwMilliseconds );
```


Next Lecture

- Windows GUI programs
 - Registering window classes
 - Creating windows
 - Message loops
- Note that both X (Unix/Linux) and Java use similar methods