

G520SC

OPERATING SYSTEMS AND

CONCURRENCY

Atomic Operations

Dr Jason Atkin

Last lecture


- Sharing memory
- Synchronising processes/threads using memory
 - Examining a variable set by another process
- Threads share an address space
- Processes have their own address space
 - But you can share memory between them
 - Have to do it manually!
 - Using a struct pointer is a simple way to manage this shared memory

This lecture


- Coordinating access to data
- Ornamental garden problem again
 - Why it went wrong
 - How it can be fixed
- Thread traces
- Atomic operations
- Critical sections
- Spin-locks

An example concurrency problem


- Ornamental garden problem:
 - Search for: "ornamental garden" concurrency
- Each should increment a **shared** counter when someone enters the garden



Two-thread update



- We create an extra thread
- We run the function twice
 - One from each thread
 - Including initial one
- Each function increments the value 1 million times
- What is the final value?



Example function

```
#include <Windows.h>
#include <stdio.h>
#define NUM_THREADS 2

// Windows name for an unsigned long - 4 bytes
volatile DWORD dwTotal = 0;

// Function called by every thread
DWORD WINAPI thread_function( LPVOID lParam )
{
    for ( int i = 0; i < 1000000; i++ )
    {
        dwTotal++;
    }
    return 0;
}
```

Contents of main()

```
HANDLE arrdwThreadHandles[NUM_THREADS];
```

```
for ( int iTN = 0; iTN < NUM_THREADS - 1; ++iTN )
```

```
    arrdwThreadHandles[iTN] = CreateThread(  
        NULL, 0,  
        thread_function, /* Name of function to call */  
        (LPVOID)iTN, /*parameter you can give*/  
        0, NULL  
    );
```

```
/* Do the last one in the current thread */
```

```
thread_function( (LPVOID)(NUM_THREADS - 1) );
```

```
WaitForMultipleObjects( NUM_THREADS - 1,  
    arrdwThreadHandles, /* Array of handles */  
    TRUE, 10000 ); /* Wait for all, for up to 10 secs */
```

What is the result when run twice?

```
// Function called by every thread
DWORD WINAPI thread_function(LPVOID lParam)
{
    for ( int i = 0; i < 1000000; i++ )
    {
        dwTotal++;
    }
    return 0;
}
```

Each function
increments the
variable one
million times

Run the function twice, at the same time

What will the result be?

Full code (for copy-paste)

```
int main()
{
    HANDLE arrdwThreadHandles[NUM_THREADS];

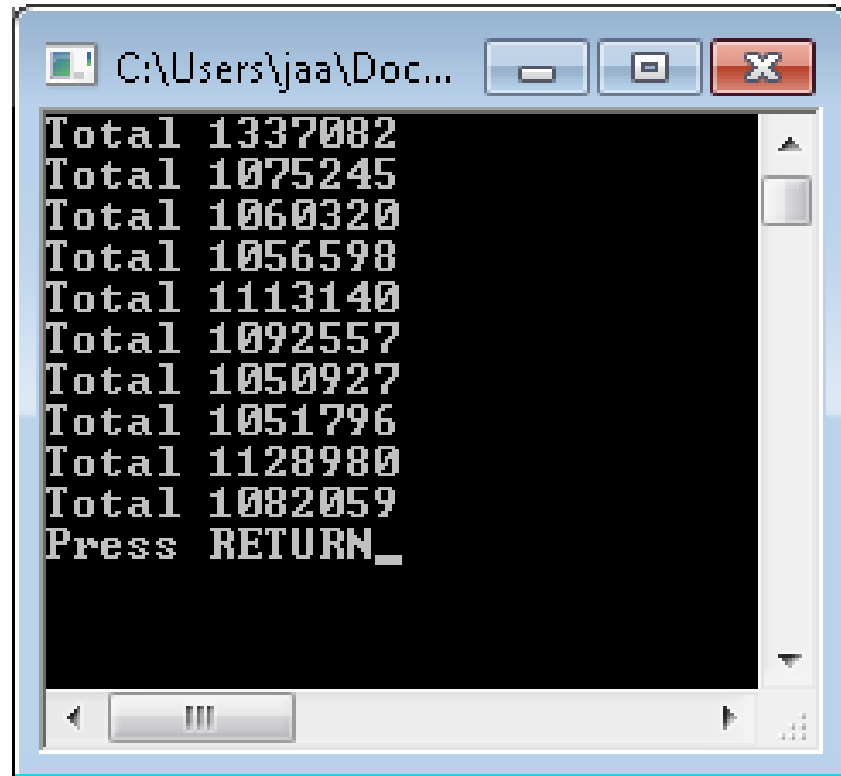
    for ( int iTN = 0; iTN < NUM_THREADS - 1; ++iTN )
        arrdwThreadHandles[iTN] = CreateThread( NULL,0,
            thread_function, /* Name of function to call */
            (LPVOID)iTN, /*parameter you can give*/ 0,NULL );

    /* Do the last one in the current thread */
    thread_function( (LPVOID)(NUM_THREADS - 1) );

    WaitForMultipleObjects( NUM_THREADS - 1,
        arrdwThreadHandles, /* Array of handles */
        TRUE,10000 ); /* Wait for all, for up to 10 secs */

    printf( "%d\n",dwTotal );
    while ( getchar() != '\n' );
}
```

Results



```

C:\Users\jaa\Doc...
Total 1337082
Total 1075245
Total 1060320
Total 1056598
Total 1113140
Total 1092557
Total 1050927
Total 1051796
Total 1128980
Total 1082059
Press RETURN_

```

- Should have been 2,000,000
- Problem:
dwTotal++;
 - Get current value
 - Increment
 - Write new value back

What went wrong?

- What does the increment do?
 1. Load current value from memory into a register
 2. Increment the register
 3. Write value back again to memory
- What could go wrong?

What should have happened?

	T1	Mem	T2	
• Thread 1				• Thread 2
• Get value	0	\leftarrow	0	
• Increment value	1		0	
• Write value	1	\rightarrow	1	
			1	\rightarrow 1
			1	2
			2	\leftarrow 2
• Get value	2	\leftarrow	2	• Get value
• Increment value	3		2	• Increment value
• Write value	3	\rightarrow	3	• Write value

Example of what can go wrong

	T1	Mem	T2	
• Thread 1				• Thread 2
• Get value	0 ← 0			
	0	0 →	0	• Get value
• Increment value	1	0	0	
	1	0	1	• Increment value
• Write value	1 → 1		1	
		1 ←	1	• Write value
• Get value	1 ← 1			
	1	1 →	1	• Get value
• Increment value	2	1	1	

Traces

- Thread 1

- Get value
- Increment value
- Write value

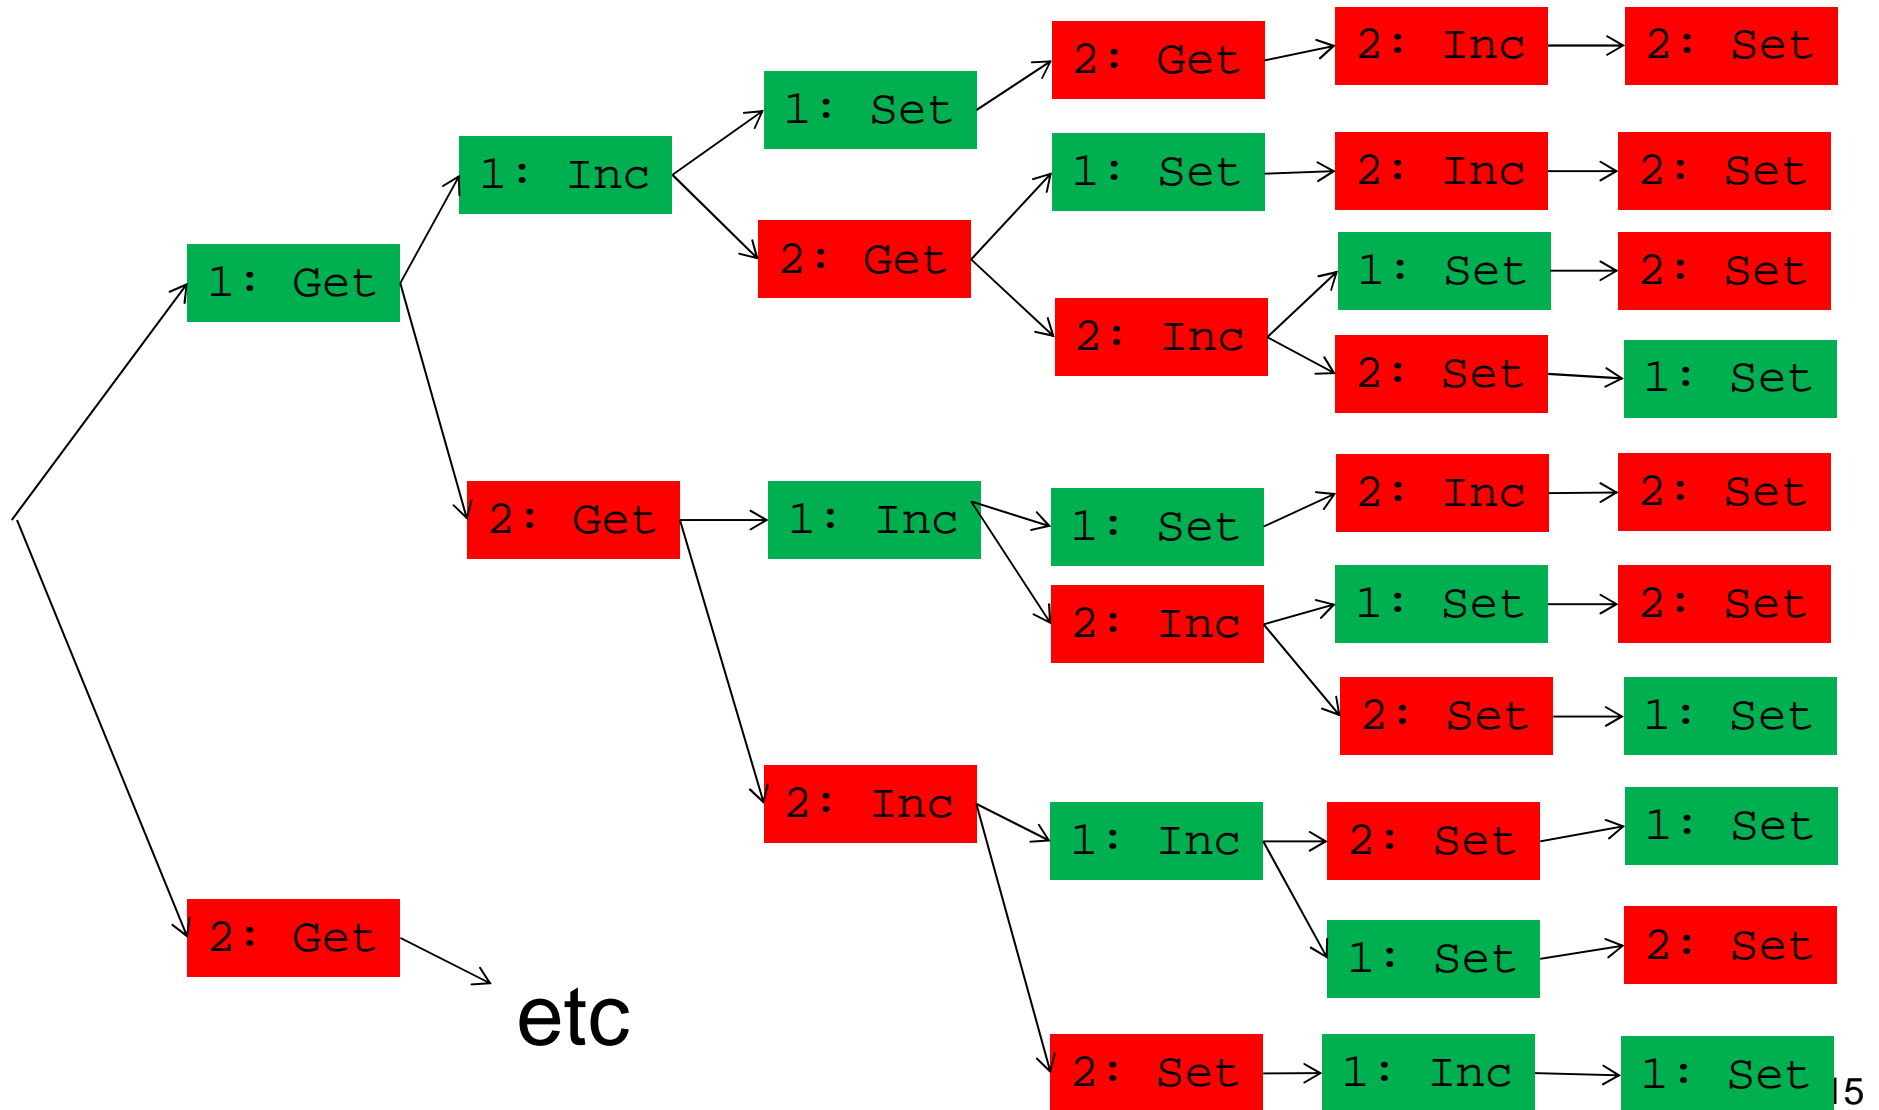
- Thread 2

- Get value
- Increment value
- Write value

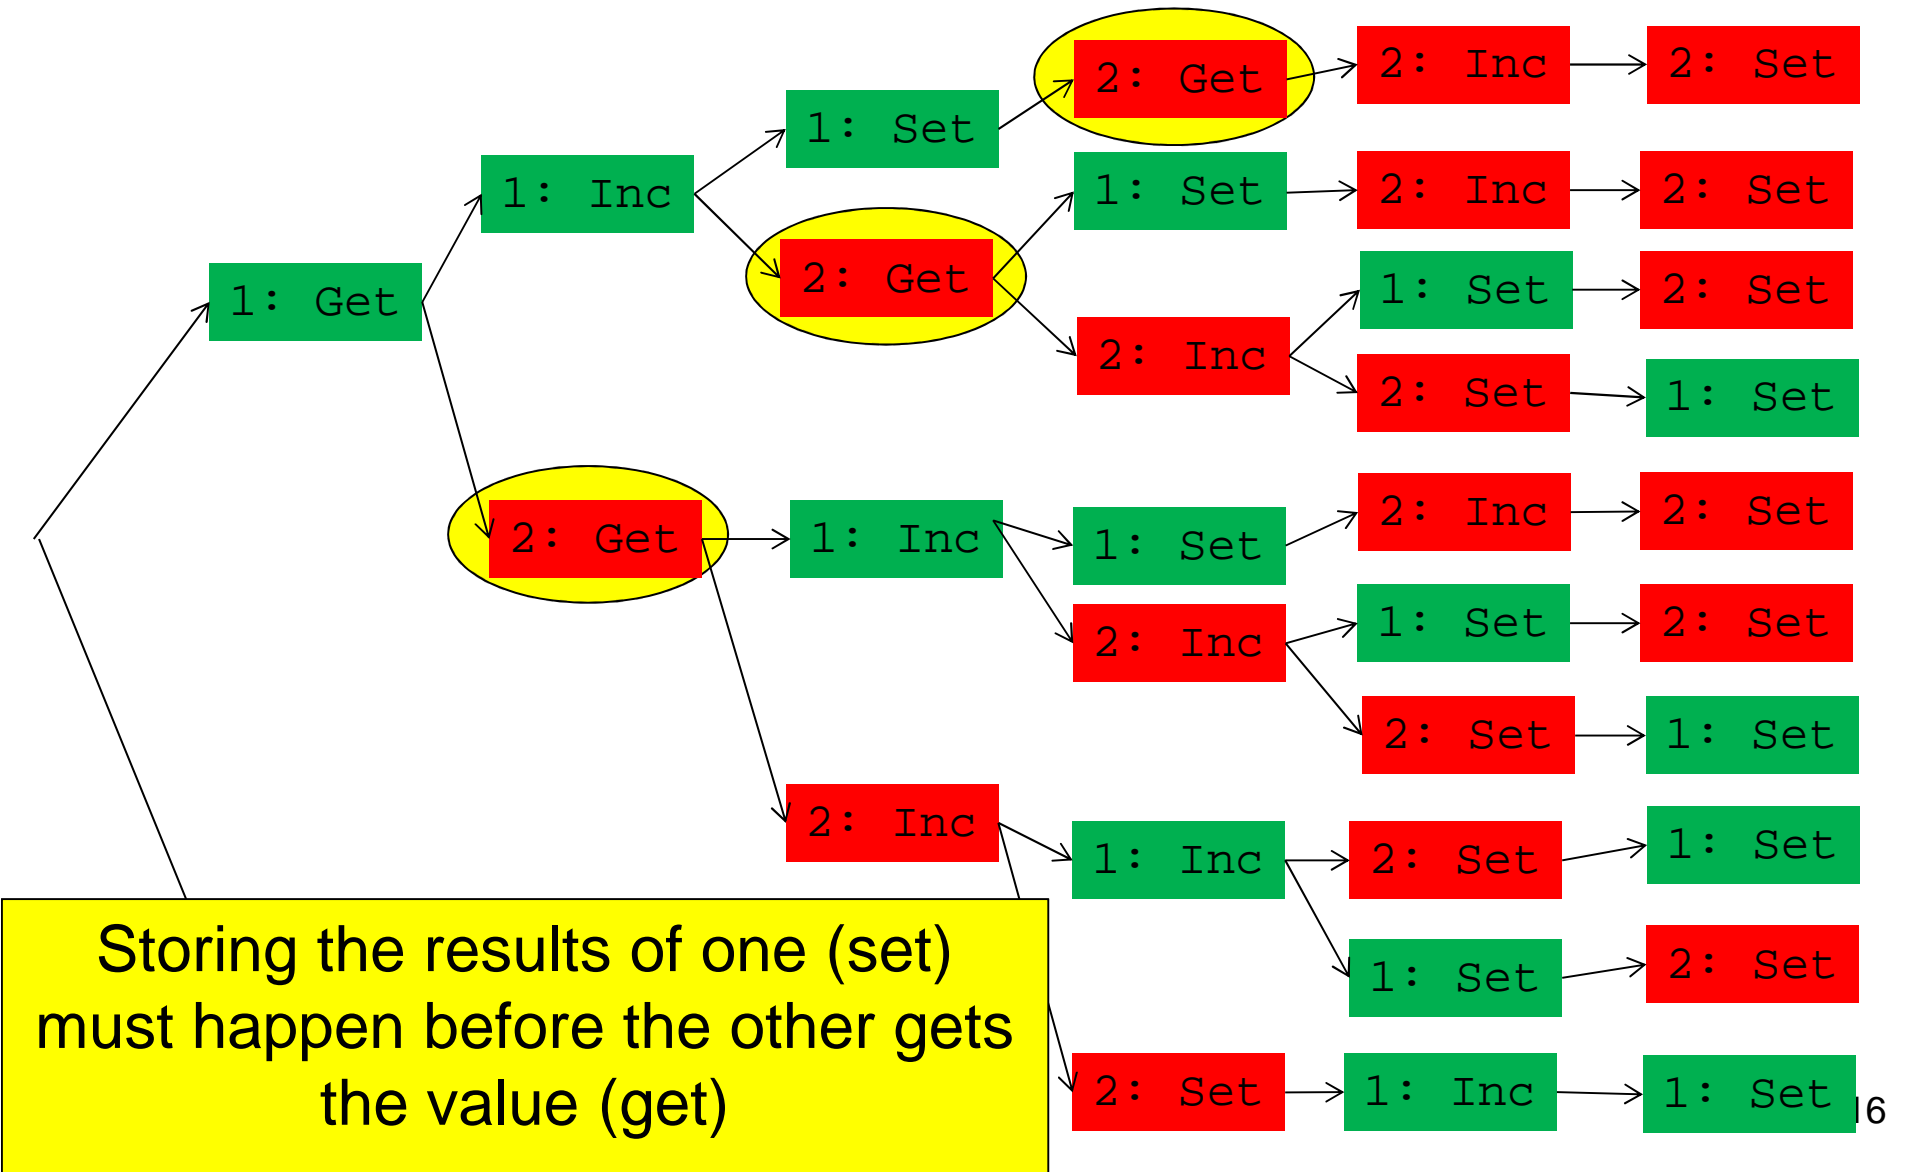
We will trace the potential execution orders for these operations to see whether the system will perform as we may expect

We will assume just one Get-Increment-Write for each, to simplify matters – the real situation is even worse

Potential traces




Which are good traces?



Methods to fix the problem

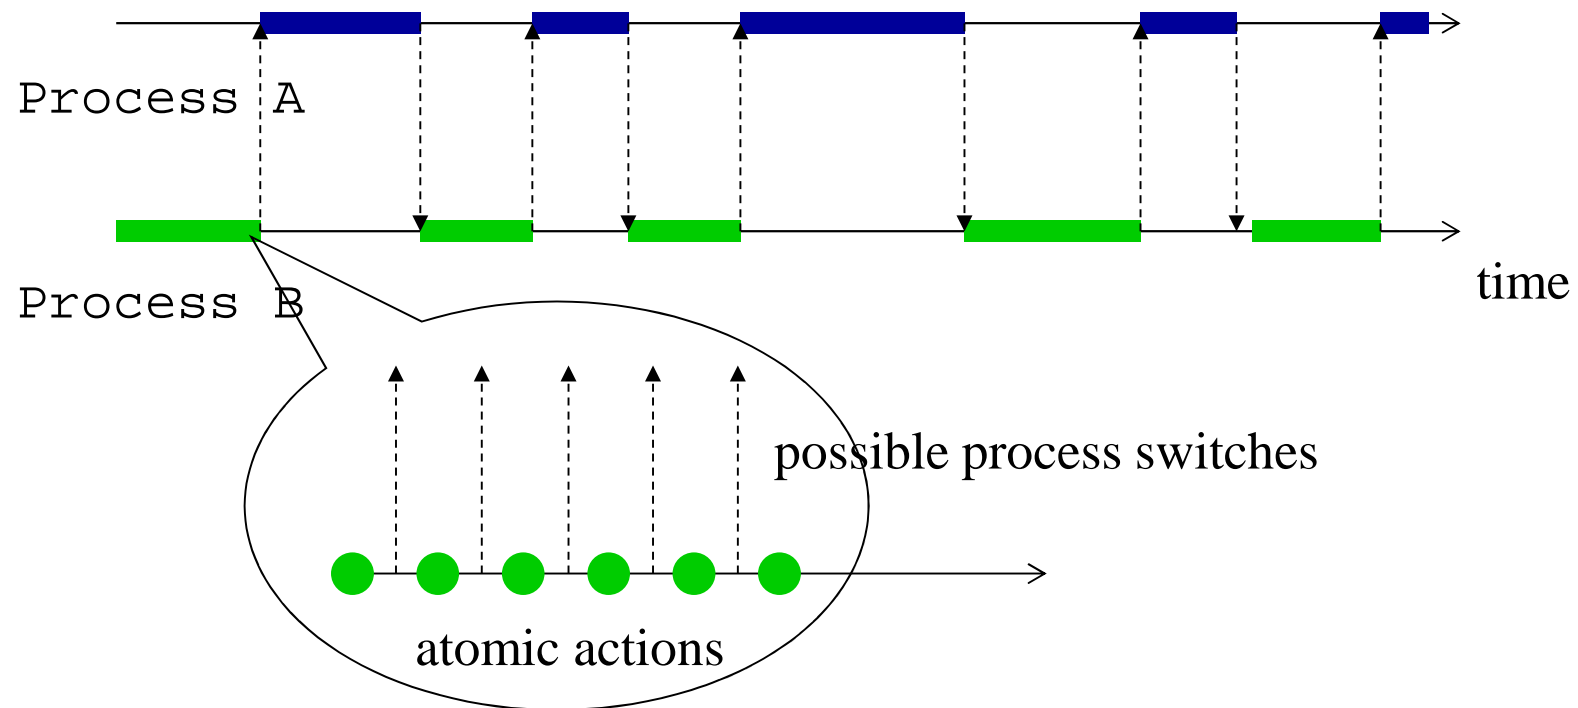
One processor (core)

- Consider a single processor situation
- Time is split between the processes



- At some point our process is interrupted

Process switches

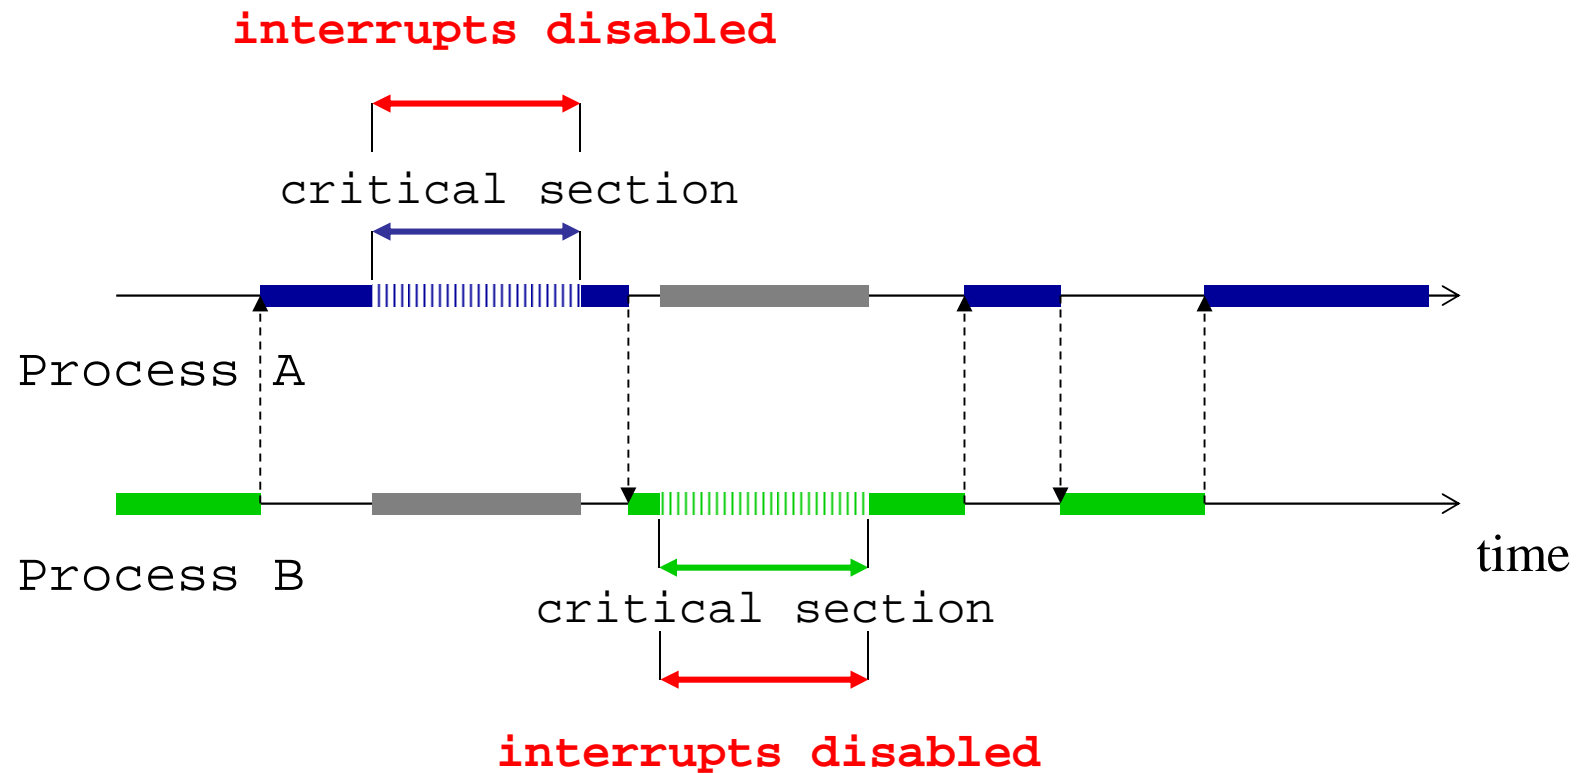


- Operations of a process consist of multiple atomic operations

Potential solution (single processor)

- The problem is that the get-increment-set is not an atomic operation
- Three obvious solutions on a single-processor system:
 1. Stop process switches happening between the operations
 - i.e. Turn off process switches/interrupts
 2. Make the get-increment-set operation atomic
 - Then it is just one operation
 3. Use hardware features (e.g. memory locks, transactional memory) – ignore for now

Turn off interrupts/switches



Problems with turning off interrupts

- Operating system must allow it
 - Makes the task scheduling harder
 - Usually very unwise
 - One process could take all of the CPU time if it doesn't turn them on again
 - Some interrupts (e.g. from devices) may really need to be handled as soon as possible
 - System-level programming only usually
- In a multi-processor system it may not work
 - Interrupt disabling usually applies to same processor
 - Other threads can change things even when you prevent interrupts on your processor

Second possibility

- Make your operation atomic
 - i.e. the whole thing is then one operation and cannot be interrupted
- Most CPUs provide atomic operations which can avoid some of these problems
 - E.g. 'test and set', 'increment', 'swap'
- Most C compilers allow you a facility to get at these atomic operations, e.g:
 - Interlocked API (windows, e.g. InterlockedIncrement)
 - `__sync_...` (gcc, e.g. `__sync_fetch_and_add()`)
- Usually platform or compiler specific options!

Atomic InterlockedIncrement

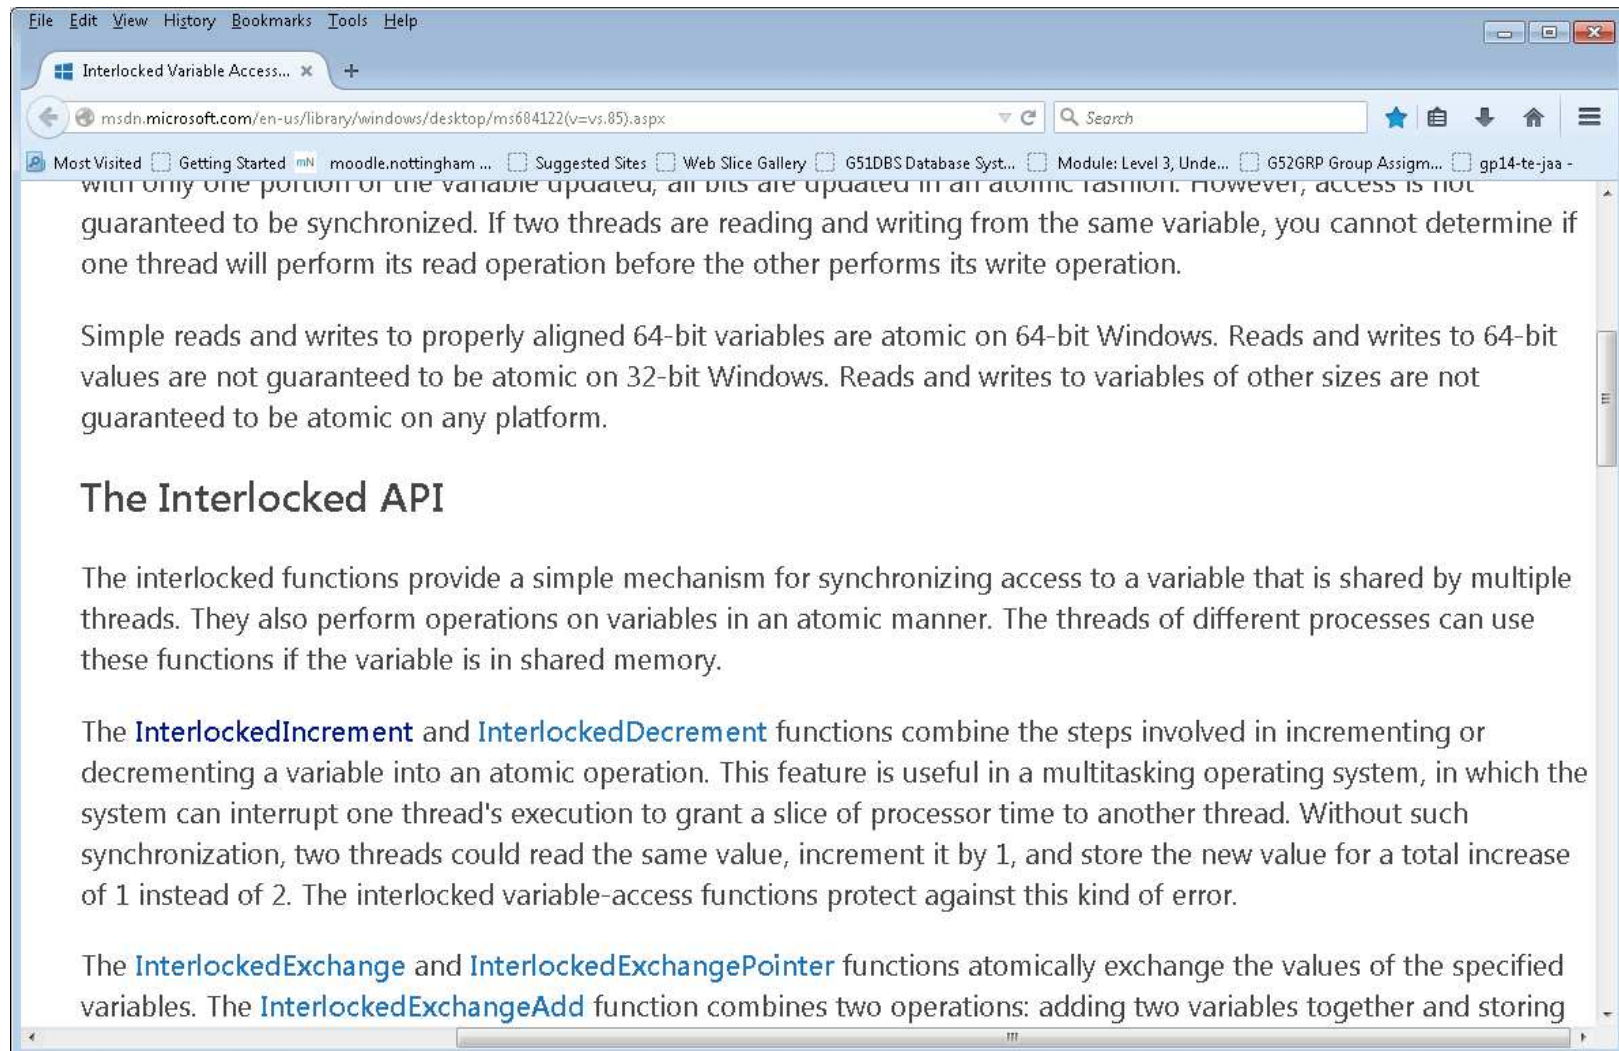
```
// Windows name for an unsigned long - 4 bytes
volatile DWORD dwTotal = 0;

// Function called by every thread
DWORD WINAPI thread_function( LPVOID lParam )
{
    for ( int i = 0; i < 1000000; i++ )
    {
        InterlockedIncrement( &dwTotal );
    }
    return 0;
}
```

Compiler-specific 'atomic' increment operation

Aside: Interlocked API

- <http://msdn.microsoft.com/en-us/library/windows/desktop/ms684122%28v=vs.85%29.aspx>



Problems with atomic operations


- Limited set of operations available
 - E.g. simple increment/decrement, test & set, compare and swap, etc
- What if you want to do something more complex?
- Need a way to prevent two things doing something at the same time, regardless of what that thing is or how long it will take...

Critical sections

- Key element: no other thread must do something which affects what we do in the critical section
 - e.g. we must not let anything else get at the data while we are half-way through a change
- Obvious solutions:
 - Only do one operation – then either done or not, cannot be interrupted (atomic operations)
 - Take it in turns – check that no other thread will access the data while you want it
 - May be fair or not

Mutex / Critical Sections

- Simple mutual exclusion:
 - Mutually exclude each other – only one at a time
- We can think of the code having a 'critical section' that only one thread can run at once
- Consider a generic structure like this:
 1. Initial code (init) – before critical section
 2. Enter critical section – **apply protocol here**
 3. Critical section code (crit)
 4. Exit critical section – **apply protocol here**
 5. Remainder (rem) – after critical section



Simple round-robin

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

Exit protocol:

turn = 2;

rem

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

crit

Exit protocol:

turn = 1;

rem

Simple round-robin : trace (1)

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

- Thread 2:

init

Simple round-robin : trace

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

Entry protocol:

while (turn != 1) ;

- Thread 2:

init

Entry protocol:

while (turn != 2) ;

BLOCKED – effectively stops
Will not move until variable
***turn* changes value**

Simple round-robin : trace

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

Still blocked

Simple round-robin : trace

Variable: *turn*: **just got set to 2**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

Exit protocol:

turn = 2;

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

Just got unblocked

Simple round-robin : trace

Variable: *turn*: has value 2

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

Exit protocol:

turn = 2;

rem

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

crit

Simple round-robin : trace

Variable: *turn*: has value 2

- Thread 1:

init

Entry protocol:

while (turn != 1) ;

crit

Exit protocol:

turn = 2;

rem

- Thread 2:

init

Entry protocol:

while (turn != 2) ;

crit

**Assume that the
crit takes a while**

Simple round-robin : trace

Variable: *turn*: has value 2

- Thread 1:

init

Entry protocol:

while (turn != 1) ;

**Now this one is blocked
Cannot get any further
yet**

- Thread 2:

init

Entry protocol:

while (turn != 2) ;

crit

**Assume that the crit
is still taking a while**

Simple round-robin : trace

Variable: *turn*: **just got set to 1 again**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

**Eventually crit finishes
in thread 2 and thread
1 can continue**

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

crit

Exit protocol:

turn = 1;

Simple round-robin : trace

Variable: *turn*: has value 1

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

and so on...

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

crit

Exit protocol:

turn = 1;

rem

Simple round-robin

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

Exit protocol:

turn = 2;

rem

Questions to consider:

- When is each thread keeping the CPU busy?
- Is this process fair across the threads?
 - Does each get a chance to act?
 - Does one get more actions than another?
- What happens if one thread stops?
- Can both be blocked at once?
- Can this be extended to more than 2?

Full round-robin algorithm

Variable: *turn*: integer variable, initialised to 1, **volatile**

- Thread 1:

init

Entry protocol:

while (*turn* != 1) ;

crit

Exit protocol:

turn = 2;

rem

- Thread 2:

init

Entry protocol:

while (*turn* != 2) ;

crit

Exit protocol:

turn = 1;

rem

Spin locks

- Sit in a tight loop waiting for a variable to take specific values
- Variable usually needs to be volatile
 - No point doing this if another thread is not going to alter it
- Thread will always be busy
 - Constantly checking the value
 - Wastes a lot of CPU time
- In the previous example they both had to take it in turns – we could do better...

Next lecture

- Mutual Exclusion
- Better entry/exit protocols for ensuring mutual exclusion and protecting critical sections