

G520SC

OPERATING SYSTEMS AND

CONCURRENCY

Semaphores

Dr Jason Atkin

Last Lectures

- Critical sections and Mutual Exclusion
 - Spinlocks
 - Round-robin
 - Decker's Algorithm
 - Peterson's Algorithm
 - Busy-wait consumed CPU time
- System support
 - CRITICAL_SECTION objects
 - Mutex objects
 - Owner by the thread – no re-entry problems
 - Actually 'sleep' and 'wake' the thread/process

This lecture

- Semaphores
 - Locks with a counting mechanism
- The theory of semaphores
- Using semaphores
- Problems where semaphores can help
 - Producer - consumer

Semaphore

- A counting mechanism
 - A counter and two atomic operations upon it
- Could think of it as a limited resource which threads can ask for
 - Wait(), p() : try to get some of the resource
 - Signal(), v() : indicate that you have finished with the resource
- Differences from mutexes:
 - Counter, not just a binary flag for locked or not
 - Not owned by the thread
 - wait()/p() twice decrements by two, can deadlock yourself
- Binary semaphore : only one can have it

More formally: Semaphores

- A *semaphore* s is an integer variable
 - which can take only non-negative values
- Once it has been given its initial value, the only permissible operations on s are the atomic actions:
- $P(s)$: if $s > 0$ then $s = s - 1$,
else **suspend** execution of the process that called $P(s)$
 - Sometimes called wait() or lock()
- $V(s)$: if some process p is suspended by a previous $P(s)$ on this semaphore then resume p , else $s = s + 1$
 - Sometimes called signal() or release()
- A *general semaphore* can have any non-negative value
- A *binary semaphore* is one whose value is always 0 or 1

P and ***V*** as atomic actions

- Reading and writing the semaphore value needs to be in a *critical section*, or atomic action:
 - ***P*** and ***V*** operations on the same semaphore must be *mutually exclusive*
 - e.g., suppose we have a semaphore, *s*, which has the value 1, and two processes simultaneously attempt to execute ***P*** on *s*:
 - only one of these operations will be able to complete before the next ***V*** operation on *s*;
 - the other process attempting to perform a ***P*** operation is suspended.
 - Semaphore operations on different semaphores do not need to be mutually exclusive

V() on binary semaphores

- Binary semaphores have values 0 or 1
- Effects of performing a **V** operation on a binary semaphore which has a current value of 1 are implementation dependent:
 - operation may be ignored
 - may increment the semaphore
 - may throw an exception
- We will assume that a **V** operation on a binary semaphore which has value 1 does not increment the value of the semaphore.
- Note: Windows applies a maximum on any semaphore, set when created

Example implementation

Integer Variable: SEM

Initialisation: SEM = <Initial count>

Implementation of p()

Lock Critical Section

If SEM > 0

SEM = SEM - 1

Unlock Critical Section

Return success

Unlock Critical Section

Spin, wait for SEM > 0 **and**
repeat (or return failure)

Implementation of v()

Lock Critical Section

SEM = SEM + 1

Unlock Critical Section

Note: With atomic
increment/decrement, the
critical section is implicit

Resuming suspended processes

- Note that the definition of V doesn't specify which process is woken up if more than one process has been suspended on the same semaphore
- This has implications for the fairness of algorithms implemented using semaphores and upon properties like Eventual Entry

Windows support : CreateSemaphore

- MSDN, Using Semaphores: <https://msdn.microsoft.com/en-us/library/windows/desktop/ms686946%28v=vs.85%29.aspx>

- Create/access the object and store the handle

```
HANDLE ghSemaphore;  
ghSemaphore = CreateSemaphore(  
    NULL,                // default security attributes  
    MAX_SEM_COUNT,       // initial count  
    MAX_SEM_COUNT,       // maximum count  
    NULL);               // unnamed semaphore
```

- Use the semaphore
 - WaitForSingleObject(), ReleaseSemaphore()
- Close the handle at the end

```
CloseHandle(ghSemaphore);
```

Windows support : wait and release

```
switch (WaitForSingleObject(  
    ghSemaphore, // semaphore handle  
    10000L) )    // Timeout - could be zero  
{  
    case WAIT_OBJECT_0:  
        printf("Thread %d: wait success\n", GetCurrentThreadId());  
        if ( !ReleaseSemaphore(  
            ghSemaphore, // handle to semaphore  
            1,           // increase count by one  
            NULL) )      // not interested in previous count  
        { /* Handle error */ }  
        break;  
  
    case WAIT_TIMEOUT:  
        printf("Thread %d: wait timed out\n", GetCurrentThreadId());  
        break;  
}
```

Maximum simultaneous executions

Using semaphores for maximum counts

- You could use a binary semaphore as a simple method to ensure mutual exclusion
 - Only one thread can have it at once
- A semaphore with initial value of n can ensure that only n threads can access a section of code (or resource) at once
 - Each calls $p()$ before it enters (or needs the resource) and $v()$ when it finishes with it

Using a semaphore to protect a resource

Semaphore Count: 2

T1

→ init

Loop:

P()

Use res

V()

rem()

T2

→ init

Loop:

P()

Use res

V()

rem()

T3

→ init

Loop:

P()

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 2

T1

init

Loop:

→ P()

Use res

V()

rem()

T2

→ init

Loop:

P()

Use res

V()

rem()

T3

→ init

Loop:

P()

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 1

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

→ init

Loop:

P()

Use res

V()

rem()

T3

→ init

Loop:

P()

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 1

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

→ P()

Use res

V()

rem()

T3

→ init

Loop:

P()

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

→ init

Loop:

P()

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P() Blocked

Use res

V()

rem()

T4

→ init

Loop:

P()

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P() Blocked

Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

Use res

→ V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P() Blocked

Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 1

T1

init

Loop:

P()

Use res

→ V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P() Blocked

Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 1

T1

init

Loop:

P()

Use res

V()

→ rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P()

Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Note: you don't know which one will get unblocked. T3 here.

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

Use res

V()

→ rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

→ Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

→ P()

Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

→ Use res

V()

rem()

T4

init

Loop:

→ P()

Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

→ P() Blocked

Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

→ Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

→ P() Blocked

Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

Use res

→ V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 1

T1

init

Loop:

→ P()

Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

Use res

V()

→ rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Note: you don't know which one will get unblocked. T1 here.

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

P()

Use res

V()

→ rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()

Using a semaphore to protect a resource

Semaphore Count: 0

T1

init

Loop:

P()

→ Use res

V()

rem()

T2

init

Loop:

P()

→ Use res

V()

rem()

T3

init

Loop:

→ P() Blocked

Use res

V()

rem()

T4

init

Loop:

→ P() Blocked

Use res

V()

rem()


Etc... The number of threads between the p() and v() is limited

Consumer-Producer

Consumer-Producer

- Assume an unlimited queue

Items in queue: 0




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 1




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 1




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 1




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 2




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 3




- Consumer cannot consume a product until it has been produced
- How can we use a (single) semaphore to prevent consumption before production?

Consumer-Producer

- Assume an unlimited queue

Items in queue: 2



- Producer creates the items
- Consumer cannot consume a product until it has been produced – **how do we stop this?**

Answer...

Answer

- Use a semaphore as a counter
- When producer produces an item, it increments the semaphore (`v()`, `signal()`)
 - Like 'release' on semaphore
- Consumer uses `p()` to consume the item
 - Blocks if count is zero
 - Decrements by 1 if the item is consumed – tracking the items still left to be consumed

Limiting the queue size


Limited queue size

- The previous example assumed an unlimited buffer size
- If we have a buffer size of n , how can we make the producer stop when the buffer gets full, to allow the consumer to empty it a bit?

Rolling buffer

- Assume a limited queue

Items in queue: 0




- When queue is filled, it rolls around, as it gets emptied

Rolling buffer

- Assume a limited queue


Items in queue: 1



Rolling buffer

- Assume a limited queue


Items in queue: 2



Rolling buffer

- Assume a limited queue


Items in queue: 3



Rolling buffer

- Assume a limited queue


Items in queue: 4



Rolling buffer

- Assume a limited queue


Items in queue: 5



Rolling buffer

- Assume a limited queue


Items in queue: 6



Rolling buffer

- Assume a limited queue

Items in queue: 6




- Assume that consumer finally got around to consuming one

Rolling buffer

- Assume a limited queue

Items in queue: 7




- Producer cannot now produce any more until the consumer has used some up
- There are no spare spaces to put the values

Rolling buffer

- Assume a limited queue

Items in queue: 6




- Consumer uses one, so there is one spare space
- Producer could continue now

Rolling buffer

- Assume a limited queue

Items in queue: 5




- Consumer uses another, so there are two spare spaces
- Producer could continue still

Rolling buffer

- Assume a limited queue

Items in queue: 4



- Consumer uses another, so there are three spare spaces now
- Producer could continue still

Question: Semaphore solution?

- Is there any way to count whether the producer can continue?
- Hint: To use a semaphore you want something that will drop to 0 when it has to stop, and be positive when it can continue

Answer

Answer...

- To use a semaphore you want something that will drop to 0 when it has to stop, and be positive when it can continue?
 - The number of spare spaces has this property
- Initially the number of spare spaces is the size of the buffer
- As items are produced, the spaces go down by one
- As items are consumed, the spaces go up again, by one at a time

Producer-Consumer Semaphores

Producer

- SpacesSemaphore.p()
 - Decrement spaces if there is at least one, or block
- Put item into buffer
- ProductSemaphore.v()
 - Record that an item is ready for production

Consumer

- ProductSemaphore.p()
 - Decrement number of product if there is one
- Remove item from buffer
- SpacesSemaphore.v()
 - Record that an extra space is now available

Two semaphores:

SpacesSemaphore: Initial = number of spaces = buffer size

ProductSemaphore: Initial = 0 (no items initially)

Summary

- Semaphores are useful for implementing counters
- Always count DOWN to ZERO
- Block if attempting to get a lock / call p() / call wait() if the counter is already zero
- You can implement them using any mutual exclusion code or atomic operations
 - Need access to the counter variable from all threads though
- Often built-in operating system support
- No guarantee of which waiting thread will be awoken when the count goes up
 - Note: Java implementation supports first-come-first-served too
- Consumer-producer model is easy to enforce using semaphores – unlimited or limited buffer space

Next lecture

- Deadlock and livelock situations
- Inter-process communication
- Synchronous vs asynchronous communication
- Window messages again
- Sockets