

# **G52OSC OPERATING SYSTEMS AND CONCURRENCY**

## Interprocess Communication

Dr Jason Atkin

# This lecture

- Locks and deadlocks
- Inter-process communication
  - Windows messages
    - Synchronous
    - Asynchronous
  - Sockets

# Reminder: Test-and-Set instruction

The Test-and-Set (atomic) instruction effectively executes the function

```
bool TS(bool lock)
{
    bool v = lock;
    lock = true; // Set true
    return v;    // Old lock value
}
```

**Stops the other thread looking at it before we set it**

See InterlockedExchange: <https://msdn.microsoft.com/en-us/library/windows/desktop/ms683590%28v=vs.85%29.aspx>

# Reminder: Test-And-Set spin-lock

```
// Process 1
```

```
init1;  
while(true) {  
    while(TS(lock)) ;  
    crit1;  
    lock = false;  
    rem1;  
}
```

```
// Process 2
```

```
init2;  
while(true) {  
    while(TS(lock)) ;  
    crit2;  
    lock = false;  
    rem1;  
}
```

Key point: this worked to avoid interference from other threads

# Two-locks, two processes

*// Process 1*

```
init1;
while(true) {
    while(TS(lock1)) ;
    crit1a;

    while(TS(lock2)) ;
    crit1b;
    lock2 = false;

    lock1 = false;
    rem1;
}
```

*// Process 2*

```
init1;
while(true) {
    while(TS(lock1)) ;
    crit2a;

    while(TS(lock2)) ;
    crit2b;
    lock2 = false;

    lock1 = false;
    rem1;
}
```

- Two-processes Test and set spinlock – does it still work?

# What about this case?

*// Process 1*

```
init1;
while(true) {
    while(TS(lock1)) ;
    crit1a;

    while(TS(lock2)) ;
    crit1b;
    lock2 = false;

    lock1 = false;
    rem1;
}
```

*// Process 2*

```
init1;
while(true) {
    while(TS(lock2)) ;
    crit2a;

    while(TS(lock1)) ;
    crit2b;
    lock1 = false;

    lock2 = false;
    rem1;
}
```

Any potential problems here?

# Livelock! Perhaps 100% CPU used

// Process 1

```
init1;  
while(true) {  
    while(TS(lock1)) ;  
    crit1a;
```

```
→ while(TS(lock2)) ;  
    crit1b;  
    lock2 = false;  
  
    lock1 = false;  
    rem1;  
}
```

// Process 2

```
init1;  
while(true) {  
    while(TS(lock2)) ;  
    crit2a;
```

```
→ while(TS(lock1)) ;  
    crit2b;  
    lock1 = false;  
  
    lock2 = false;  
    rem1;  
}
```

Why did this happen? Can we avoid it?

# Livelock: potential solution

// Process 1

```
init1;
while(true) {
    while(TS(lock1)) ;
    crit1a;

    while(TS(lock2)) ;
    crit1b;
    lock2 = false;

    lock1 = false;
    rem1;
}
```

// Process 2

```
init1;
while(true) {
    while(TS(lock1)) ;
    while(TS(lock2)) ;
    crit2a;
    while(TS(lock1)) ;
    crit2b;
    lock1 = false;

    lock2 = false;
    rem1;
}
```

**Simple trick: always lock in the same order**



# Self-livelock example

```
init1;  
while(true) {  
    while(TS(lock1)) ;  
    crit1a;  
    func();  
    lock = false;  
    rem1;  
}
```

---

```
int func()  
{  
    while(TS(lock1)) ;  
    crit1b;  
    lock = false;  
}
```

- Some function gets a lock so that it can go into a critical section
- While in there it needs to do some other (minor?) thing, so calls a function to do so
- The function it calls needs to alter some resource which needs protecting from interference between threads
- So needs to be put into a critical section
- So the function attempts to get a lock before modifying the value
- The test-and-set lock will spin forever

# Deadlock and Self-deadlock

- Exactly the same case can happen with deadlock as for livelock
  - Use a lock which blocks the thread rather than a spinlock
- Self-deadlock is often harder to achieve
  - Mutual exclusion objects often check the owner for this reason
- Two (or more) process deadlock is as easy to achieve accidentally as with livelock
- Consider using “WaitForMultipleObjects”

# Self-deadlock

```
init1;
while(true) {
    LOCK
    crit1a;
    func();
    crit1c;
    UNLOCK
    rem1;
}
```

---

```
int func()
{
    LOCK
    crit1b;
    UNLOCK
}
```

- Lock mechanism tries to get a lock in main function
- And asks again in the sub-function
- **Locks that know about ownership:**
  - E.g. CRITICAL\_SECTION or Mutex objects in windows
  - Lock automatically gained
  - Have to match locks and unlocks to avoid unlocking it too early
    - Before *crit1c* in this example
- **Locks with no concept of ownership:**
  - E.g. binary semaphore
  - Thread deadlocked at second p()

# Inter-process communication

# Interprocess Communication So Far

- We have so far seen various methods of interprocess communication
- Windows messages
  - PostMessage, SendMessage
- Shared data
  - Map the memory into both address spaces
  - Make it volatile
- Mutex locking
- Semaphore counts

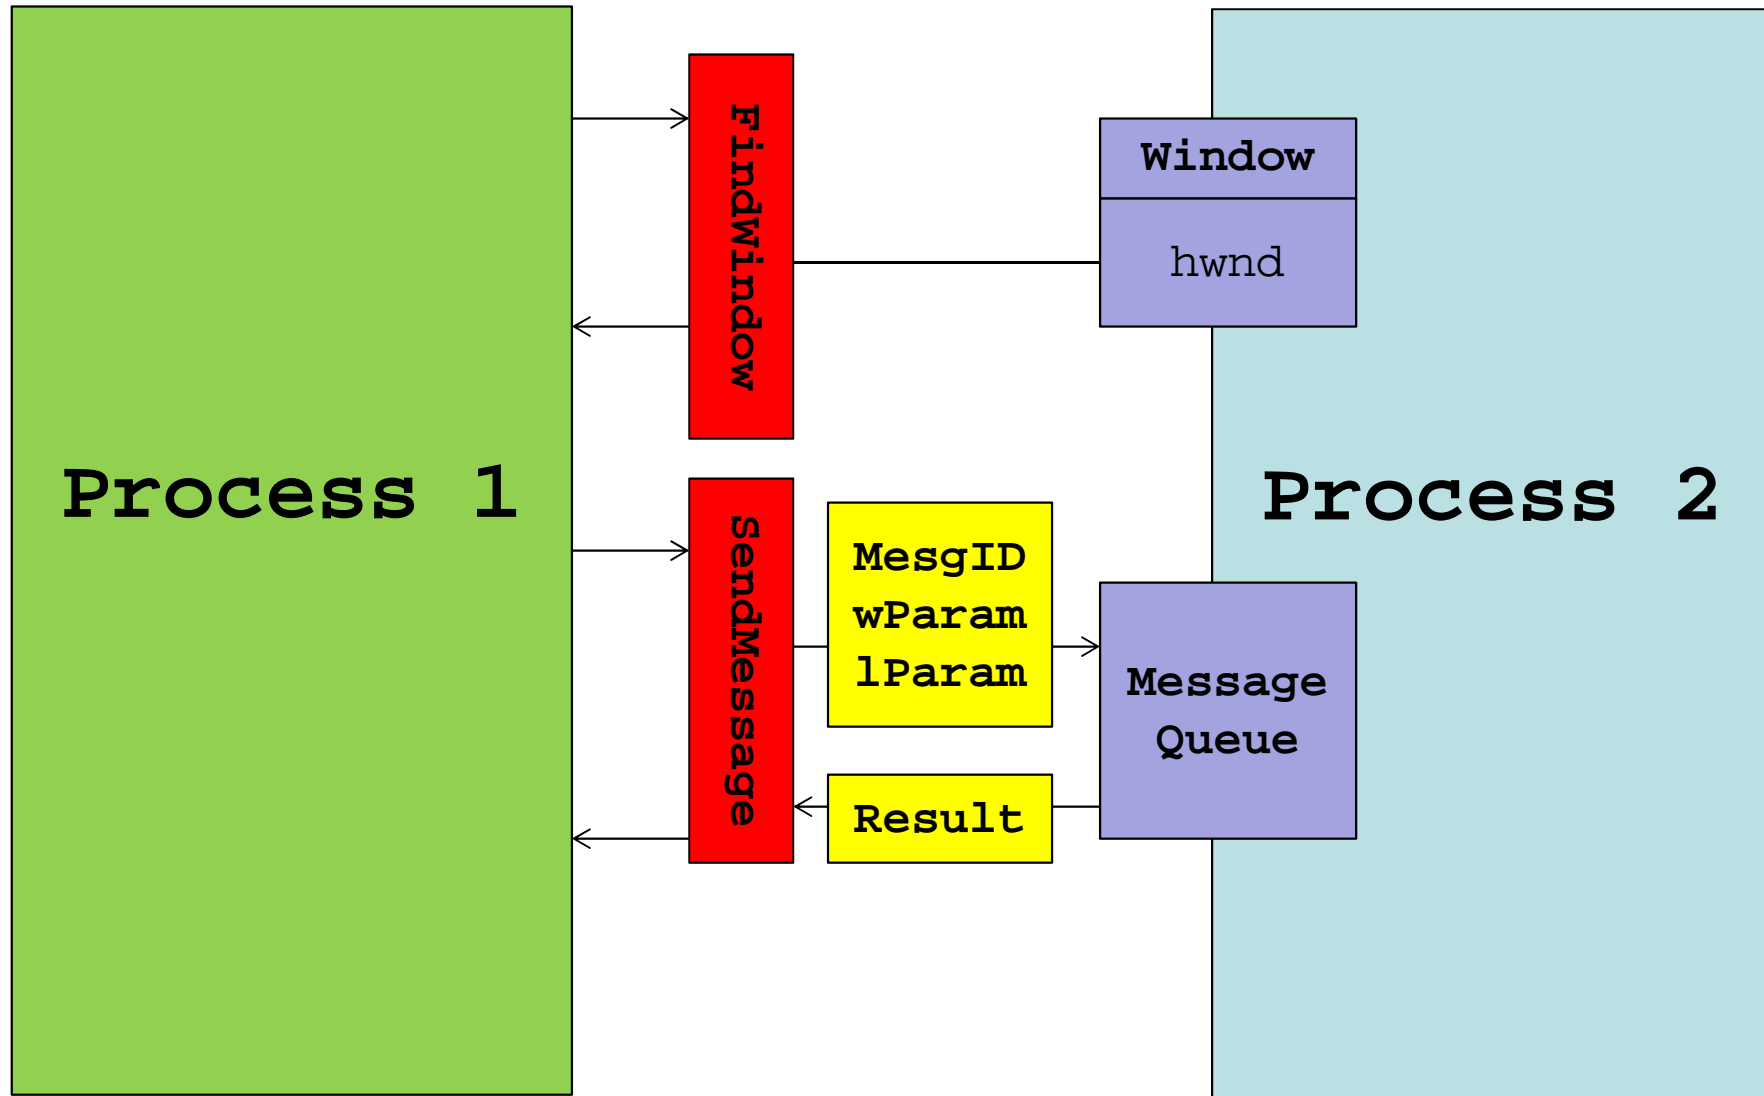
# Inter Process Communication

- Two conceptual types
- Asynchronous communication
  - Leave a message
  - It may get back to you later (or not)
- Synchronous communication
  - Ask a question, get a response
  - Make a request to do something, wait for it to happen

# Windows Messages for communication

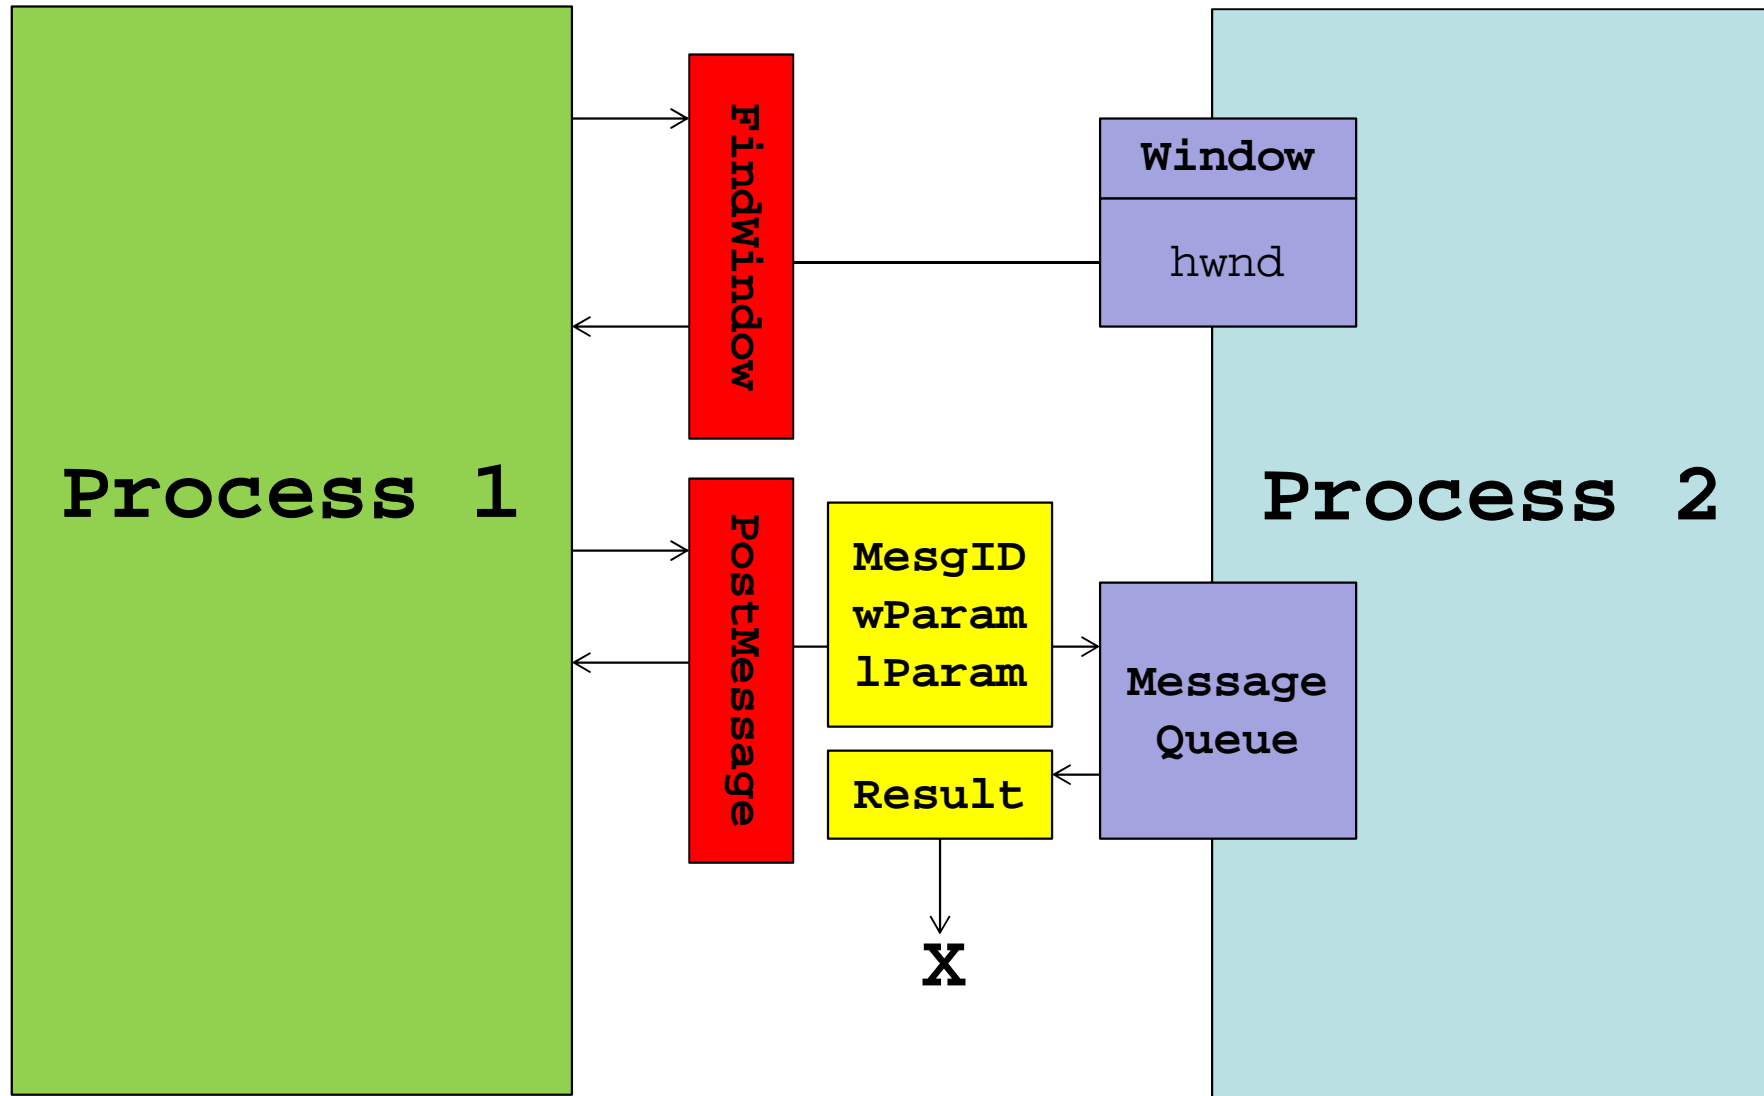
- SendMessage for synchronous communication
- Sends a message to a window and will wait for the thread which owns the window to handle the message
  - And to send a response back
- SendMessage blocks the thread which calls it until the message is handled
- If the caller thread owns the window, it will call the window procedure directly

# SendMessage

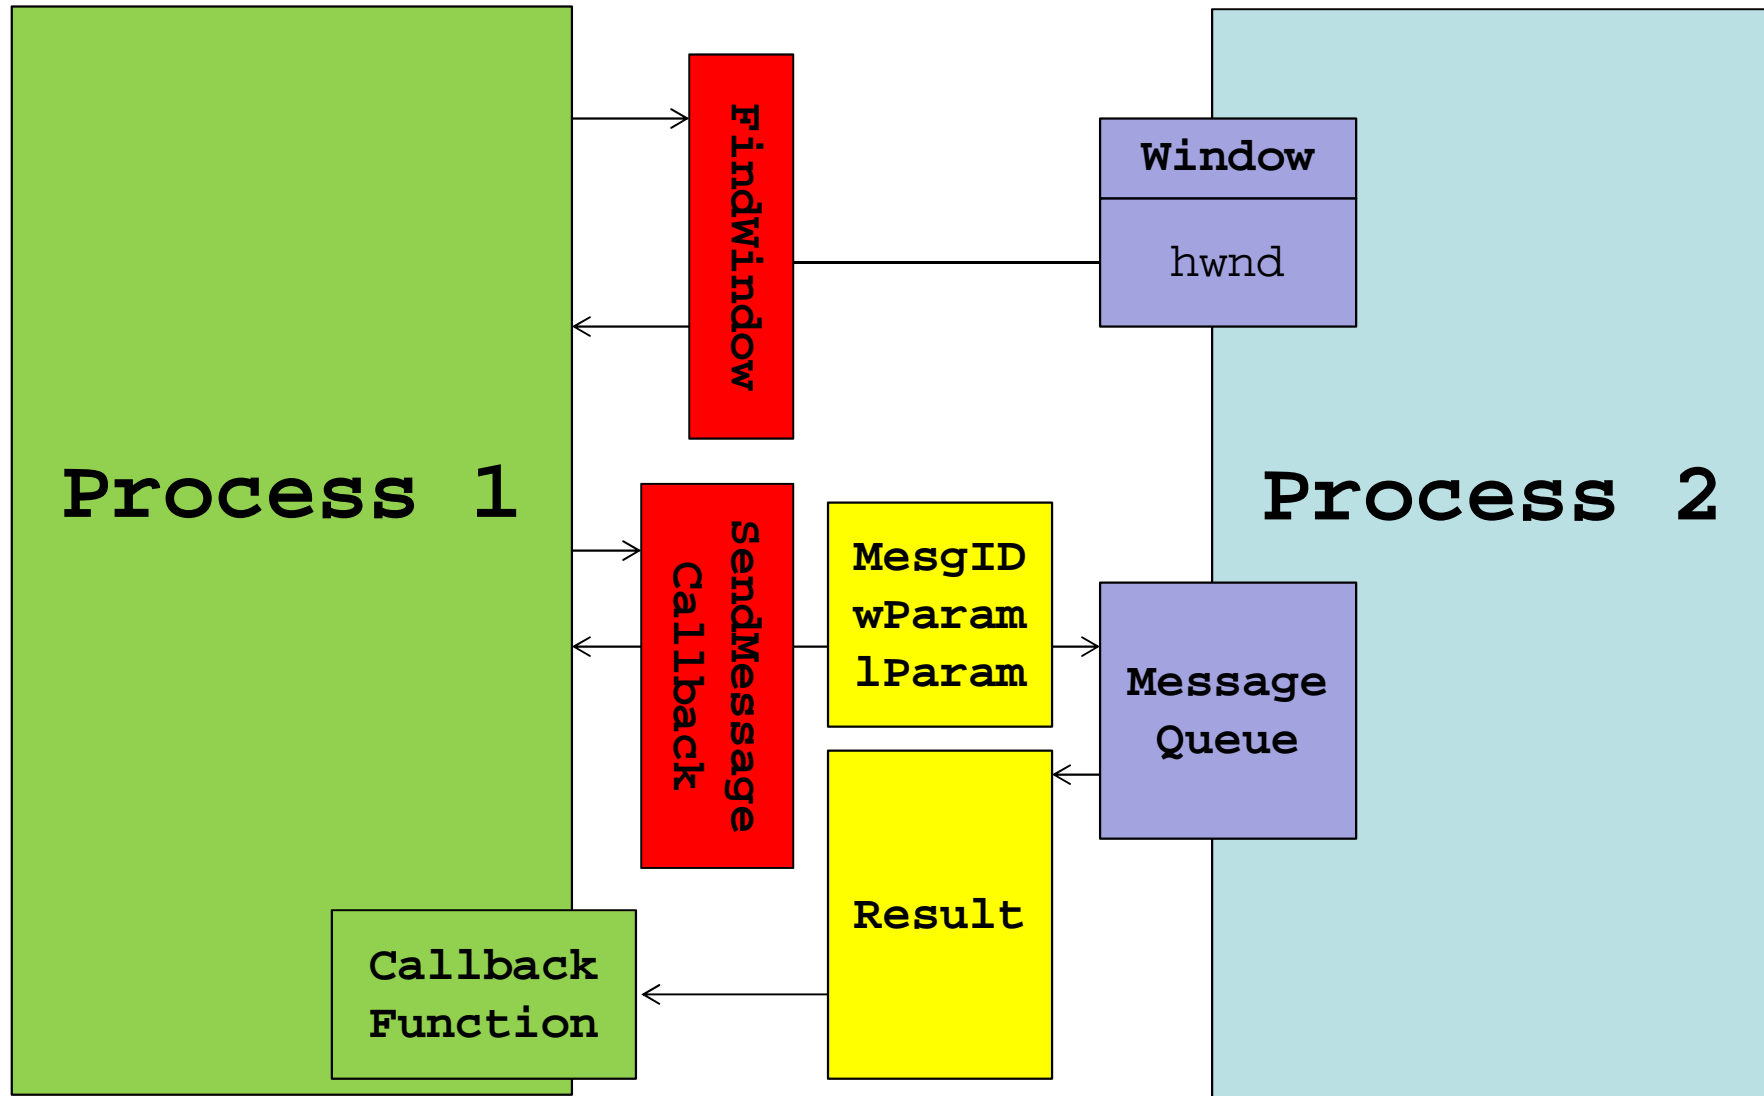




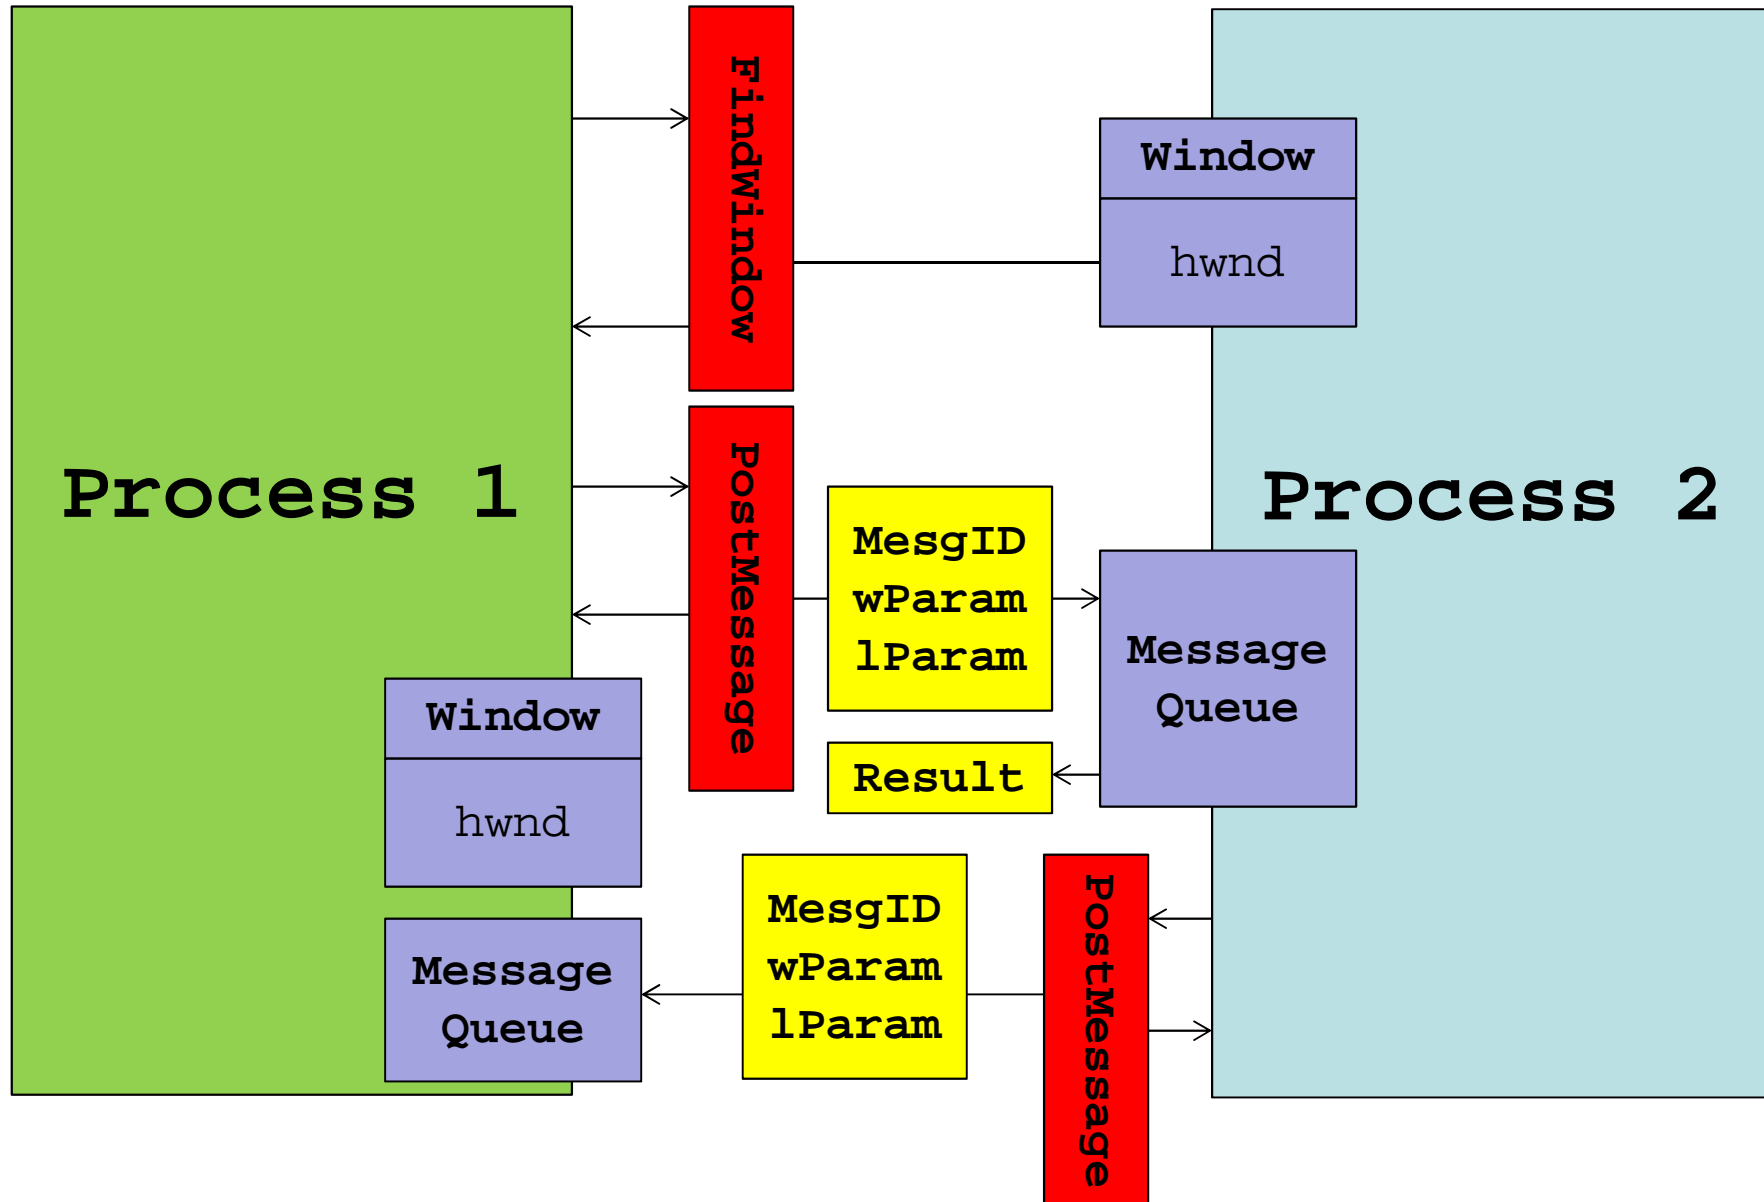
# PostMessage



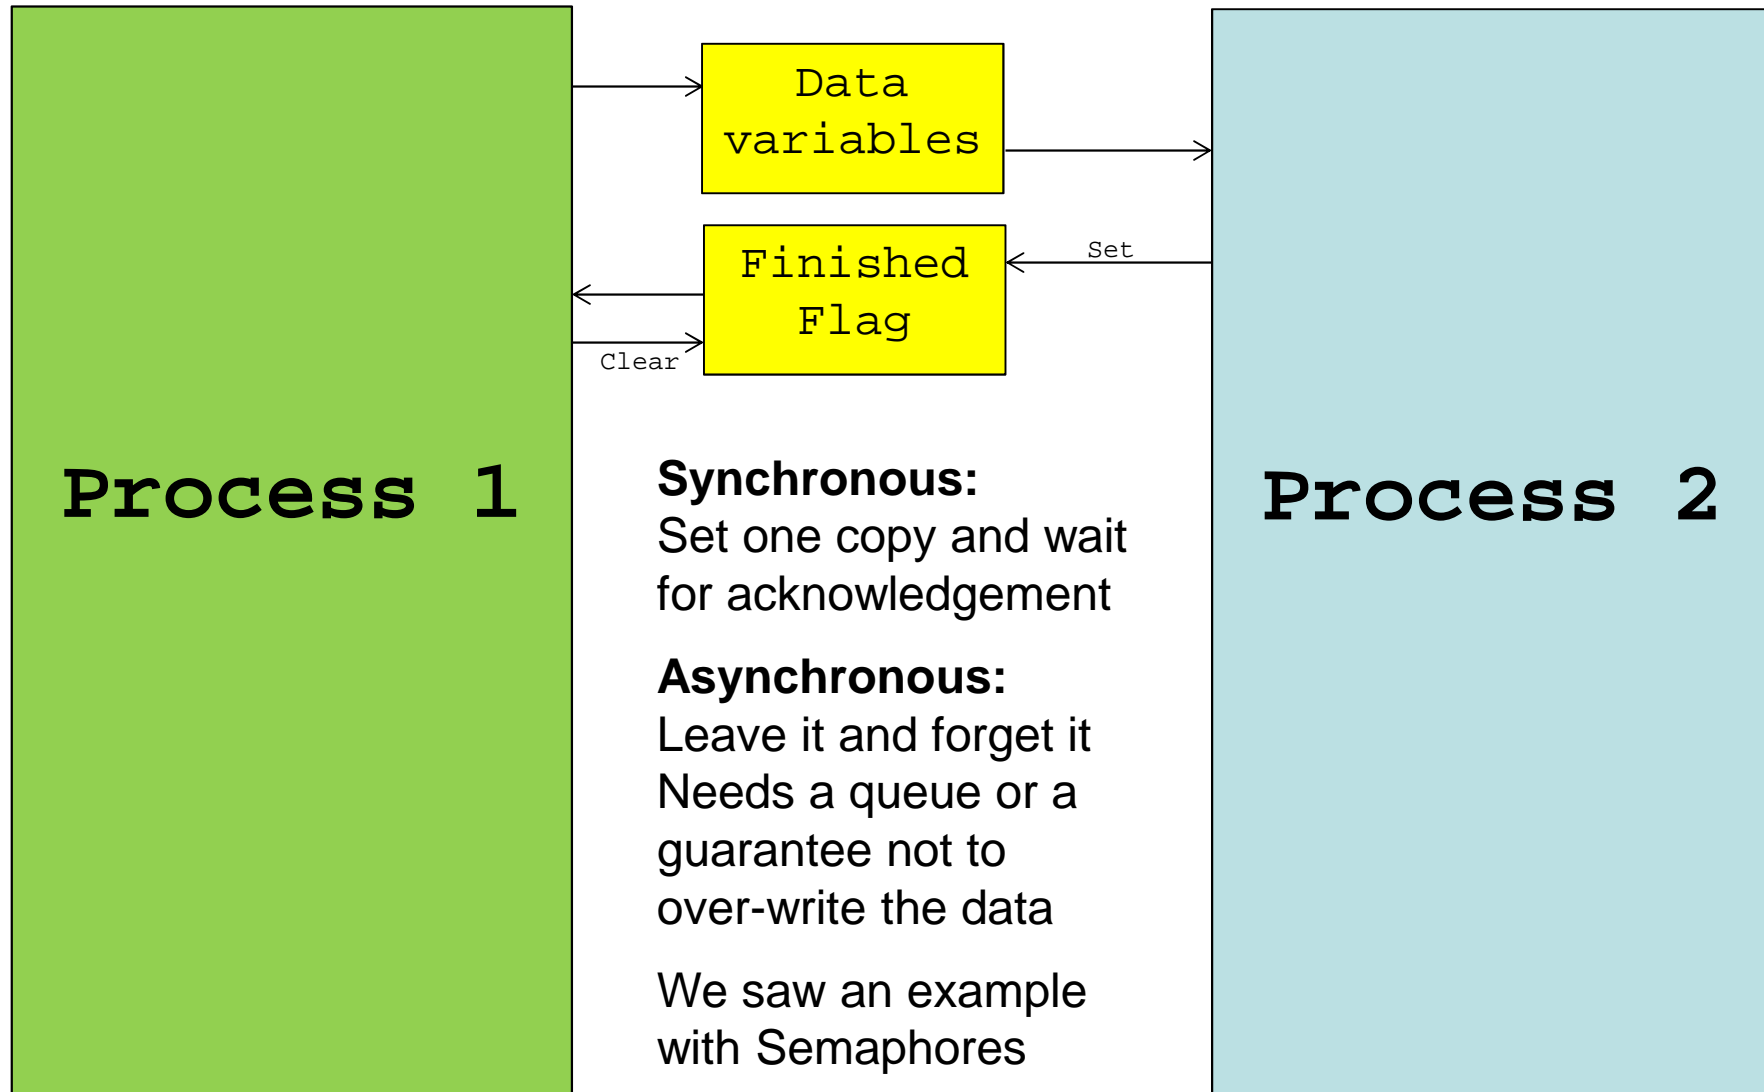
# SendMessageCallback



# Two-way PostMessage



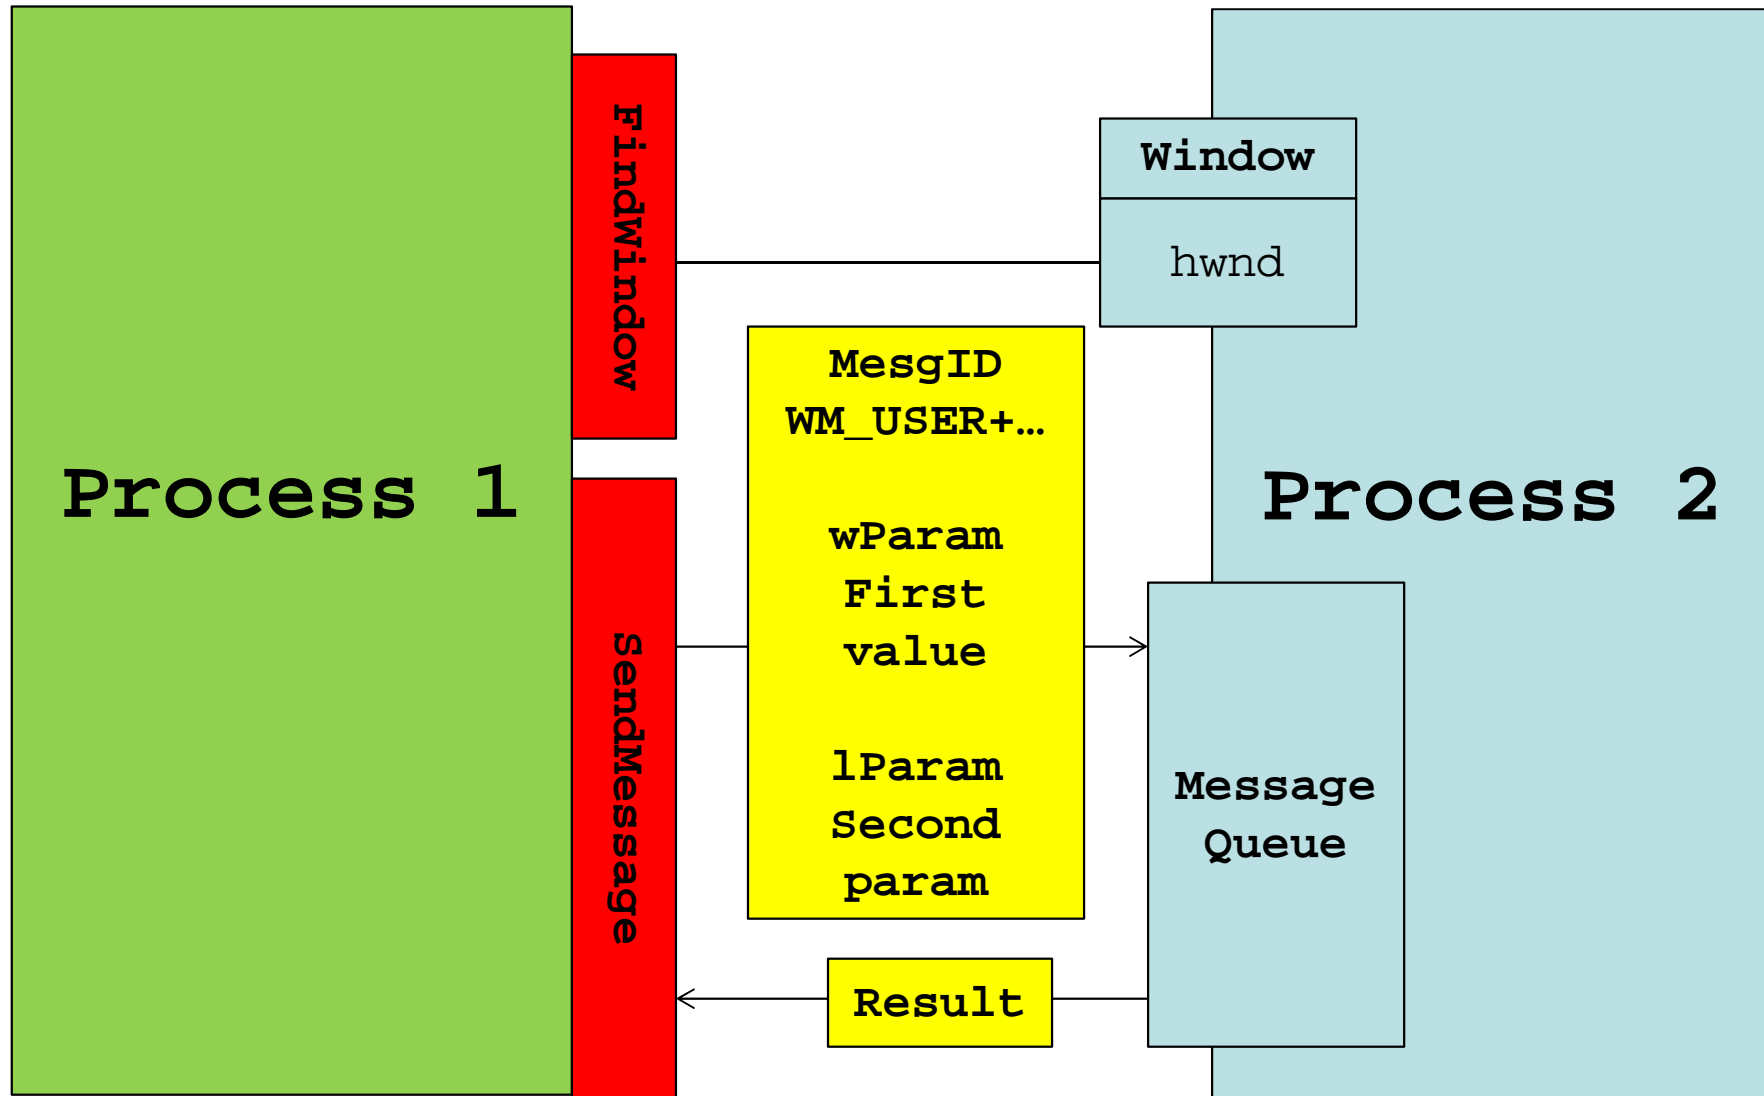
# Using shared memory to communicate



# Options for communications

- Synchronous
  - Easiest if you need a response, but caller has to wait (doing nothing/blocking)
- Asynchronous
  - Easy if you do **not** need a response
  - Need a sure-delivery method to be sure it is received
  - ‘Message’ must exist until the receiver handles it
  - Similarities to producer-consumer
    - Producer creates things which the consumer later uses
    - Things have to stay on the queue until picked up
  - Windows message queues meet this criterion
    - Limitations on payload though
  - MessageQueues exist on other OSes as well

# Example



# Example : Caller

```
// Name matches what I used for a class name
HWND hwnd = FindWindow( "MyClassName",NULL );

while ( TRUE )
{
    printf( "Type two numbers with a space between:" );
    if ( scanf( "%d %d",&i,&j ) == 2 )
    {
        total = SendMessage( hwnd, WM_USER + 1, i, j );
        printf( "%d + %d = %d\n",i,j,total );
        PostMessage( hwnd,WM_USER + 2,total,0 );
    }
    ...
}
```

# Example: Server/receiver


```
LRESULT CALLBACK WndProc( HWND hwnd, UINT msg,
    WPARAM wParam, LPARAM lParam )
{
    switch ( msg )
    {
        ...
        case WM_USER+1: // Do a plus
            printf( "%d plus %d = %d\n",wParam,lParam,wParam+lParam);
            ...
            return wParam + lParam;
        case WM_USER + 2: // Set the caption
        {
            char buf[20]; sprintf( buf,"Calculated: %d",wParam );
            SetWindowText( hwnd,buf );
            return 0; // Not looking for this anyway
        }
    }
}
```



# Windows Sockets Introduction


# Sockets (Windows Sockets API)

- Stream-based
  - Even if underlying datagram-based
  - Sure-delivery – resends failed packets
- Numbered sockets (2 byte value)
- Server-side:
  - Listen for connections on a socket
  - Accept connections
    - Handle the connection (send/receive data)
  - Listen for next connection
- Client-side
  - Connect to socket on server (machine IP/socket #)
  - Send/receive data



# Windows Sockets : server-side

- **socket ( )**: allocates a socket handle to listen on
- **bind ( )**: grab a socket, associating a local address with it
- **listen ( )**: listen for incoming connections on the socket
- **accept ( )**: will block until a connection has occurred (or error)
  - **Start a new thread to handle connection?**
- **send ( ) , recv ( )**: send and receive data using the open connection
- **closesocket ( )**: close the socket, release resources



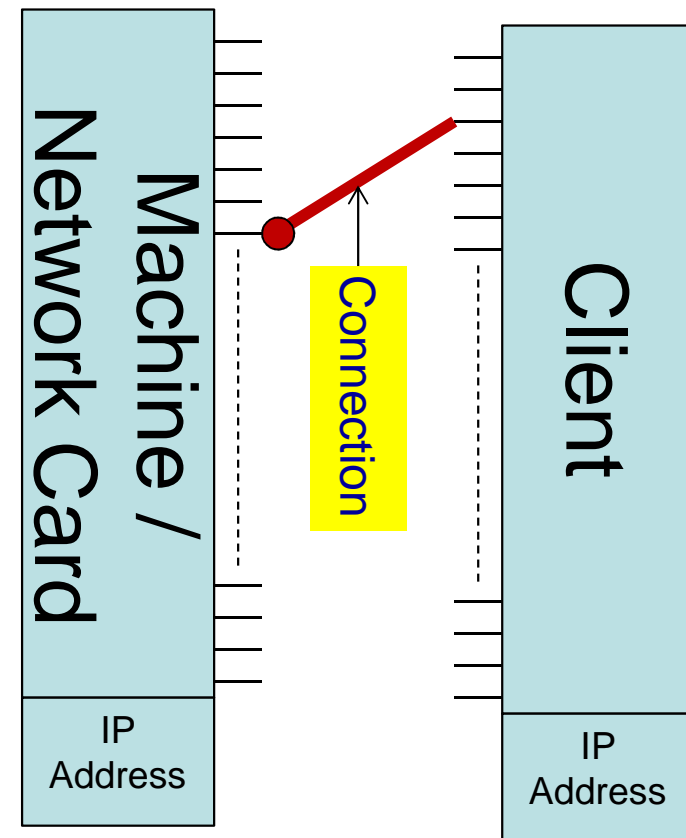
See the sample for an example of use, and for an overview, see:

<https://msdn.microsoft.com/en-us/library/windows/desktop/ms740673%28v=vs.85%29.aspx>


<https://msdn.microsoft.com/en-us/library/windows/desktop/ms741394%28v=vs.85%29.aspx>

# Windows Sockets : Client

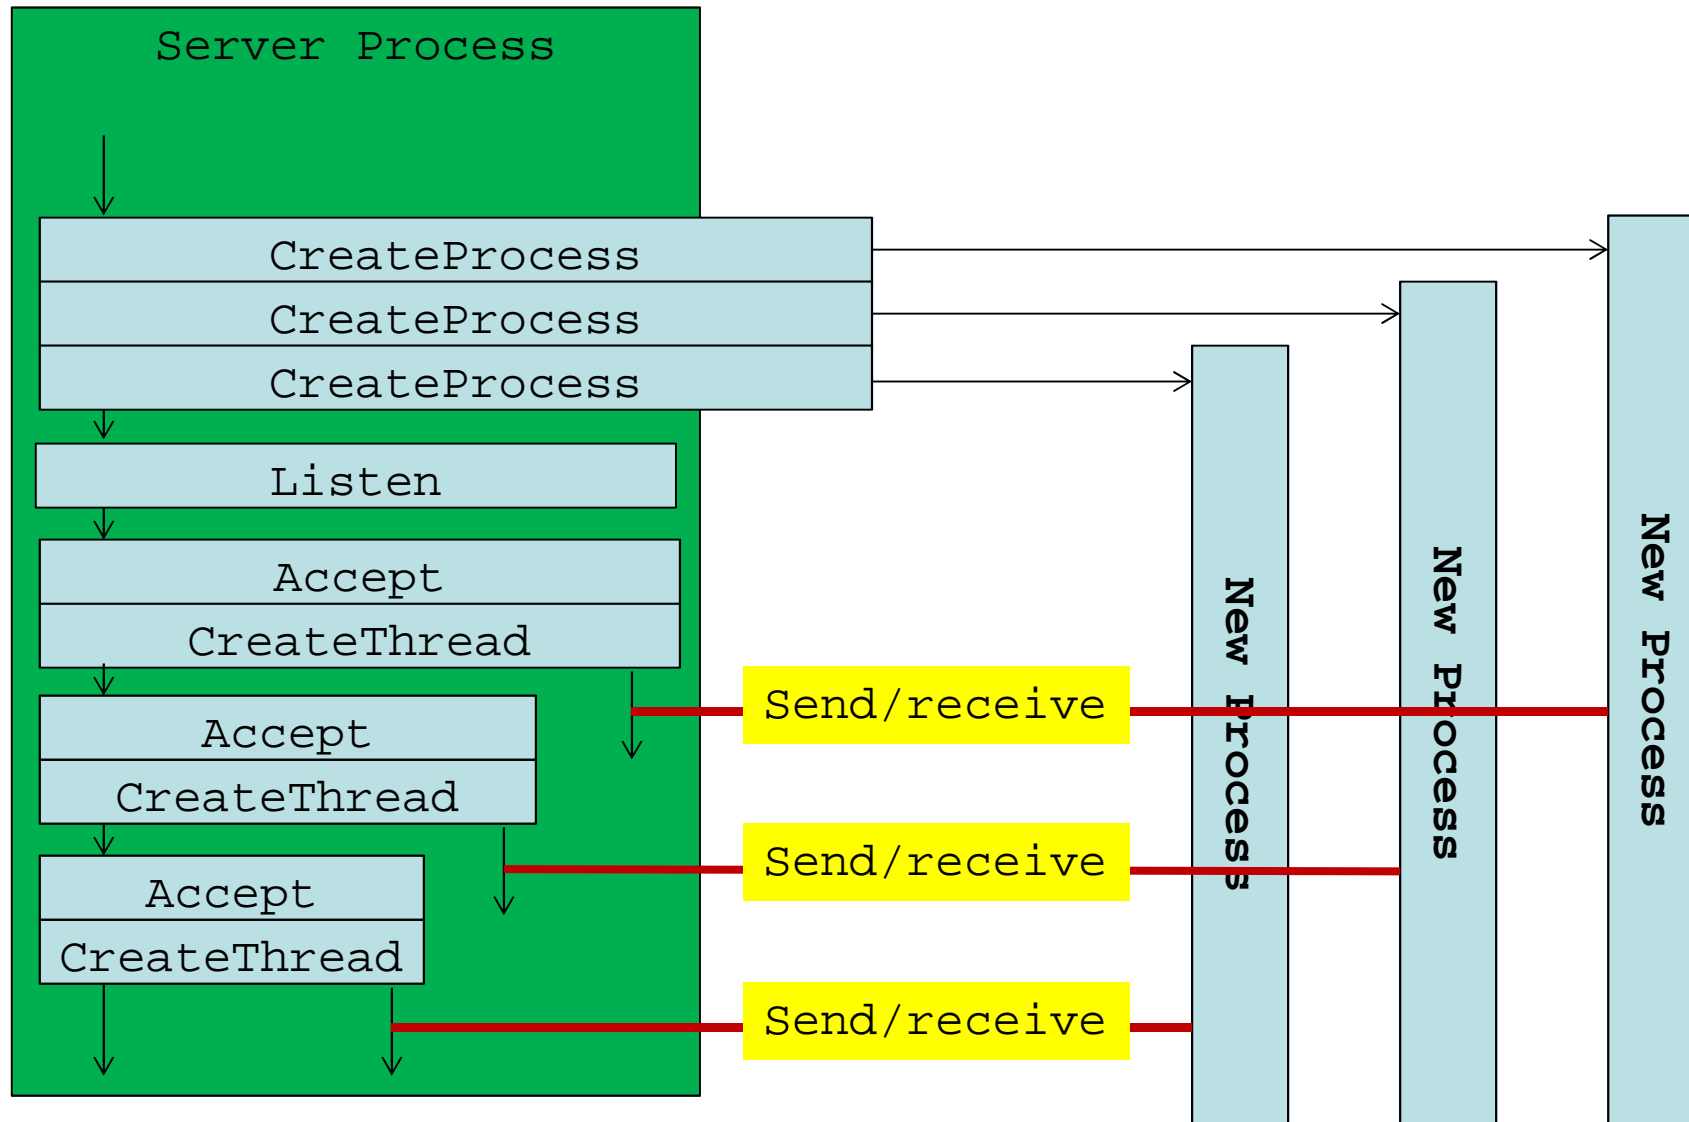
- **socket ( )**: Allocates a socket handle on local machine (protocol and address specified)
- **connect ( )**: Attempt to connect to server, supplying server address and port.  
Note: this will 'bind' an unbound local socket
- **send ( ) , recv ( )**: send and receive data
- **shutdown ( )**: shutdown send or receive operations. Can shutdown send to let the server know that you have finished
- **closesocket ( )**: close the socket, release resources



# Multiple connections possible



# Example program




# Sample code

1. Server opens socket and listens for connections
2. Client connects to server
3. Client sends and receives data
  - Client sends data
  - Server receives data from client
  - Server sends back whatever bytes it received
  - Client displays the messages
4. Eventually client has finished sending all data
  - Client calls shutdown to shutdown output
5. Server detects socket shutdown when it tries to read and closes the connection
6. Server thread dies

# Structured messages

- Streams of bytes are sent down the connection
- You can send structured data if you wish
- E.g.

```
struct DataClientToServer
{
    char op;
    int first;
    int second;
};
```



- Even if you send the whole struct in one go, it may be split up when you receive it
  - E.g. you may receive the first two bytes, then 5 more, then 2 more, etc.
- You may need to reconstruct your data before you interpret it
- Example code in lab 4 exercise 4 does this for you (copy-paste-adapt)



# Next lecture

- Concurrency in Java
- Monitors