

G520SC

OPERATING SYSTEMS AND

CONCURRENCY

Java Critical Sections and Monitors

Dr Jason Atkin

Last Lecture

- Java
 - Thread Creation
 - Anonymous classes
 - Memory model
 - Atomic variables
 - Volatile
 - Memory gates

Spin locks with `volatile`

```
// Process 1
init1;
while(true) {
    // entry protocol
    c1 = true;
    turn = 2;
    while (c2 && turn == 2);
    crit1;
    // exit protocol
    c1 = false;
    rem1;
}

// shared variables
bool c1 = c2 = false;
integer turn = 1;
```

```
// Process 2
init2;
while(true) {
    // entry protocol
    c2 = true;
    turn = 1;
    while (c1 && turn == 1);
    crit2;
    // exit protocol
    c2 = false;
    rem2;
}
```

Try the sample and you should find that it worked, without having to add any explicit memory barriers

Equivalent behaviour - optimising

```
// Process 1
```

```
init1;
```

```
while(true) {
```

```
    // entry protocol
```

```
    c1 = true;
```

```
    turn = 2;
    while (c2 && turn == 2);
```

```
    crit1;
```

```
    // exit protocol
```

```
    c1 = false;
```

```
    rem1;
```

```
}
```

```
// shared variables
```

```
bool c1 = c2 = false;
```

```
integer turn = 1;
```

```
// Process 1
```

```
init1;
```

```
while(true) {
```

```
    // entry protocol
```

```
    turn = 2;
```

```
    c1 = true;
```

```
    c1 = false;
```

```
    while (c2 && turn == 2);
```

```
    crit1;
```

```
    // exit protocol
```

```
    rem1;
```

```
}
```

Equivalent behaviour
In a single-threaded process

Volatile

- In C volatile stops caching of the variable in local memory
 - Forces the value to be written to memory and/or read from memory immediately
- In Java it does more
 - Prevents the local caching
 - AND tells compiler not to re-order the accesses to volatile variables
 - AND writes op codes into the program to stop the processor re-ordering the operations

Coursework and Labs

- Note: the four labs will take a while to finish
 - This is deliberate, but will make the coursework shorter
- A 15% coursework for a 20 credit module (equivalent 30% coursework for 10 credit) could take a while to complete
- I included the time to do the labs, which may take you longer than the coursework itself
- You learn how to do things in the labs, then apply your knowledge to the coursework (part 1)
 - **Learning and understanding takes time**
 - **Do not leave doing the labs until the last minute!!!**

This Lecture

- Fairness and spin locks
- Synchronisation
 - Locks
 - Wait/notify
- Monitors

Reminder: Peterson's algorithm

```
// Process 1
init1;
while(true) {
    // entry protocol
    c1 = true;
    turn = 2;
    while (c2 && turn == 2);
    crit1;
    // exit protocol
    c1 = false;
    rem1;
}
```

```
// shared variables
bool c1 = c2 = false;
integer turn = 1;
```

```
// Process 2
init2;
while(true) {
    // entry protocol
    c2 = true;
    turn = 1;
    while (c1 && turn == 1);
    crit2;
    // exit protocol
    c2 = false;
    rem2;
}
```

To work the algorithm depends upon variables being set in this order and no reordering occurring

Fairness

- A *weakly fair* scheduling policy guarantees that if a process requests to enter its critical section (and does not withdraw the request), the process will *eventually* enter its critical section
 - i.e. wait long enough – don't stop waiting
- A *strongly fair* scheduling policy guarantees that if a process makes enough requests it will eventually succeed
 - Does not have to keep waiting
 - Can do something else between requests

Properties of the Java scheduler

- Java makes **no** promises about scheduling or fairness, and does not even strictly guarantee that threads make forward progress:
 - Most Java implementations display some sort of weak, restricted or probabilistic fairness properties with respect to executing runnable threads
 - **However you can't depend on this**
- There may be no guarantee that the any specific thread has a chance to run

Thread priorities

- **Threads** have *priorities* which **heuristically** influence schedulers:
 - Each thread has a priority in the range `Thread.MIN_PRIORITY` to `Thread.MAX_PRIORITY`
 - By default, each new thread has the same priority as the thread that created it
 - The initial thread associated with a main method by default has priority `Thread.NORM_PRIORITY`
 - The current priority of a thread can be accessed by the method `getPriority` and set via the method `setPriority`, e.g.:
`Thread.currentThread().setPriority(Thread.MAX_PRIORITY)`
- When there are more runnable threads than CPUs, a scheduler is generally biased in favour of threads with higher priorities
 - But it may not make much difference

Archetypical mutual exclusion

- We assumed that:
 - the initialisation, critical sections and remainder may be of any size and may take any length of time to execute—each may vary from one pass through the `while` loop to the next;
 - **the critical sections must execute in a finite time; i.e., each process must leave its critical section after a finite period of time**
 - the initialisation and remainder of each process may be infinite
- **If the critical sections don't execute in finite time, the scheduling policy cannot be weakly fair (others could wait for infinite time)**

Java **synchronized**

Mutual Exclusion in Java

- Every object in Java can be used as a Mutex
 - And as a condition variable – see later
- Use the keyword **synchronized** for this purpose
 - Notice the American spelling!!! z not s.
- You 'synchronize' on an object
 - Either the current object (this)
 - Or you give it an object explicitly
- Mutual exclusion is ensured between any two pieces of code which are **synchronized** on the same object (functions synchronize on **this**)

synchronized functions

```
synchronized void procedure1()
{
    for ( int i = 0 ; i < 1000000 ; i++ )
    {
        ++v;
    }
    System.out.println( "Procedure 1 Finished\n" );
}
```

```
synchronized void procedure2()
{
    for ( int i = 0 ; i < 1000000 ; i++ )
    {
        ++v;
    }
    System.out.println( "Procedure 2 Finished\n" );
}
```

synchronized blocks

```
void procedure1()  
{  
    for ( int i = 0 ; i < 1000000 ; i++ )  
    {  
        synchronized ( this )  
        {  
            ++v;  
        }  
    }  
    System.out.println( "Procedure 1 Finished\n" );  
}
```

... etc for the other procedure

- The code within the block is synchronized, rather than the entire function

Java **synchronized**

- You can consider a synchronized block/function to work as a mutex lock
 - Obtain lock in order to enter the block/function
 - Release the lock automatically when leaving the block/function
- Same problems as mutexes
 - E.g. deadlock issues (next slide)
- Avoids you having to remember to unlock the mutex
- You can use different objects if you want different mutexes to protect different critical sections

Deadlock example

```
class SynchronizedClass {  
    synchronized void f1( SynchronizedClass ob )  
    {    ob.f2(this);    }  
  
    synchronized void f2( SynchronizedClass ob )  
    {    i = 2;    }  
}
```

f1() on the first
object calls f2() on
the second object

Both methods are
synchronized

```
SynchronizedClass o1 = new SynchronizedClass();  
SynchronizedClass o2 = new SynchronizedClass();
```

```
public void go()  
{  
    new Thread() { public void run() {  
        for ( int i = 0 ; i < 100 ; i++ ) o1.f1(o2);  
    } }.start();  
    new Thread() { public void run() {  
        for ( int i = 0 ; i < 100 ; i++ ) o2.f1(o1);  
    } }.start();  
}
```

o1 is first
o2 is second

o2 is first
o1 is second

Reasons for deadlock

Thread 1

```
for ( int i = 0
      ; i < 100
      ; i++ )
    o1.f1(o2);
```

- Calls f1 on o1
 - Locks o1
- Calls f2() on o2
 - Locks o2

Thread 2

```
for ( int i = 0
      ; i < 100
      ; i++ )
    o2.f1(o1);
```

- Calls f1 on o2
 - Locks o2
- Calls f2() on o1
 - Locks o1

Reasons for deadlock

Thread 1

```
for ( int i = 0
      ; i < 100
      ; i++ )
  o1.f1(o2);
```

- Calls f1 on o1
 - Locks o1
- Calls f2() on o2
 - Locks o2

Thread 2

```
for ( int i = 0
      ; i < 100
      ; i++ )
  o2.f1(o1);
```

- Calls f1 on o2
 - Locks o2
- Calls f2() on o1
 - Locks o1




Data encapsulation

- Advantage of synchronized functions:
 - You can make your data private
 - Make all methods which access it synchronized
 - Then you will never have issues with interference between threads
- Java goes further than this, allowing simple wait and signal/notify facilities
 - And more complex facilities to implement more realistic monitors – later
- **Example: Producer-Consumer problem**

Rolling buffer

- Assume a limited queue

Items in queue: 0




- When queue is filled, it rolls around, to start again
- As long as there is enough room
 - i.e. consumer has consumed some by then

Rolling buffer

- Assume a limited queue


Items in queue: 1



Rolling buffer

- Assume a limited queue


Items in queue: 2



Rolling buffer

- Assume a limited queue


Items in queue: 3



Rolling buffer

- Assume a limited queue


Items in queue: 4



Rolling buffer

- Assume a limited queue


Items in queue: 5



Rolling buffer

- Assume a limited queue


Items in queue: 6



Rolling buffer

- Assume a limited queue

Items in queue: 6




- Assume that consumer finally got around to consuming one

Rolling buffer

- Assume a limited queue

Items in queue: 7




- Producer cannot now produce any more until the consumer has used some up
- There are no spare spaces to put the values

Rolling buffer

- Assume a limited queue

Items in queue: 6




- Consumer uses one, so there is one spare space
- Producer could continue now

Rolling buffer

- Assume a limited queue

Items in queue: 5




- Consumer uses another, so there are two spare spaces
- Producer could continue still

Rolling buffer

- Assume a limited queue

Items in queue: 4



- Consumer uses another, so there are three spare spaces now
- Producer could continue still

Example source code extract

```
final int BUFFER_SIZE = 8;
volatile int iCount = 0;
volatile int iInsertionPosition = 0;
volatile int iRemovalPosition = 0;
volatile int[] arrayItems = new int[BUFFER_SIZE];

synchronized boolean ProduceItem( int i )
{
    if ( iCount >= BUFFER_SIZE ) // Is there any space?
        return false;

    arrayItems[iInsertionPosition] = i; // Store item
    ++iInsertionPosition; // Update next insert posn
    if ( iInsertionPosition >= BUFFER_SIZE ) // Wrap around
        iInsertionPosition = 0; // Back to start

    ++iCount; // Increment count of items
    return true;
}
```

Returns failure if no space

Note: could loop, checking the variable rather than return but then we don't release the lock!

Similarly for the consumer...

```
synchronized int ConsumeItem()
{
    if ( iCount <= 0 )
        return -1; // No item to consume
    // Get item
    int iThisItem = arrayItems[iRemovalPosition];

    ++iRemovalPosition; // Increment removal position
    // Wrap around if we get to the end
    if ( iRemovalPosition >= BUFFER_SIZE )
        iRemovalPosition = 0;

    --iCount; // Reduce count of items available

    return iThisItem;
}
```

Returns failure if no item

Overview of usage

Producer

- **While there is no space**
 - Wait and then try again
- Produce item
- Increment count of items

Consumer

- **While there are no items**
 - Wait and then try again
- Retrieve/consume item
- Decrement count of items

- Producer can only produce items when there is space to do so
- Consumer can only consume items when there are items on the queue
- Problem is the spinning at the beginning
- **It would be useful to be able to wait for an item to be produced (or space to be available) rather than to keep trying**

Wait and notify/signal let us do that

- Enter the synchronized block
- Check your condition to continue
- If you have to wait, call wait()
 - wait() releases the lock that you have until it is awoken
 - When it 'wakes up' again it will have to wait to get that lock before continuing
- Anything which changes the condition it is waiting on must call notify, to wake it up again
 - When it wakes up, it must re-check the condition
 - If it is still not true then go back to sleep (wait())
- You can have multiple condition variables (queues)
 - Basic Java implementation assumes just one condition variable
 - Condition variable is associated with an object – wait on object
 - You have to wait() on the same object you synchronized on

Consumer

```
synchronized int ConsumeItem()
{
    // Do the wait-notify block first
    while ( iCount <= 0 )
        try { wait(); } catch (InterruptedException e) {}
    // Note: doing this.wait() to wait on current object

    int iThisItem = arrayItems[iRemovalPosition]; // Get item

    ++iRemovalPosition; // Increment removal position
    if ( iRemovalPosition >= BUFFER_SIZE ) // Wrap around?
        iRemovalPosition = 0;

    --iCount; // Decrement count of items stored

    // Tell any waiting producers that it is worth carrying on now
    notifyAll();

    return iThisItem;
}
```

Producer

```
synchronized boolean ProduceItem( int i )
{
    // Do the wait-notify block first
    while ( iCount >= BUFFER_SIZE )
        try { wait(); } catch (InterruptedException e) { }

    arrayItems[iInsertionPosition] = i; // Insert item

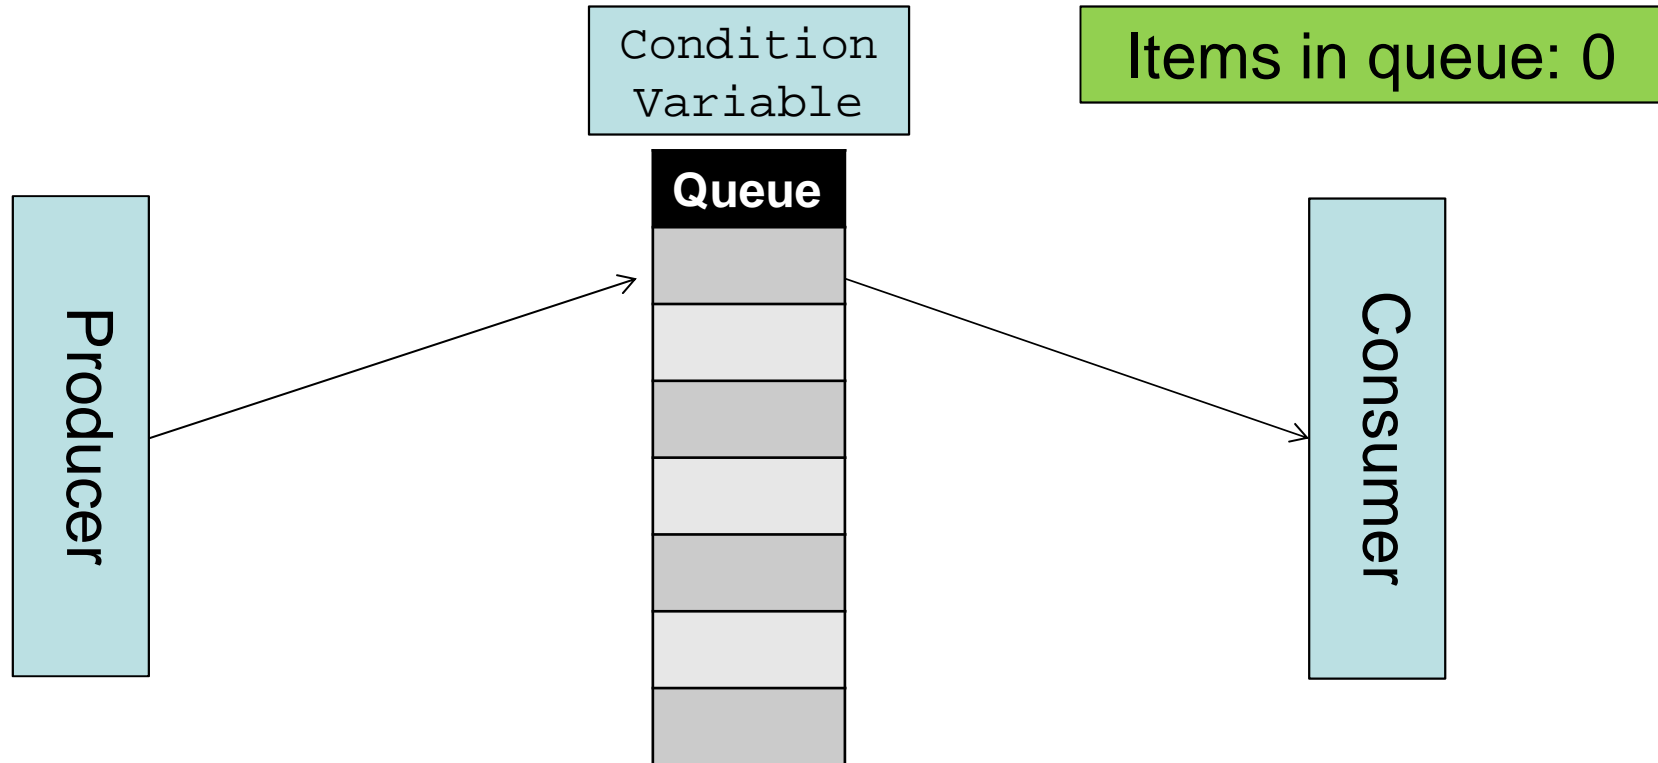
    ++iInsertionPosition; // Increment insertion position
    if ( iInsertionPosition >= BUFFER_SIZE ) // Wrap around?
        iInsertionPosition = 0;

    ++iCount; // Increment count of items

    // Tell any waiting consumers that it is worth carrying on now
    notifyAll();

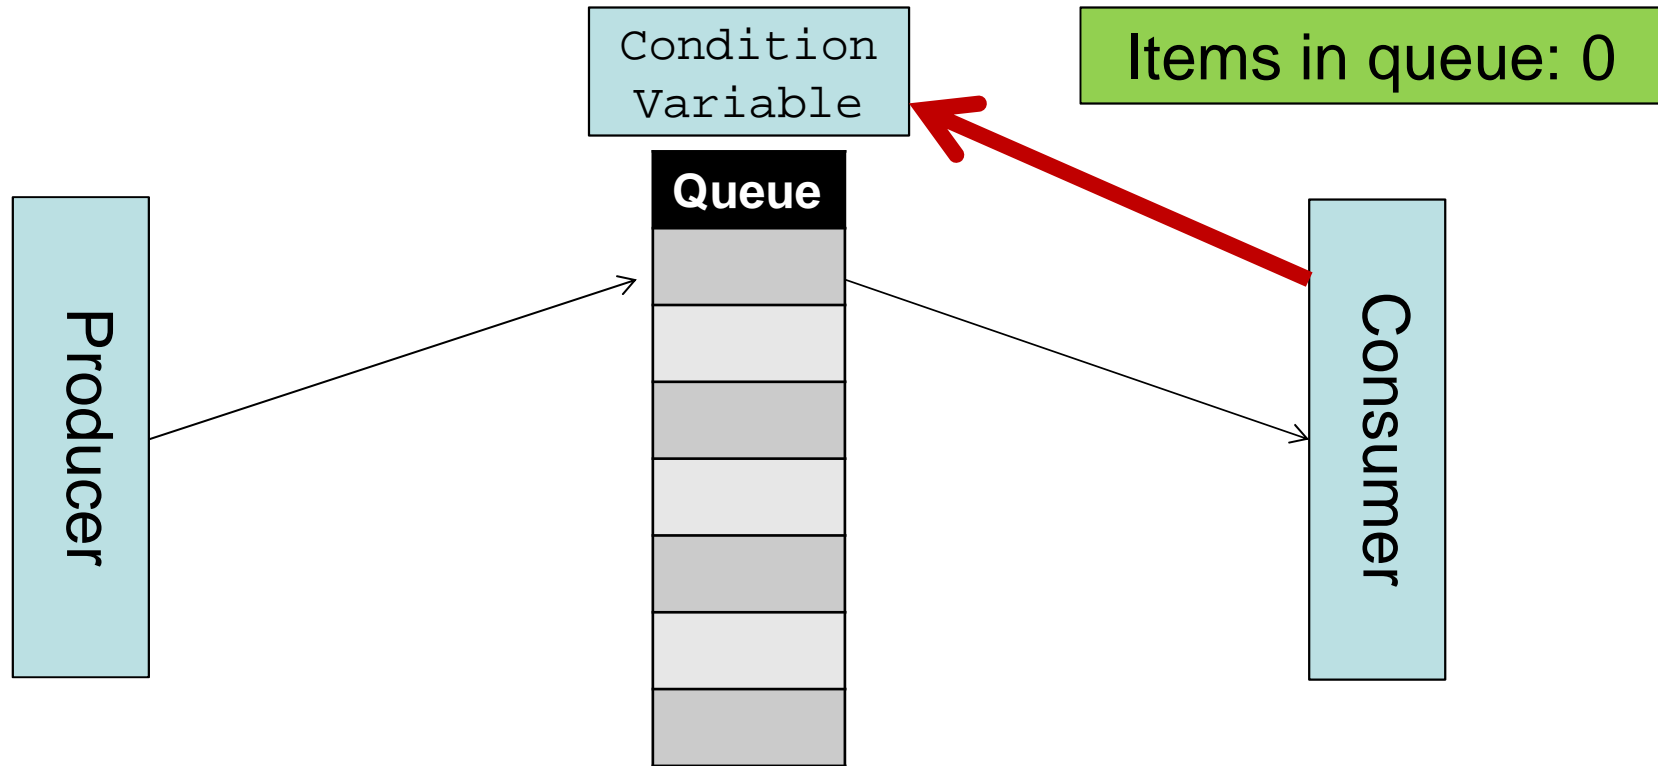
    return true;
}
```

Producer-Consumer



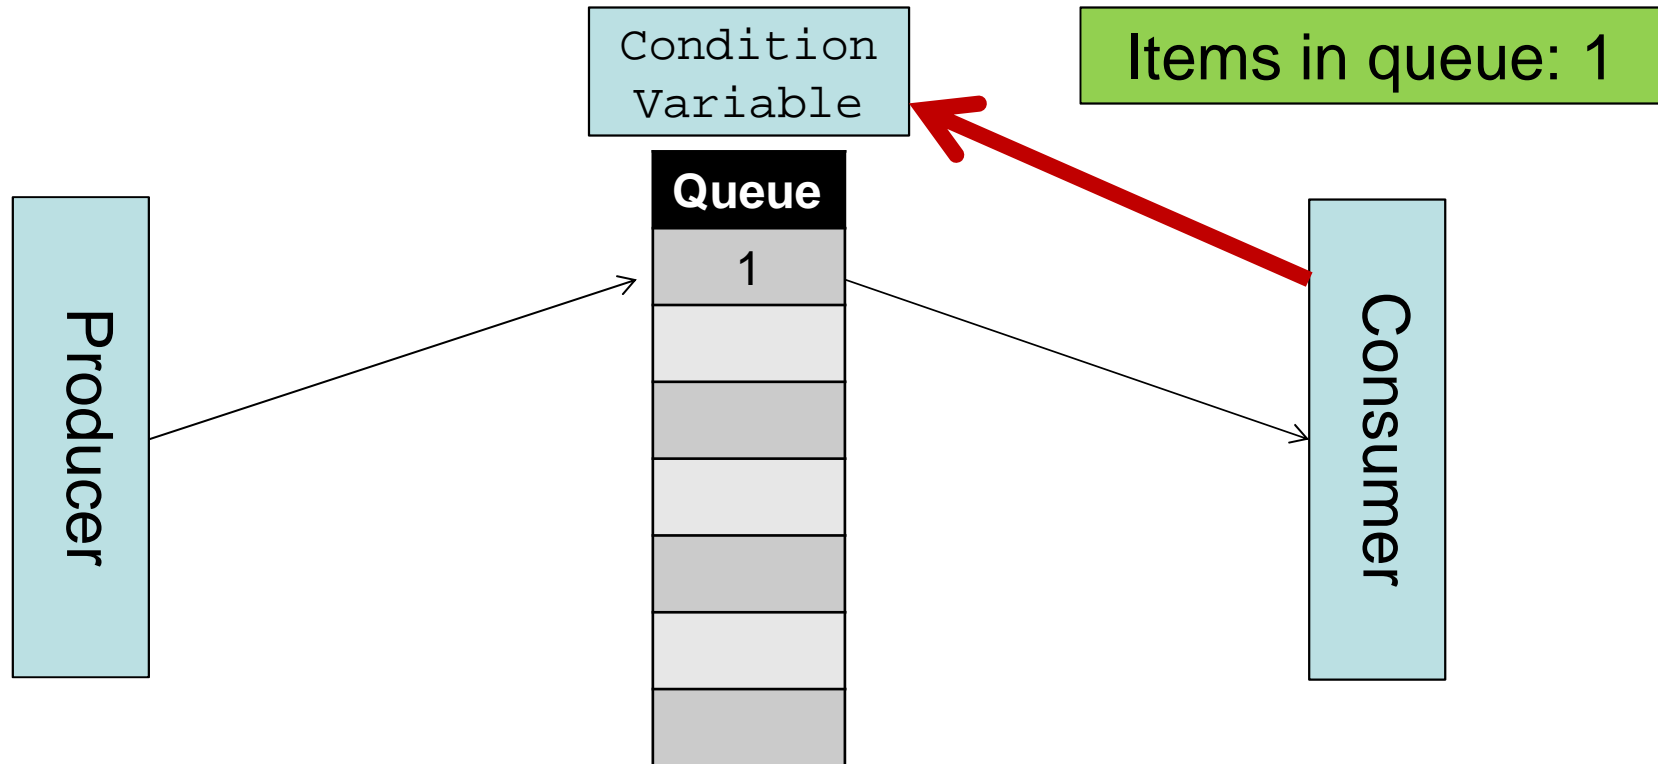
- Consumer tries to consume
- Nothing in the queue ...

Producer-Consumer



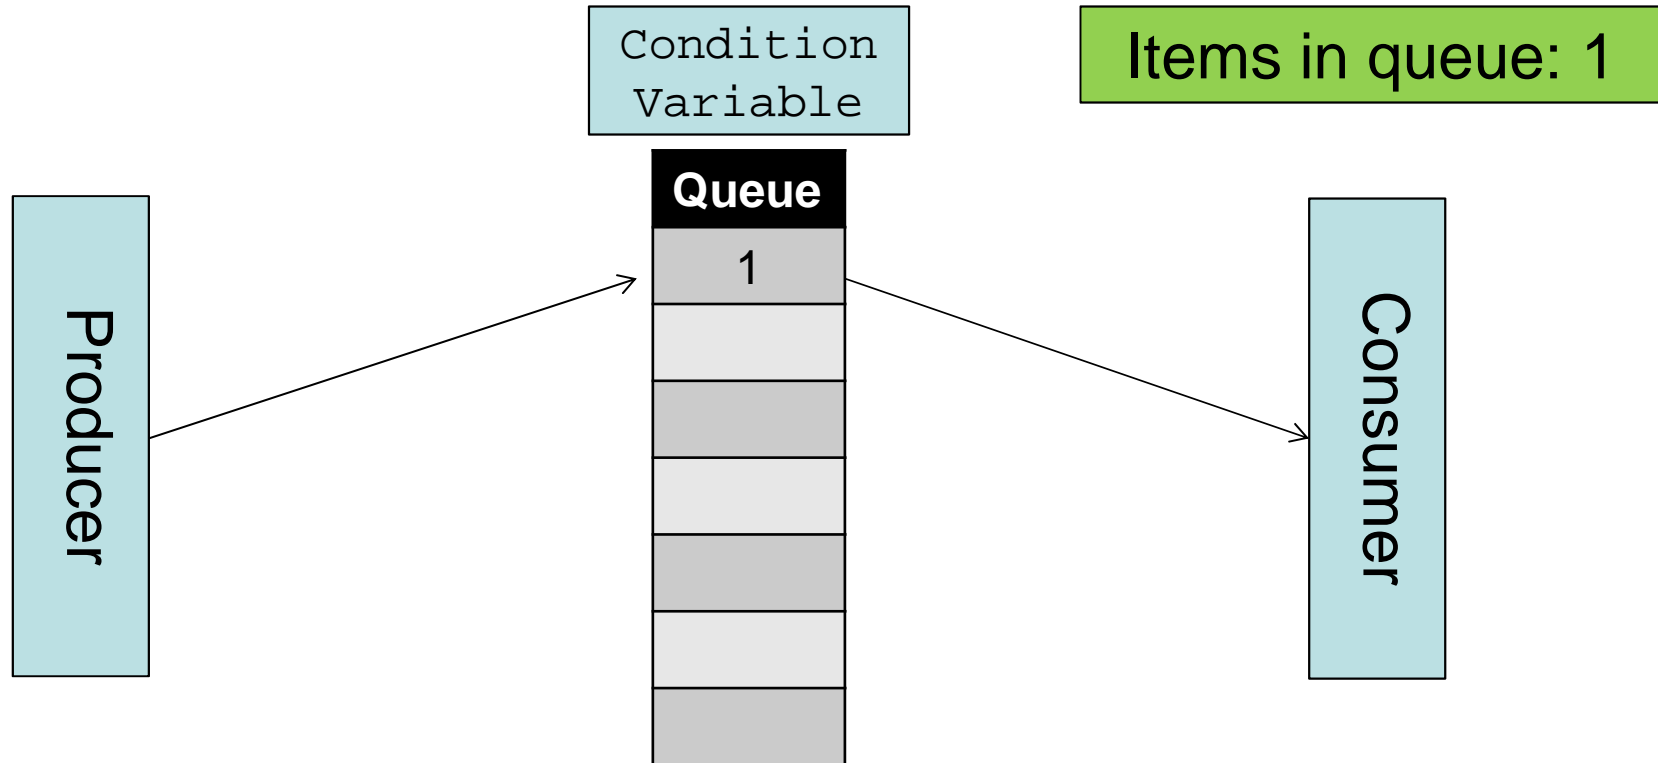
- Consumer tries to consume
- Nothing in the queue
- Consumer issues `wait()` on the condition variable

Producer-Consumer



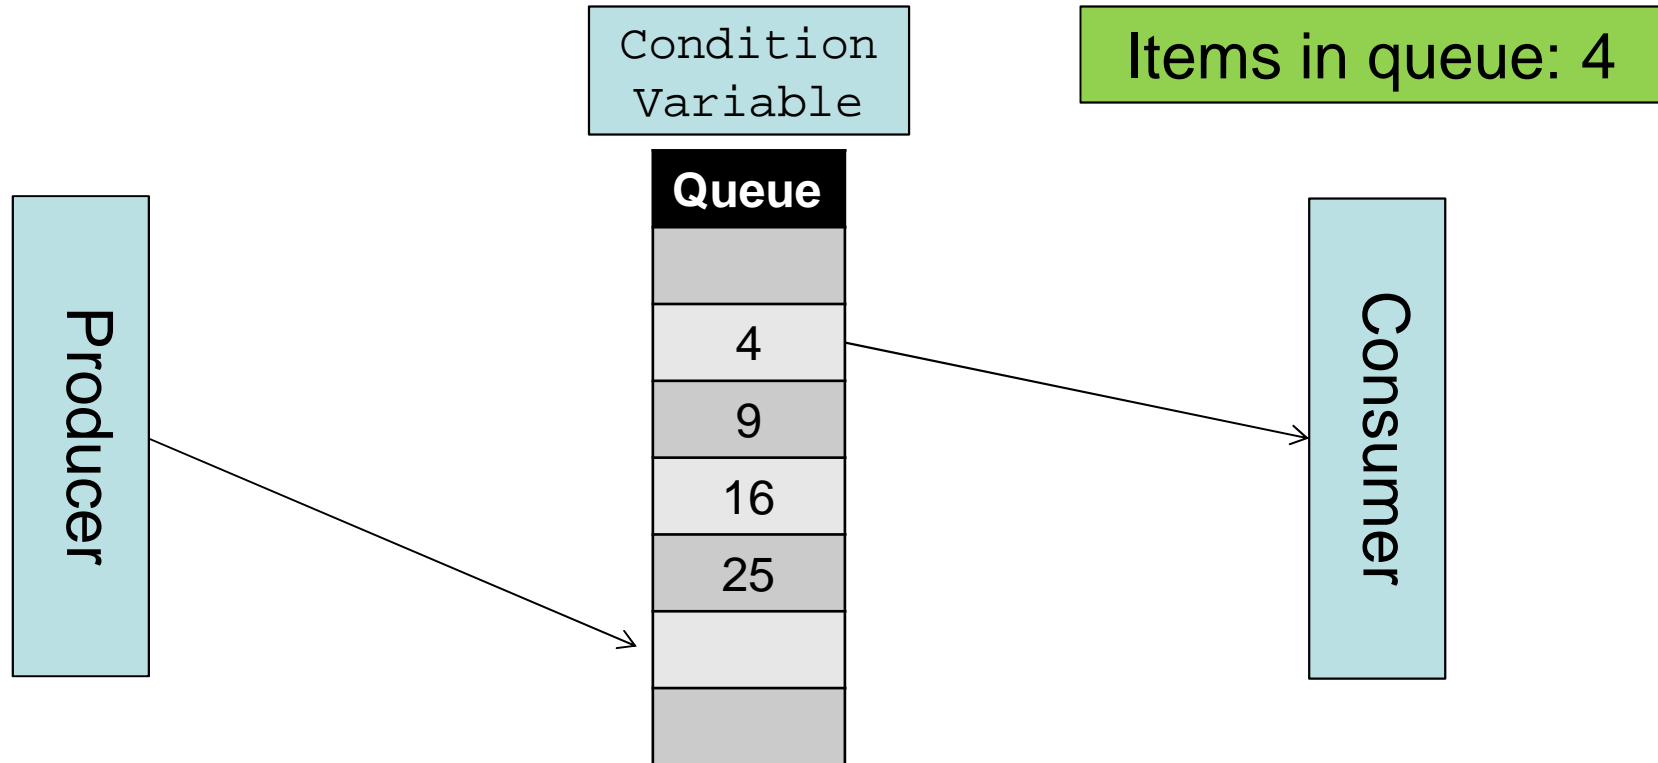
- Producer produces an item
- Producer calls `notify()` on the condition variable
 - Often called 'signal' rather than notify

Producer-Consumer



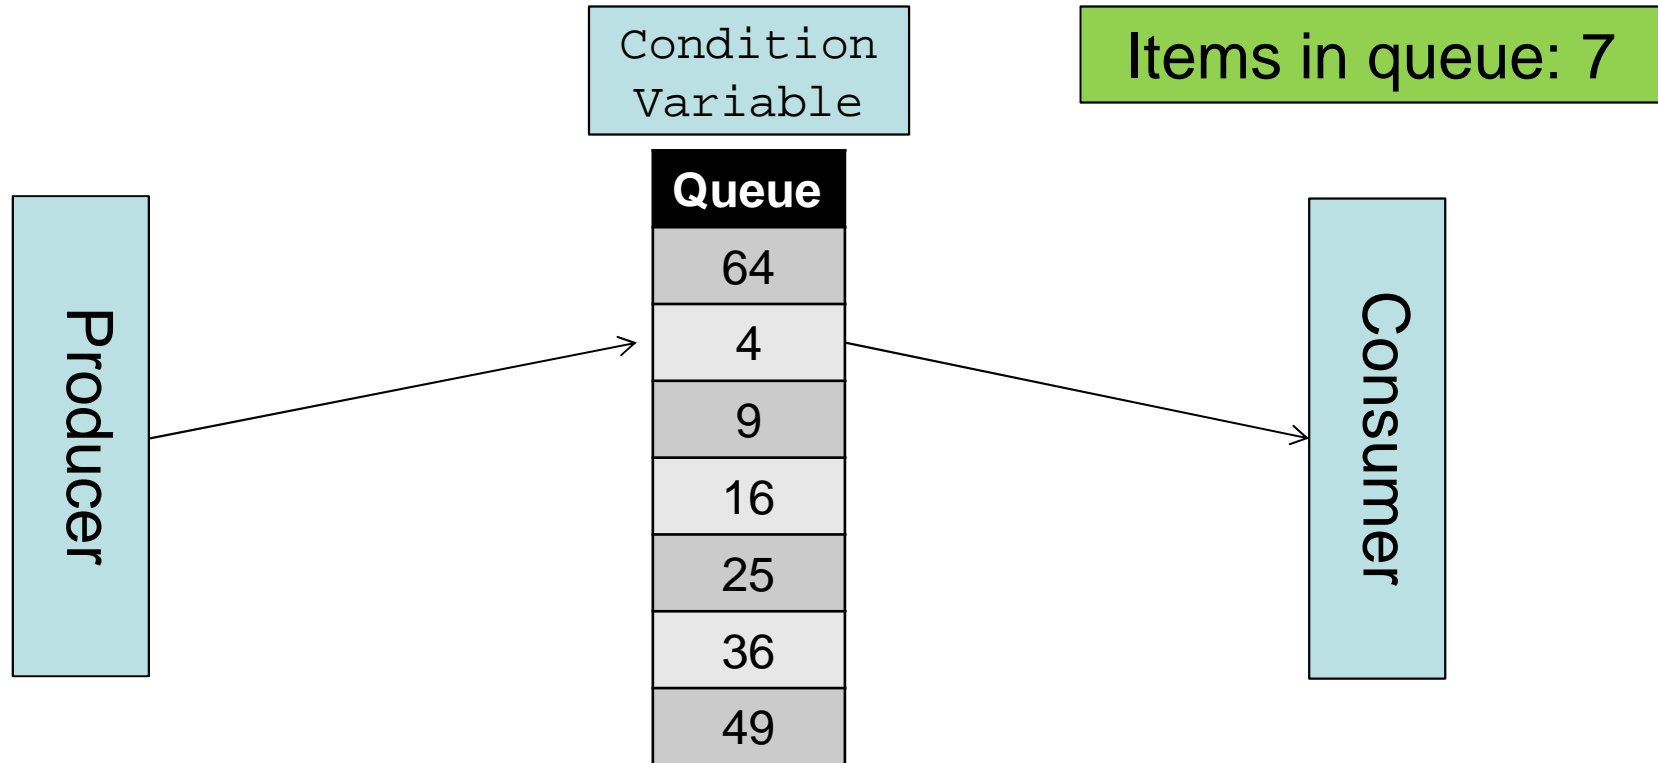
- Consumer wakes up and wants the lock so that it can continue
- When Producer leaves the synchronized function, consumer can enter its own, and consume the item

Producer-Consumer




- Producer keeps producing...

Producer-Consumer




- Producer keeps producing...
- At a later point, the Producer finds that the queue is full...

Producer-Consumer




- Producer keeps producing
- At a later point, the Producer finds that the queue is full
- Producer calls Wait() on the condition variable

Producer-Consumer



- Consumer will eventually consume a product, making room for a new one
- **Every** time consumer has consumed a product it calls `notify()` on the Condition Variable

Producer-Consumer



- Consumer calls notify
- Producer is awakened and will check for space
 - When it can get the lock / enter the synchronized section again
 - If still no space it will wait() again

wait() and notify()

- The thing you wait on is called a **condition variable**
 - In Java any object can be used for this
- In the basic Java implementation you need to have locked the object that you are using as a condition variable (i.e. synchronized on it)
 - i.e. you cannot have more than one condition variable associated with the object
- The wait will unlock it, awakening will re-lock it
- More Java complex concurrency classes allow this though (e.g. ReentrantLock)
- Now we are ready to understand monitors...

Monitors

Monitors as abstract data types

- A *monitor* is an abstract data type representing a shared resource and operations to protect and manipulate it
- Monitors (conceptually) *encapsulate* the shared resource
- A monitor implements a shared data structure together with the operations which manipulate the data structure
- Think “private data and public access methods”
- Monitors have four components:
 - A set of *private variables* which represent the state of the resource (the data to protect)
 - A set of *monitor procedures* which provide the public interface to the resource (the functions/methods you can call)
 - A set of *condition variables* used to implement condition synchronisation (e.g. a queue of waiting threads)
 - *Initialisation code* which initialises the private variables

Next Lecture

- More Monitors
 - The theory
 - Implementing full monitors
- More concurrency in Java