

G520SC

OPERATING SYSTEMS AND


CONCURRENCY

ThreadPools and Problems

Dr Jason Atkin

Last lecture: Monitors

- Private data
- Public methods
- Locks
- Condition variables




This Lecture

- ThreadPools
 - And Futures
- Dining Philosophers Problem


Threadpools

- It is common for us to need to do things in parallel for a short time
- Then to do so again
- And again
- And again...
- But creating threads has an overhead



Threadpools

- Threadpools are pools of threads which are available for you to use
- They will run the functions for you
 - At some point, when they get around to it
- And block when not running something
 - No CPU usage
- Decide how many you need in the pool



Java threadpool example

- Create the thread pool and specify number of threads

```
ExecutorService executorConsumers =  
    Executors.newFixedThreadPool(3);
```

- Or:

```
= newCachedThreadPool()  
= newSingleThreadExecutor()
```

etc (there are others, see **Executors** class)

- **Then use the executor service to execute a runnable:**

```
executorConsumers.execute( myRunnable );  
executorConsumers.execute( new Runnable()  
    { public void run() { doConsumer(); } } );
```

C/Windows API threadpool example

As usual, using C and the windows API is a bit more tricky, but the same principles apply and I have created a sample for you to illustrate this.

- Create the threadpool:

```
PTP_POOL pool = CreateThreadpool( NULL );
```

- Specify number of threads to run:

```
SetThreadpoolThreadMinimum( pool, 1 );
```

```
SetThreadpoolThreadMaximum( pool, 4 );
```

- Tell it to run the function:

```
SubmitThreadpoolWork( work[i] );
```

- If you want to know when some work finishes:

```
WaitForThreadpoolWorkCallbacks( work[i], FALSE );
```

- Close it the work when you have finished with it (before or after it runs):

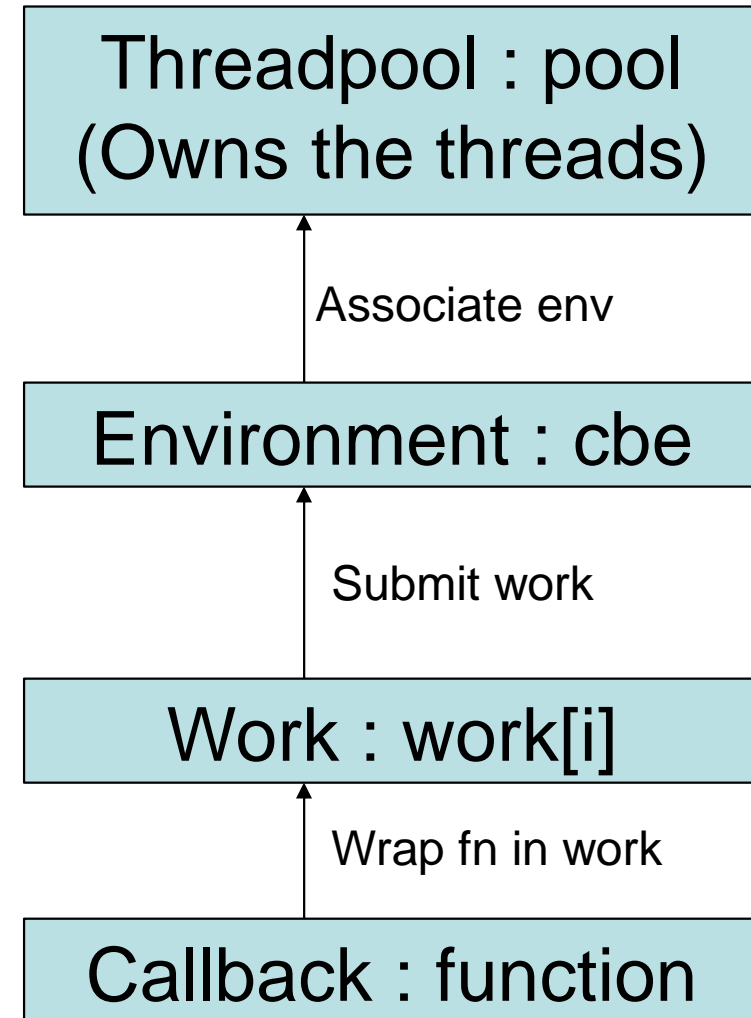
```
CloseThreadpoolWork( work[i] );
```

- And close the threadpool when you will not use it any more:

```
CloseThreadpool( pool );
```

ThreadPool Environment and Work

- Create the pool
- Associate environment with it
`InitializeThreadPoolEnvironment()`
`SetThreadPoolCallbackPool()`
- For every function you wish to execute, create a work object to wrap it up and associate it with the environment
`CreateThreadPoolWork(...)`
- Then you submit the work object, not the function:
`SubmitThreadPoolWork(work[i]);`



Back to Java: futures

Java futures example

- Posting to a threadpool, you may want to know:
 - When has the work finished?
 - What was the result?
- In C you would probably give it a pointer to some result structure to fill in (using the parameter that you can pass to a thread function/work)
- But there are nicer/more elegant ways to do this via futures
 - Java and C++ (since C++11) both support futures
 - So do other languages, or at least the concepts
 - We are ignoring C++ so will look at the Java version

Submitting the callback

- Create the ThreadPool as usual

```
ExecutorService executor =  
    Executors.newSingleThreadExecutor();
```

- Use submit() rather than execute()
 - Give it an object which implements Callable<> (call() function)

```
Future<String> myFutureString =  
    executor.submit( MyCallableObject );
```

```
Future<String> myFutureString = executor.submit(  
    new Callable<String>() {  
        public String call() {  
            try { Thread.sleep(1000); }  
            catch (InterruptedException e) { }  
            return "Finished...";  
        }  
    } );
```

Submitting the callback

- Create the ThreadPool as usual

```
ExecutorService executor =  
    Executors.newSingleThreadExecutor();
```

- Use submit() rather than execute()
 - Give it an object which implements Callable<> (call() function)

```
Future<String> myFutureString =  
    executor.submit( MyCallableObject );
```

```
Future<String> myFutureString = executor.submit(  
    new Callable<String>() {  
        public String call() {  
            try { Thread.sleep(1000); }  
            catch (InterruptedException e) { }  
            return "Finished...";  
        }  
    } );
```

Using the Future


- This will create a future and capture various information in the future (for 'future' use by you)
- The future is returned from the submit() function:
`Future<String> myFutureString = executor.submit(`
- You can use the future to see what is happening
 - Check whether the function call has completed:
`myFutureString.isDone()`
 - Retrieve the returned value – type depends on the future type
`myFutureString.get()`
 - The value from get() will be whatever you **returned** from the function which was **submitted**

Dining Philosophers Problem

Dining Philosophers Problem

- The *Dining Philosophers* problem illustrates mutual exclusion between processes which compete for overlapping sets of shared variables
 - five philosophers sit around a circular table
 - each philosopher alternately thinks and eats spaghetti from a dish in the middle of the table (think...eat...think...eat)
 - the philosophers can only afford five forks—one fork is placed between each pair of philosophers
 - to eat, a philosopher needs to obtain mutually exclusive access to both the fork on their left and right
- The problem is to avoid *starvation* / deadlock
 - e.g., each philosopher acquires one fork **and refuses to give it up until they have eaten something with it**, can you define a protocol such that all of them will get to each

Dining Philosophers Problem



Examples: Each tries to take the one on their right then the one on their left, or each tries a random one each time then the other one.
Which policies will work to solve the problem (to the end of the meal)?

Deadlock in the Dining Philosophers

- The key to the solution is to avoid *deadlock* caused by circular waiting
- E.g. everyone grabs the fork to their right then tries to grab the one on their left:
 - process 1 is waiting for a resource (i.e. a fork) held by process 2
 - process 2 is waiting for a resource held by process 3
 - process 3 is waiting for a resource held by process 4
 - process 4 is waiting for a resource held by process 5
 - process 5 is waiting for a resource held by process 1
- No process can make progress and all processes remain deadlocked

Some slides from 1st year
Database Systems module
on Concurrency/Deadlocks

Deadlocks

- A deadlock is an impasse that may result when two or more ~~transactions~~ **threads** are waiting for locks to be released which are held by each other. E.g.:
 - T1 has a lock on X and is waiting for a lock on Y
 - T2 has a lock on Y and is waiting for a lock on X
 - **Both wait for each other**
- We can detect deadlocks that will happen in a schedule using a *wait-for graph* (WFG)
- We are tracking what is 'waiting for' something else
 - If there is any **cycle** of things waiting for other things, then they will **deadlock**

Using Wait-For Graphs

If you have a specific trace then you can use a wait-for graph to look for deadlock (already occurred)

Operating systems and databases do this a lot


What you would actually like to know though is whether **any** trace can result in deadlock later

Wait-for Graph

- Aim to find out if we (will) have a deadlock
- Each ~~transaction~~ thread is a vertex
- Looks at locks
- Edge from T2 to T1 if
 - T1 **read-locks** X then T2 tries to **write-lock** it
 - T1 **write-locks** X then T2 tries to **read-lock** it
 - T1 **write-locks** X then T2 tries to **write-lock** it

Databases: Example


Schedule	Locks
T1 Read(X)	write-lock(X)
T2 Read(Y)	read-lock(Y)
T1 Write(X)	
T2 Read(X)	tries read-lock(X)
T3 Read(Z)	write-lock(Z)
T3 Write(Z)	
T1 Read(Y)	tries write-lock(Y)
T3 Read(X)	tries read-lock(X)
T1 Write(Y)	



Wait for
graph

Observations

- To get deadlock (i.e. to have a problem) you need each thread which is involved to:
 - Have a lock that another thread wants
 - Attempt to gain a lock that another thread has
- If you only ever want one lock at a time, this cannot happen




- If threads use different locks/resources you cannot have a problem
- If you have only one thread it is OK as long as locks have owners

Informal: can we use a graph
to help us to spot potential
problems in advance?

INFORMAL similar graph for locks

- Build a graph showing which threads may lock which resources, in which order
- Thread 1:
 - Lock A, Lock B, Unlock B, Unlock A
- Thread 2:
 - Lock B, Lock C, Unlock C, Unlock B
- More complicated rules, so only an informal analysis
 - This is an **informal** way to visualise what you would do with traces


- Switch arcs and vertices
 - Vertices of the graph are the locks
 - Arcs of the graph are the order in which a thread requests the locks




- Note that this transformation makes no difference topologically

INFORMAL similar graph for locks

- Add the **paths** for the order in which locks can occur for each thread
 - If a process 1 requests lock A then lock B, add an arc from A to B (label it 1)
 - If a process 2 requests lock B, then lock A then add an arc from B to A
- Cycles represent the possibility that deadlock **could** occur if **all of the locks and requests can occur at the same time**





- Classic problem:
 - T1: A then B
 - T2: B then A



INFORMAL similar graph for locks



- If there is no cycle there can be no deadlock
- If every thread locks everything in the same order, you cannot possibly get a cycle
- If there is a cycle then you **may** be able to get deadlock (if all of the locks **can occur at once**)
 - Really you need to check **all of the traces** but this may be easier visually



Applying this to the Dining Philosophers Problem



Dining Philosophers Problem

- When we trace the potential paths, what does it look like?
- Philosophers all take the fork on their right then the fork on their left



Dining Philosophers Problem



- Philosophers all take the fork on their right then the fork on their left
- Each gets one fork and has to wait for the other ... forever



Dining Philosophers Problem

- Philosophers all take the fork on their right then the fork on their left
- Each gets one fork and has to wait for the other ... forever



How could we fix this?



Dining Philosophers Problem

Switching the order for one philosopher fixes the problem

Order: 1 2 3 4 5



Semaphore Solution

```
// Philosopher i, i == 1-4

while(true) {
    //get right fork then left
    P(fork[i]);
    P(fork[i+1]);
    // eat ...
    V(fork[i]);
    V(fork[i+1]);
    // think ...
}
```

```
// Philosopher 5


while(true) {
    //get left fork then right
    P(fork[1]);
    P(fork[5]);
    // eat ...
    V(fork[1]);
    V(fork[5]);
    // think ...
}
```

```
// Shared variables
binary semaphores fork[5] = {1, 1, 1, 1, 1};
```


Further examples


Example Locks

- T1:
 - Lock A
 - Lock D
 - Lock E
- T2:
 - Lock E
 - Lock C
 - Lock B
- T3:
 - Lock C
 - Lock D
 - Lock B




Can this deadlock?

- T1:
 - Lock A
 - Lock D
 - Lock E
- T2:
 - Lock E
 - Lock C
 - Lock B
- T3:
 - Lock C
 - Lock D
 - Lock B




Can this deadlock?

- T1:
 - Lock A
 - Lock D
 - Lock E
- T2:
 - Lock E
 - Lock C
 - Lock B
- T3:
 - Lock C
 - Lock D
 - Lock B



Re-entrant locks


- Re-entrant lock: if you lock it again you still have it
 - E.g. Mutex
- Non-reentrant: If you lock it again you count as a new locker
 - Deadlock if it is a binary lock
 - E.g. binary semaphore



Returning to example locks


- T1:
 - Lock A
 - Lock D
 - Lock E
- T2:
 - Lock E
 - Lock C
 - Lock B
- T3:
 - Lock C
 - Lock D
 - Lock B

Can we fix this problem?
While still keeping the locks
when necessary?



Assume reentrant locks

- **Add T3 lock D first**
- T1:
 - Lock A
 - Lock D
 - Lock E
- T2:
 - Lock E
 - Lock C
 - Lock B
- T3:
 - **Lock D**
 - Lock C
 - ~~Lock D~~
 - Lock B



Order: ?

Assume reentrant locks

- **Add T3 lock D first**

- T1:


- Lock A
- Lock D
- Lock E

- T2:

- Lock E
- Lock C
- Lock B

- T3:

- **Lock D**
- Lock C
- ~~Lock D~~
- Lock B




Order: A D E C B

More examples (1)

- T1:
 - Lock A
 - Lock B
 - Lock C
- T2:
 - Lock A
 - Lock C
 - Lock B


Order: ?



- Remember that the question is: could any trace end in deadlock?

More examples (2)

- T1:
 - Lock A
 - Lock B
 - Lock C
- T2:
 - Lock A
 - Lock C
 - Lock B
- T3:
 - Lock C
 - Lock B



- Remember that the question is: could any trace end in deadlock?

More examples (3)

- Remove lock A from T2:


- T1:

- Lock A
- Lock B
- Lock C

- T2:

- ~~Lock A~~
- Lock C
- Lock B

Order: ?



- Remember that the question is: could any trace end in deadlock?

Using graphs

- Wait-for-graphs are a common means for detecting deadlock
- We can use an informal variant to try to find whether deadlock is possible
- Eliminating cycles is useful
 - In which case you have an order for your locks
- You can use whatever method you wish to use in the coursework or exam
 - I only showed you the informal method I use

Summary

- In general, putting an order on your locks is very useful
- You can **guarantee** that you do not get deadlock
- Only having one lock at a time is a very simple example of this, but often not practical
- Sometimes there are advantages to being able to lock out-of-order (for efficiency)
 - It is risky – a small change could break it

Next week

- Another classic pattern :
Readers and Writers
- Summary/review
 - The changing face of concurrency
 - Overview: communication and data protection
 - Why was this important?
 - How do we apply what we have seen?
 - Hardware support, now and the future?