

G520SC

OPERATING SYSTEMS AND

CONCURRENCY

Readers and Writers

Dr Jason Atkin

This lecture

- Readers-Writers problems
- Fairness vs correctness

An example

- A simple library catalogue has two kinds of users:
 - *Borrowers* who use the catalogue to search for books or to find out whether a particular book is on loan (readers)
 - *Library staff* who update the catalogue when new books are added to the library stock or to record which books are on loan or returned (writers)
- If we allow all users unrestricted access to the catalogue, a borrower may see the catalogue when it is an inconsistent state, e.g.:
 - A borrower is looking for a book which has just arrived at the library and not all details have been added
 - A librarian is updating the catalogue to include the book, e.g., by copying and editing an existing entry
 - The borrower sees an inconsistent state of the catalogue, e.g., the entry contains an incorrect shelf number


Correctness vs efficiency

- **Readers and Writers problem: special case** of general problem where processes may all **both** read **and** write:
 - Any solution to the general problem will also be a **correct** solution to the Readers and Writers problem
 - However (much) more **efficient** solutions are possible for the Readers and Writers problem
- In general, an efficient solution depends on the specifics of the problem – can we take advantage of structure?
- One solution to the readers and writers problem would be to make **all** accesses to the file **mutually exclusive**:
 - At most one process/thread (reader or writer) would be able to access the file at a time (e.g. binary semaphore, monitor, mutex)
 - **This is simple and correct** but, in general, the performance of such an approach is unacceptable (only one borrower/reader could search the catalogue at a time)

Synchronisation requirements

- To ensure correctness **and for maximum efficiency** (concurrency) and we require that:
 - If a writer is writing to the file:
 - No other writer may write to the file
 - No reader may read it
 - Any number of readers may simultaneously read the file (if no writers)

How could we do this?



Tools:

Monitor

Semaphore


Mutex /
Critical
Section

Condition
Variable /
Event

- How could we build a protocol to enforce mutual exclusion for a single writer, without being inefficient?


Semaphore

- A counter with a waiting mechanism
- $p()$ wait() : 'lock' it
 - If count is zero then wait on the semaphore
 - Otherwise decrement count and continue
- $v()$ notify() signal() :
 - If there are waiting processes, wake one up
 - Otherwise increment the count

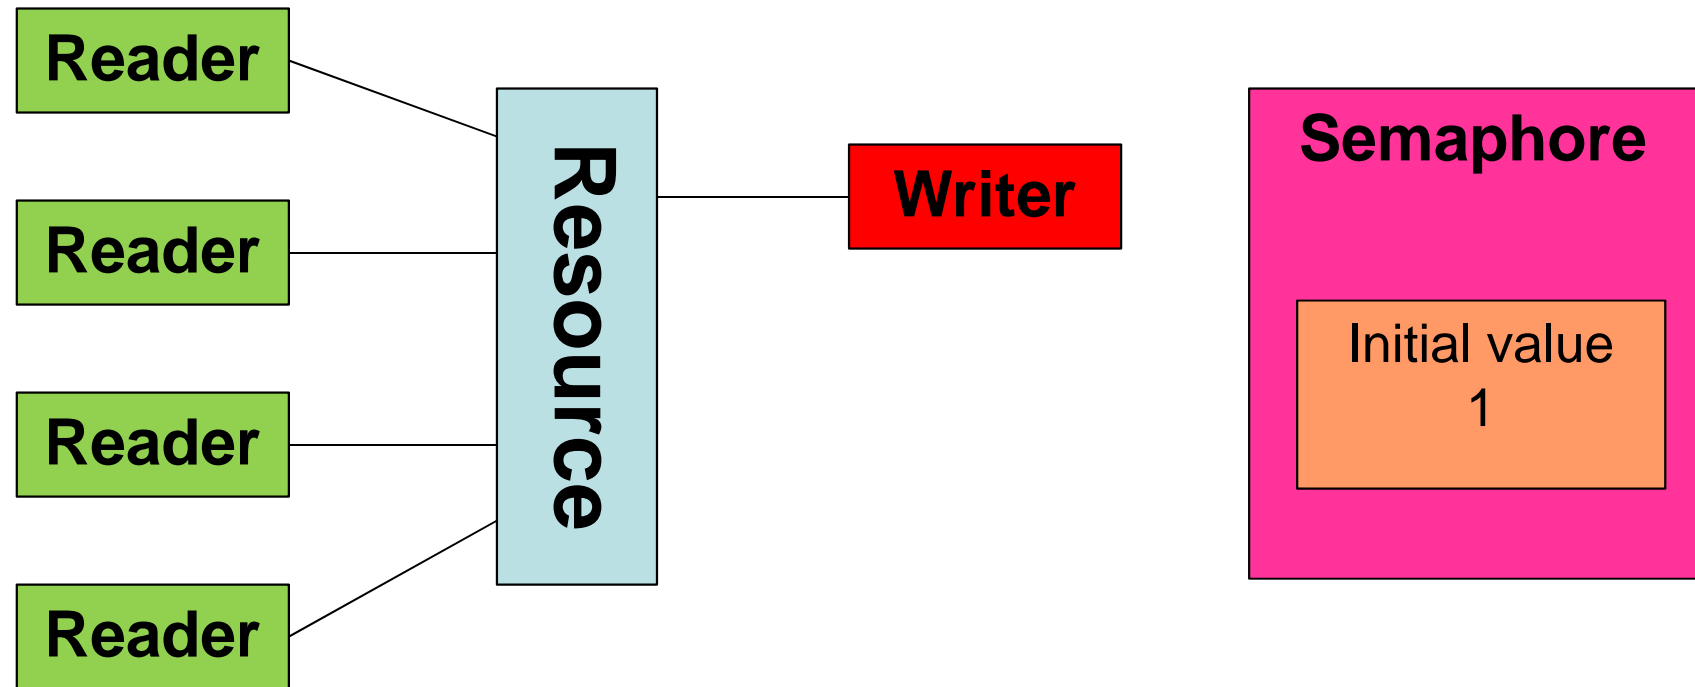


Equivalent Semaphore Implementation

- A counter with a waiting mechanism
- $p()$ wait() : 'lock' it
 - Decrement count
 - If count is now negative then wait on the semaphore
- $v()$ notify() signal() :
 - Increment semaphore value
 - If count was negative **before** the increment (i.e. a process was waiting) then wake up a waiting process




Single writer binary semaphore?



- Assume four readers and one writer
- What can we do with a BINARY semaphore?

Single writer binary semaphore

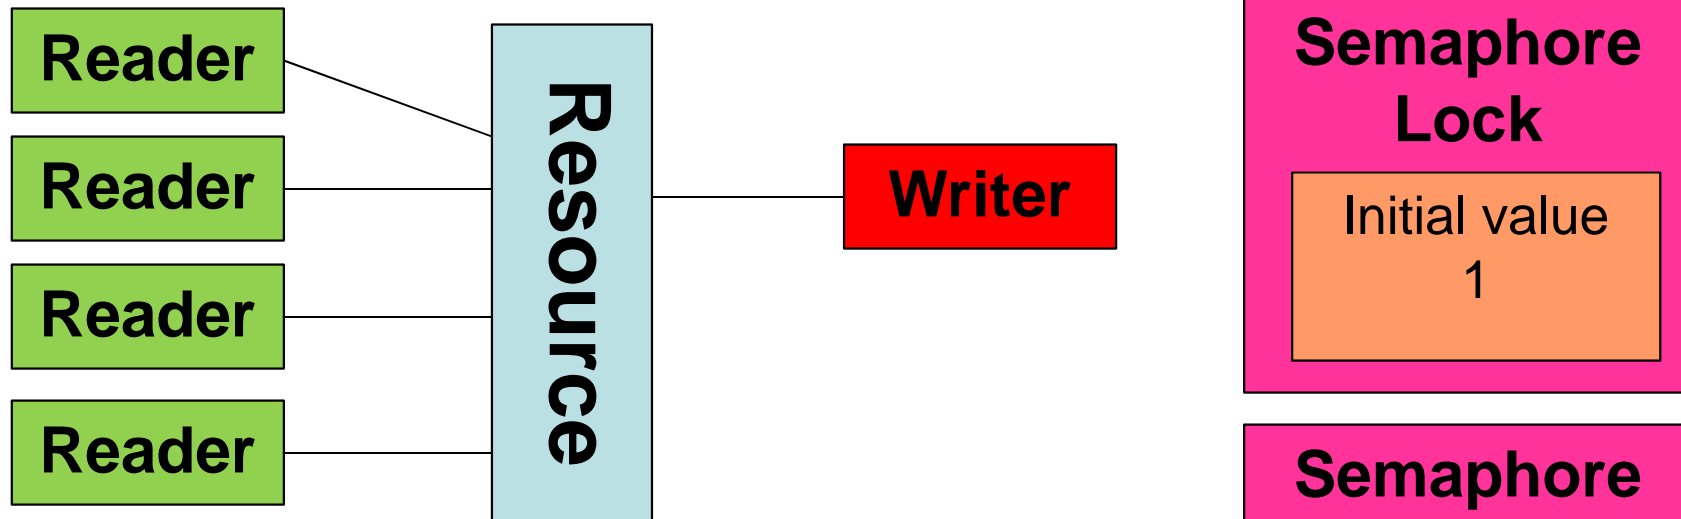


- Each thread does $p()$, uses it, then $v()$
- Mutual exclusion between everything - **inefficient**

Could we allow multiple
readers?

Hint: using a second binary
semaphore

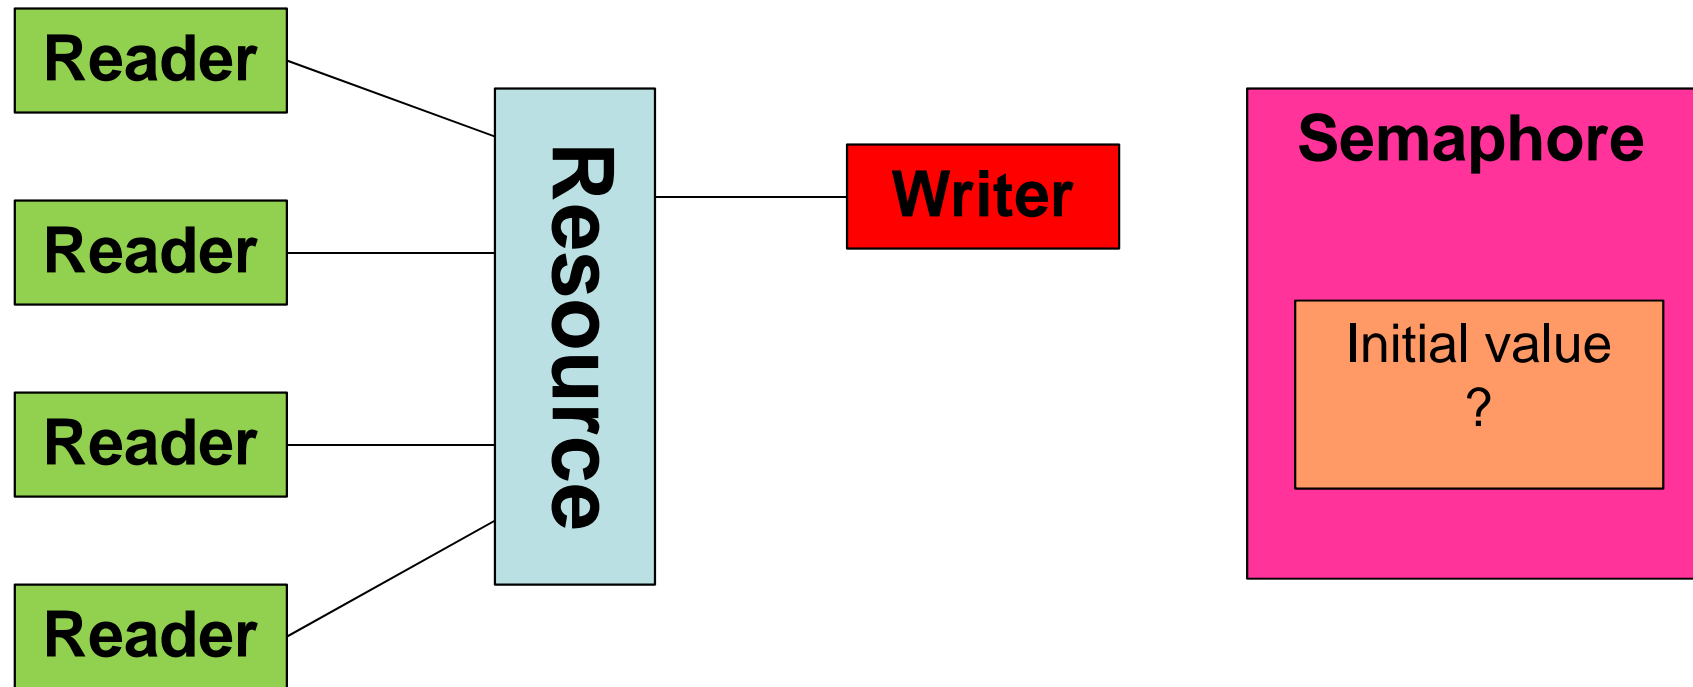
Allowing multiple readers



- Writer: `lock.p() lock.v()`
- Reader: `reader.p(), readcount++,`
`if (readcount==1) lock.p(),`
`reader.v()`
`...`
`reader.p(), readcount--,`
`if (readcount==0) lock.v(),`
`reader.v()`

Using a generic semaphore

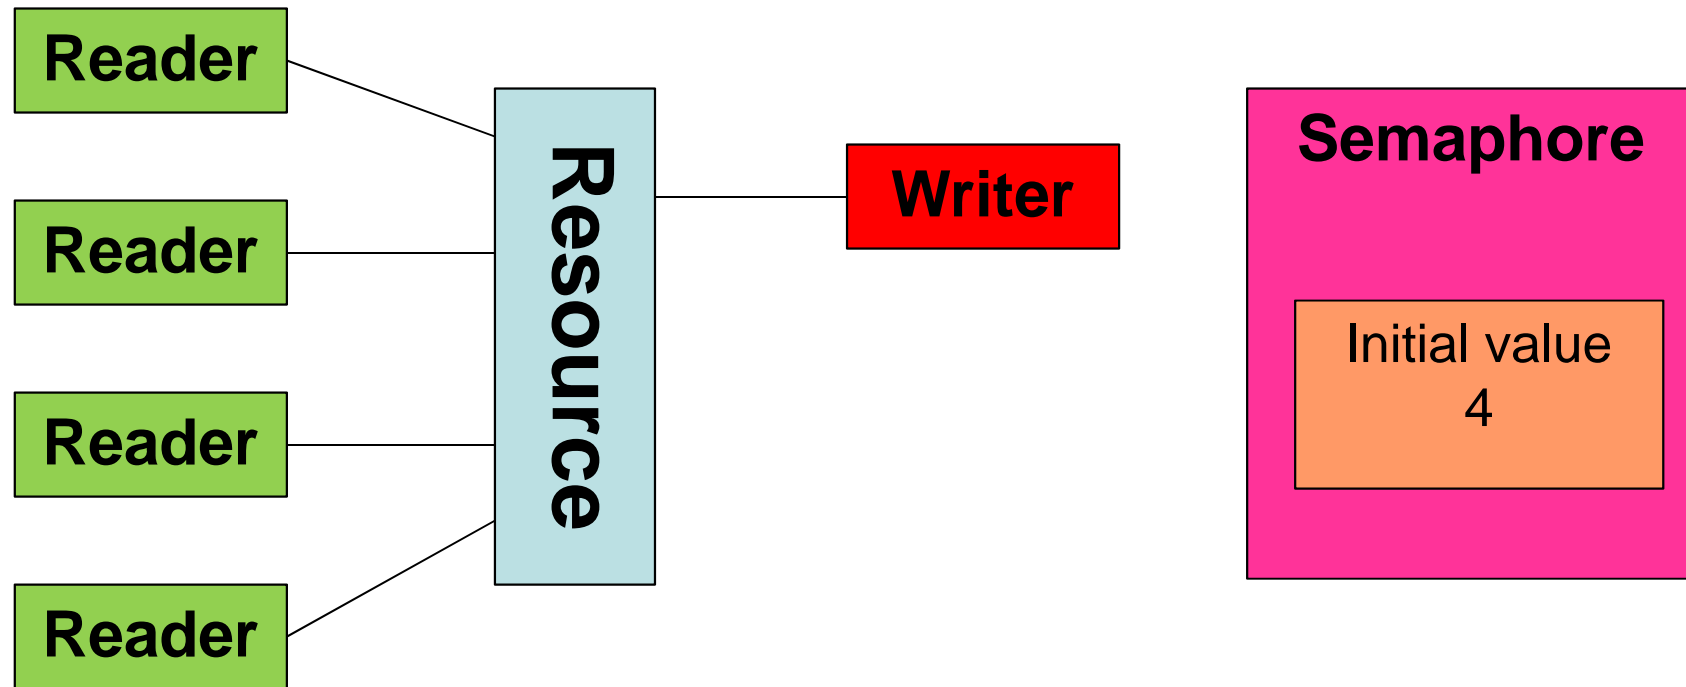
Single writer semaphore solution?



- Assume four readers and one writer
- What can we do with a generic semaphore?

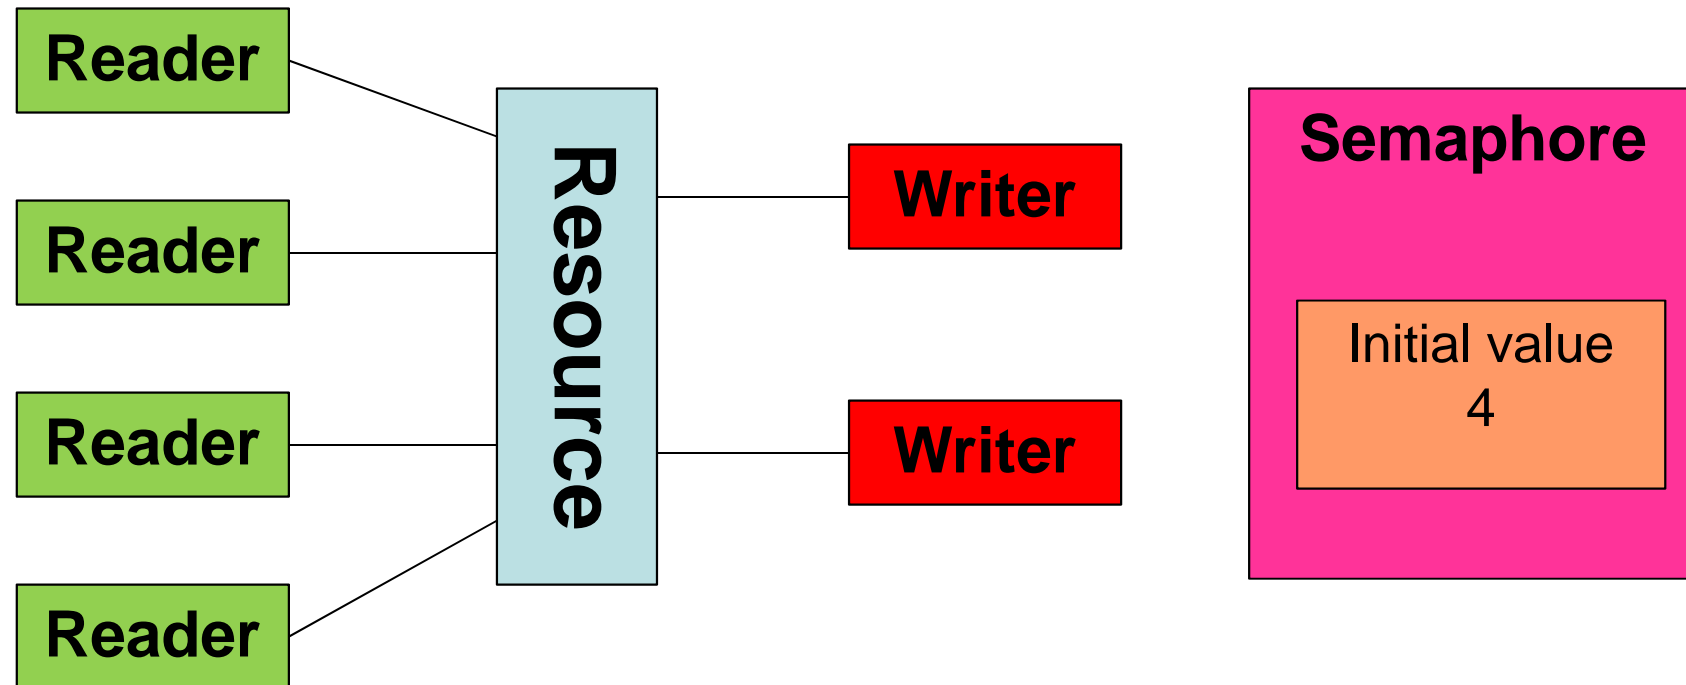
ANSWER

Simple solution



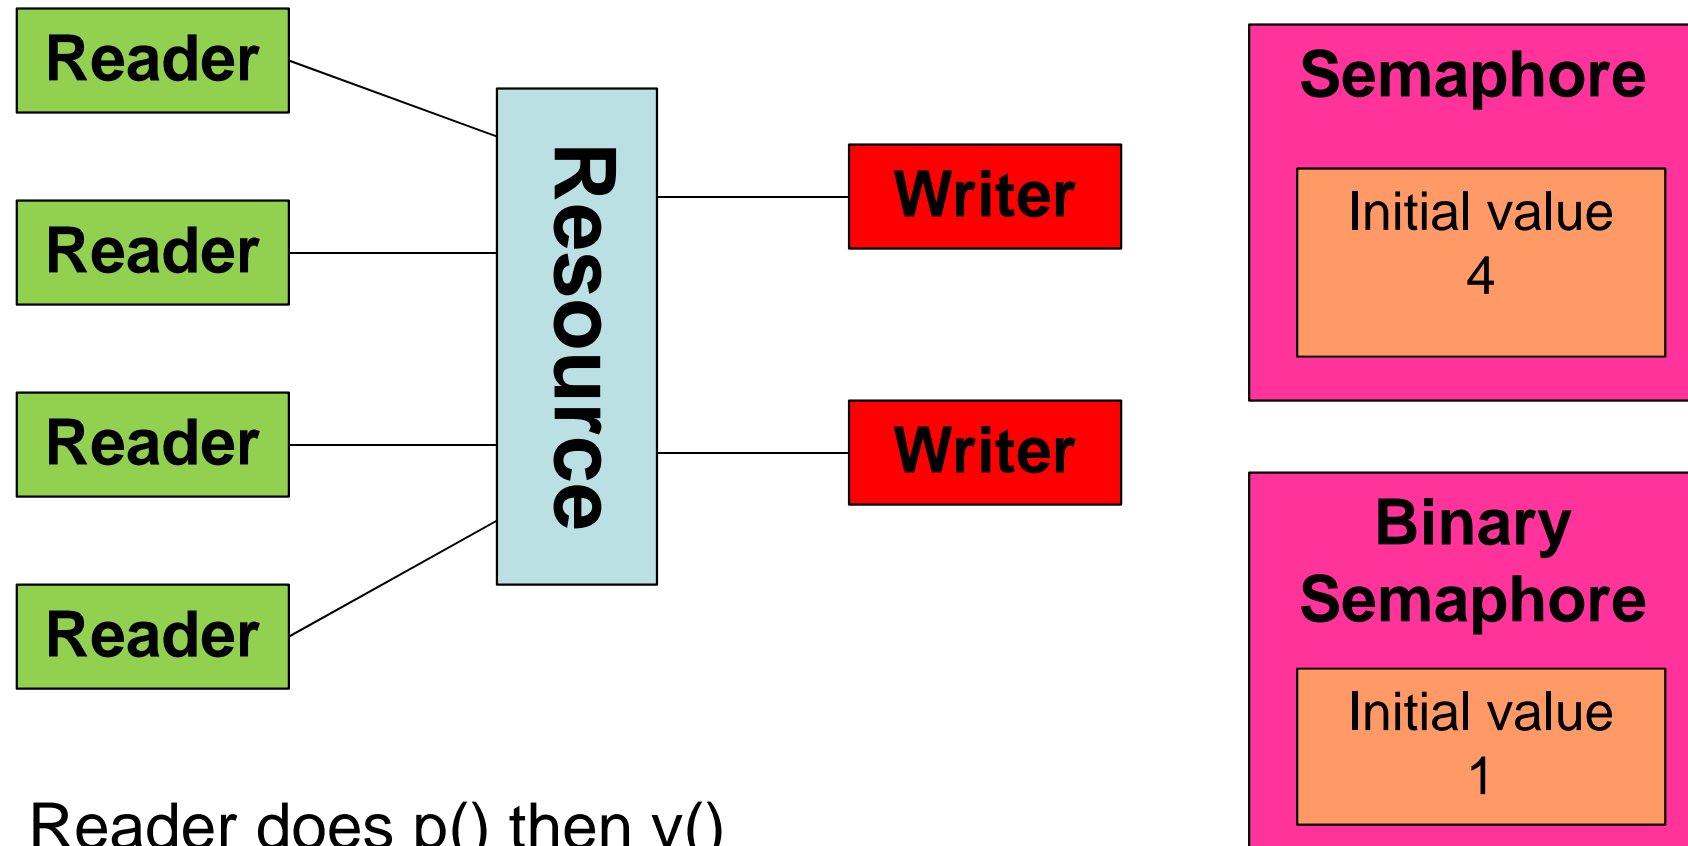
- Initialise semaphore to 4 (or $n = \text{\#readers}$)
- Each reader does $p()$ then $v()$
- Writer does $p()$ 4 times, then $v()$ 4 times

What if there were 2 writers?



- Initialise semaphore to 4 (or n)
- Reader does p() then v()
- Each writer does p() 4 times, then v() 4 times
- **Does it work? If not how could we fix it?**

Potential fix: lock for the writers?

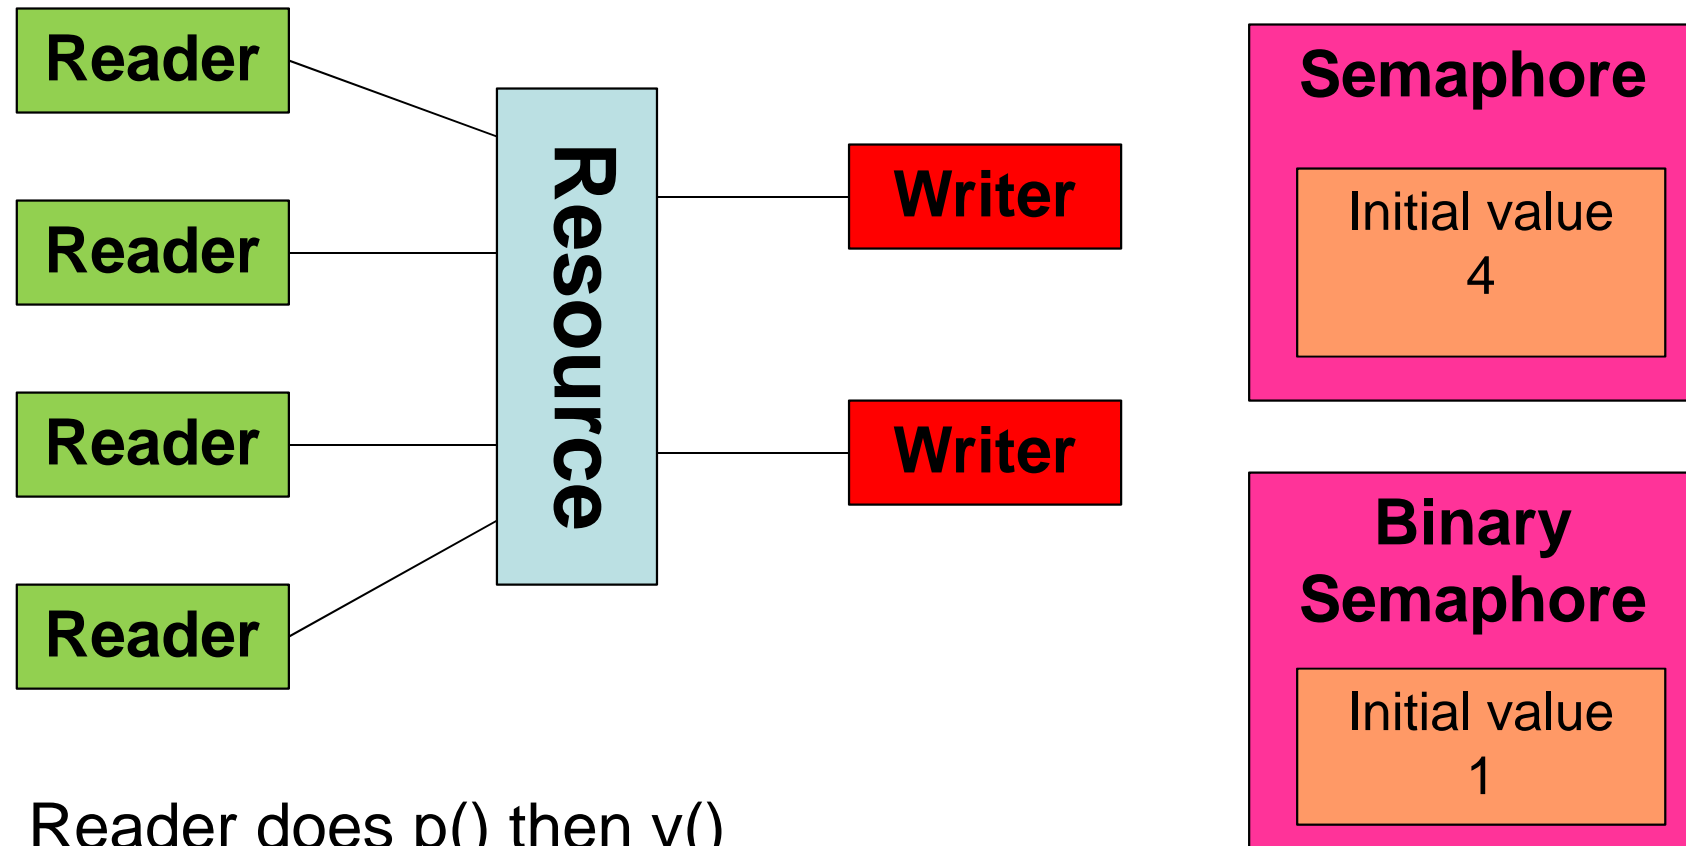


- Reader does $p()$ then $v()$
- Writer does $p()$ on binary semaphore, then $p()$ 4 times on original, then $4 \times v()$, then $v()$ on binary semaphore
- Any problem with this?

Consider various cases

- Like traces at the thread rather than operation level
- **W1 R1 R2 R3 R4 R5 R6 W2**
- **W1 W2 R1 R2 R3 R4 R5 R6**
- **R1 R2 R3 W1 W2 R4 R5 R6**
- Assume FIFO queue on the semaphore
 - Which is probably not realistic in most implementations

Reverse the unlock order



- Reader does $p()$ then $v()$
- Writer does $p()$ on binary, then $p()$ 4 times on original, then $v()$ on binary semaphore, then 4x $v()$ on original
- Consider: **W1 W2** R1 R2 R3 R4 R5 R6


Overview of semaphore solution

- With one writer: a single semaphore works for multiple readers or one writer at once
- As soon as you have multiple writers there are problems
 - E.g. They can get half of the semaphore count each
- You could implement mutual exclusion around the writers, so that only one can go for it at once
 - But that can deprioritise the multiple writers


Monitor solution?

Structure of a monitor

- Private data
- Public methods
- Locks
- Condition variables




Could we use a monitor for this?



- Simple implementation:
 - Data = the resource to protect
 - Mutual exclusion is guaranteed by the monitor
 - Could we do better?

Use the monitor to arbitrate access



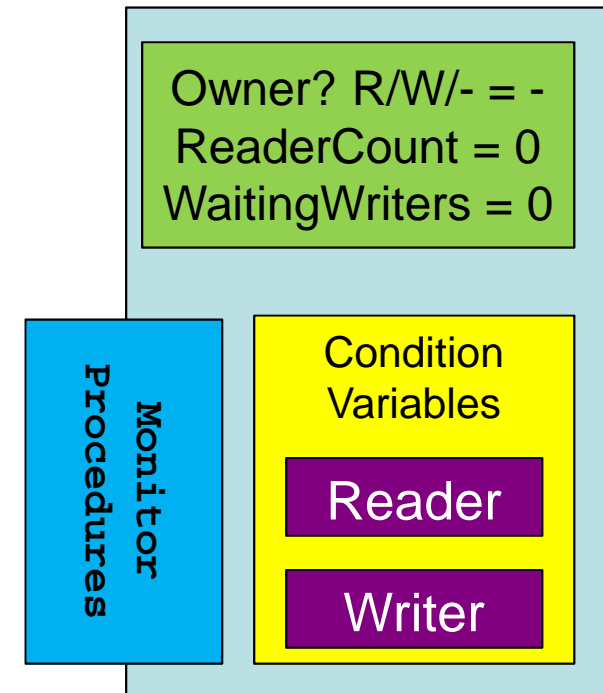
- Better implementation – monitor just arbitrates access
- Resource is global and shared
- All access has to go through the monitor at the start and end, to request access and to release access

Using a monitor to arbitrate access

- All of these are within the monitor procedures, to ensure mutual exclusion
- Four protocols/procedures are needed:
 - Reader entry protocol – reader wants to read
 - Reader exit protocol – reader finishes reading
 - Writer entry protocol – writer wants to write
 - Writer exit protocol – writer finishes writing

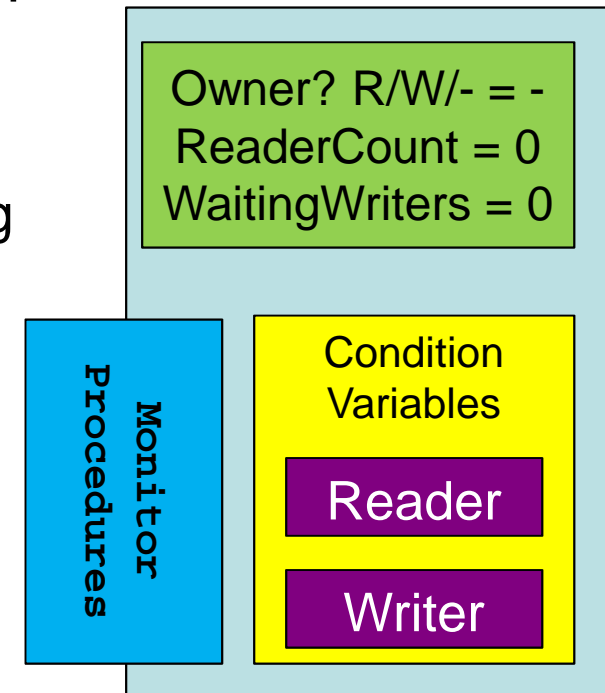
Reader protocols

- Reader requests access: (entry protocol)
 - If a writer currently has the resources:
 - Wait on condition variable for readers
 - When awoken re-check this condition
 - If no reader currently has the resources:
 - Record that a reader has it now
 - Increment the count of active readers
- Reader finishes: (exit protocol)
 - Decrement count of readers
 - If count is now zero then:
 - Clear that a reader has it
 - Notify the writer condition variable (may be a writer waiting)



Writer protocols

- **Writer requests access: (entry protocol)**
 - If a reader or writer currently has the resource:
 - Increment count of waiting writers
 - Wait on condition variable for writers
 - When awoken, decrement count of waiting writers then re-check the condition above
 - If we got here then no thread has access
 - Record that a writer has it
- **Writer finishes: (exit protocol)**
 - Clear that a writer has it
 - Notify (all) reader and writer condition variables



- **Note:** We could prefer to signal one condition variable over the other

Consider various cases

- **W1 R1 R2 R3 R4 R5 R6 W2**
- **W1 W2 R1 R2 R3 R4 R5 R6**
- **R1 R2 R3 W1 W2 R4 R5 R6**

Readers' Preference protocol

- Given a sequence of read and write requests which arrive in the following order:
 - $R_1 R_2 W_1 R_3 \dots$
- In a *Readers' Preference* protocol: R_3 takes priority over W_1
 - $R_1 R_2 \textcolor{red}{R_3} W_1 \dots$
- Writer may have to wait arbitrarily long time before getting access

Next Lecture

- Overview of what we have seen so far
 - And of hardware support for features
- Considering larger problems
 - And how we can solve them using what we have seen so far