

Introduction to Operating Systems

OPS Lecture 2, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

Remarks


Questions after the last lecture

- A **permissions problem** on Moodle meant that the slides were inaccessible to G52OSC \Rightarrow resolved
- We will make the **documentation for the G52OSC lab sessions** and **coursework** available to **G53OPS**
- For G52OSC, 3 out of 5 questions (2 CON, 3 OPS) for the exam with **at least one concurrency question \Rightarrow at least one OPS question**

Recap

Last Lecture

- 1 The **code** we write heavily **uses operating system functionality**
- 2 The **operating systems provides abstractions** (e.g., hides away hardware details) and **manages resources**
- 3 **Multi-programming** results in “a few” complications, e.g. **context switches** and **deadlocks**



Goals for Today

Overview

- 1 **G51CSA** from an OPS point of view
- 2 Interrupts, address spaces and context switches
- 3 **OS Structures/implementation**

Computer Hardware

G51CSA in a nutshell

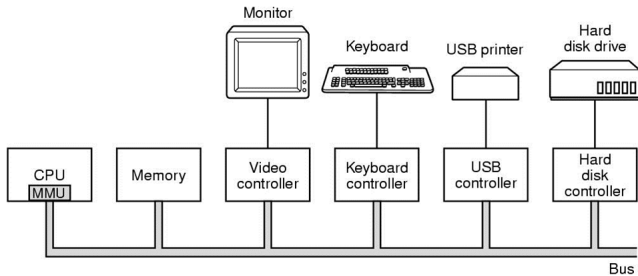


Figure: Simplified computer model (Tanenbaum, 2014)

Computer Hardware

CPU Design

- A CPU's basic cycle consist of **fetch**, **decode**, and **execute** (pipelines, or superscalar), and uses its own **instruction set**




Figure: Processor Architectures (Tanenbaum, 2014)

- Each CPU has a set of **registers** used for **data** or for special functions (e.g. **program counter**, **program status word** – mode bit)
 - The compiler/programmer decides what to keep in the registers
- **Context switching** must save and restore the CPU's internal state, including its **registers**

Computer Hardware

Memory Management Unit

```
1  #include <stdio.h>
2
3  int iVar = 0;
4  main() {
5      int i = 0;
6      while(i < 10) {
7          iVar++;
8          sleep(2);
9          printf("Address:%x; Value:%d\n",&iVar, iVar);
10         i++;
11     }
12 }
```

Computer Hardware

Memory Management Unit

- Modern computers use **logical** and **physical** memory addresses:
 - Every process has a logical address space – $[0, MAX_{64}]$ (theoretically)¹
 - The machine has a physical address space – $[0, MAX]$ (MAX determined by the amount of physical memory)
- **Address translation** takes place in the memory management unit (MMU)
 - $physical = f(logical)$




Figure: Address Translation

¹ $MAX_{64} = 16$ exbibyte, 1 exbibyte = 260 bytes = 1152921504606846976bytes = 1024pebibytes

Computer Hardware

Memory Management Unit (Cont'ed)

- **Context switching** invalidates the MMU (and registers, and cache) → they must **save** and **restore** the CPU's internal state





Figure: Address Translation

Computer Hardware

Context Switching: Timer Interrupts

- Interrupts are responsible for **temporarily pausing** a process's normal operation
- Different types of interrupts exist, including:
 - Timer interrupts by **CPU clock**
 - **I/O interrupts** for **I/O completion** or error codes
 - **Software generated**, e.g. errors and exceptions
- **Context switches** can be initiated by **timer interrupts** after a “**set time**”



Defining Operating Systems

A Resource Manager (Timer Interrupts)

1. Timer generates an interrupt
2. CPU finishes current instruction and tests for interrupt
4. Transfer to interrupt service routine
 - Hardware saves current process state (PSW, program counter)
 - Set program counter to interrupt service routine
 - Save registers and other state information
6. Carry out interrupt service routine (scheduler)
7. Restore next process to run

Computer Hardware

Moore's "law"

Moore's law

"The number of transistors on an integrated circuit (chip) doubles roughly every two years"

- Closely linked, but **not necessarily related to performance** (performance roughly doubles until 2003)
- Moore's still continuing, but the **"power wall"** slows performance improvements of single core/single processor systems
 - A few cores for multiple "programs" is easy to justify
 - How to use **massively parallel** computers/CPU's/many core machines
 - Can we **extract parallelism automatically**, can we implement **parallelism at the lowest level** (similar to multiprogramming)

Computer Hardware

General Motors vs. Microsoft

- Bill Gates:

*"If GM had kept up with technology like the computer industry has, we would all be **driving \$25.00** cars that got **1,000 miles to the gallon**"*

- General Motors:

- *"For no reason whatsoever, your car would **crash twice a day**"*
- *"The airbag system would ask **"Are you sure?"** before deploying. "*
- *"Occasionally, for no reason whatsoever, your **car would lock you** out and refuse to let you in until you simultaneously lifted the door handle, turned the key and grabbed hold of the radio antenna"*
- *"You'd have **to press the "Start" button to turn the engine off.**"*
- More at: <http://mistupid.com/jokes/msvgm.htm>




Figure: Volkswagen XL1 (Wikipedia)

Computer Hardware

Multi-core, hyperthreaded processors

- Modern CPUs contain **multiple cores** and are often **hyper-threaded**
- **Evolution in hardware** has implications on **operating system design**
 - XP did not support **multi processor architectures**
 - **Process scheduling** needs to account for **load balancing** and **CPU affinity**²
 - Cache **coherency** becomes important



²Last year's exam: Describe how, in your opinion, recent developments in computer architecture and computer design have influenced operating system design

Computer Hardware

Memory

- **Memory hierarchies** used to balance **cost** and **performance**
 - Fast and expensive memory is used for **caching**
 - Slow and inexpensive memory is used for **long term storage**
 - Memory includes, registers, L1/L2 cache, **main/core memory**, **disk**, etc.
- **L2 Cache** can be **shared** or **dedicated** to individual cores
- **Cache management** is mainly done by **hardware**
- The CPU can only **access main memory directly** (i.e. files have to be brought into memory first) **context switches**

Computer Hardware

I/O Devices

- **Device driver** interacts with the **controller**, controller interacts with the **device** (e.g., disk controller)
- The operating system/device driver typically **communicates with the controller through registers**
- I/O can take place through:
 - **Busy waiting**
 - **Interrupt** based
 - Direct memory access (using **DMA** chip)

Micro Kernels

Operating System Structure

- Operating Systems contain **a lot of functionality**, including:
 - Processes, process management, process scheduling, context switching, etc.
 - Memory management, virtual memory, etc.
 - File Systems
 - I/O Management
- \Rightarrow How are they **logically organised**?
 - Micro kernels
 - Monolithic

Micro Kernels

Minimal Kernel Functionality

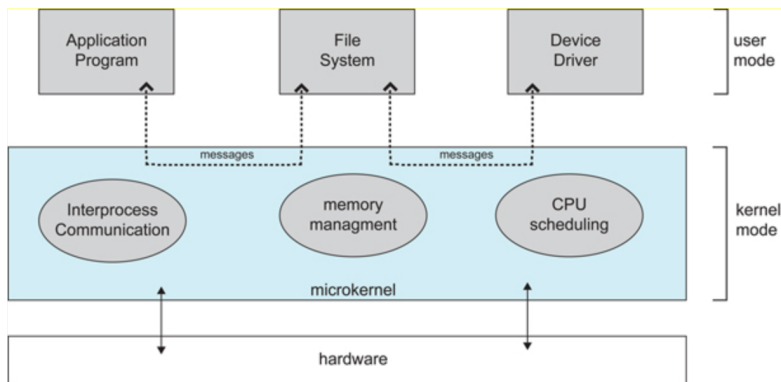


Figure: Structure of a Micro Kernel (Silberschatz, 2008)

Micro Kernels

Minimal Kernel Functionality

- All **non-essential functionality** are **extracted** from the kernel
 - **Communication, memory management** and **CPU scheduling** are likely to be included
 - The **file system, GUI, device drivers** are likely to be user processes
- Micro kernels are more **easy to extend**, more **portable**, and usually more **reliable**
- Frequent **system calls** and **kernel traps** cause significant **overhead**
- Some Unix version, Mac OS X, Minix, and early versions of Windows (NT4.0) were (partially) micro kernels

Monolithic Systems

The big mess

- All procedures are **linked together** into one **single executable** running in **kernel mode**
- Monolithic kernels are **difficult to maintain**
- Current versions of Windows, Linux are implemented as monolithic Kernels

Recap

Take-Home Message

- Operating Systems are heavily linked to **computer architecture**
- **Interrupts, address translation**
- **OS structures**

Next Lecture

Content

- Processes!