

Threads and Multi-processor/core Scheduling

OPS Lecture 6, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

1 Multi-level **feedback queues**

2 Processes vs **Threads**:

- Parallel **execution traces** (i.e., must be scheduled) that **share resources**
- Thread control blocks, containing less info than process control blocks
- **Faster** to create/switch, alternate **I/O and CPU bound, parallelism**
- They **abstract parallelism** from the programmer

Goals for Today

Overview

- 1 Different thread implementations
- 2 Multi-processor and multicore scheduling

OS Implementations

Thread Types

- **User** threads
- **Kernel** threads
- **Hybrid** implementations

User Threads

Many-to-One

- **Thread management** (creating, destroying, scheduling, thread control block manipulation) is carried out **in user space** with the help of a **user library**
- The process maintains a **thread table** managed by the **runtime system** without the **kernel's knowledge**
 - Similar to **process table**
 - Used for **thread switching**
 - Tracks thread related information
- They can be implemented on **OS** that **does not support multi-threading**

User Threads

Many-to-One

- Advantages:

- Threads are in user space (i.e., **no mode switches** required)
- **Full control** over the thread scheduler
- **OS independent** (threads can run on OS that do not support them)
- The runtime system can **switch local blocking threads** in user space (program counter, store and reload registers)

- Disadvantages:

- **Blocking system calls** suspend the entire process (User threads are onto a single process, managed by the kernel)
- **Page faults** result in blocking the process
- **No true parallelism** (a process is scheduled on a single CPU)
- **Clock interrupts** are non-existent (i.e. user threads are non-preemptive)
- Remember: threads are particularly **useful** when **code often blocks**!

User Threads

Many-to-One




Figure: User threads (Tanenbaum 2014)

Kernel Threads

One-to-One

- The **kernel manages** the threads, user application accesses threading facilities through **API** and **system calls**
 - **Thread table** is in the kernel, containing thread control blocks (subset of process control blocks)
 - **Thread blocks**, kernel chooses thread from same or different process (\leftrightarrow user threads)
- Advantages:
 - **True parallelism** can be achieved
 - **No non-blocking** system calls needed
 - No run-time system needed
- Frequent **mode switches** take place, resulting in lower performance
- Windows and Linux apply this approach

Kernel Threads

One-to-One




Figure: Kernel threads (Tanenbaum 2014)

Hybrid Implementations

Many-to-Many

- User threads are **multiplexed** onto kernel threads
- Kernel sees and schedules the kernel threads (a limited number)
- User application sees user threads and creates/schedules these (an “unrestricted” number)






Figure: Kernel threads (Tanenbaum 2014)

Comparison


Thread Implementations



(a) Pure user-level



(b) Pure kernel-level



(c) Combined

Figure: Comparison (Stallings)

Exam 2013-2014: In which situations would you favour user level threads? In which situation would you definitely favour kernel level threads?

Performance

User Threads vs. Kernel Threads vs. Processes

- Null fork: the overhead in creating, scheduling, running and terminating a null process/thread
- Signal wait: overhead in synchronising threads

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Figure: Comparison, in μs (Stallings)

Scheduling on Multi-core/Processor Systems

Einstein's influence on multi-core processors :-)

- Electrical signals cannot propagate any faster than the **speed of light** (30cm/nsec - in copper 20cm/nsec)
- A 10GHz clock speed, i.e. 10×10^9 cycles per second, means the **maximum distance is 2cm**, and **2mm for 100GHz**
- The faster a computer runs, the more **heat** it dissipates
- This imposes a **fundamental limit on clock speed**, hence, **parallelism** is exploited to increase computational power


Multi-processor Scheduling

Scheduling Decisions

- **Single processor** machine: **which thread** to run next (one dimensional)
- Scheduling decisions on a **multi-processor/core** machine include:
 - Which process/thread to run **where**, i.e., which CPU?
 - Which process/thread to run **when**?
- Threads can be **related** or **unrelated** (multiple users)
 - **Related** threads: e.g. the **same process** creates **multiple threads** that **communicate** with one another and **ideally run** together (e.g. search algorithm)
 - **Unrelated** threads: e.g. threads that belong to **different processes**, possibly started by **different users** running **different programs**

Multi-processor Scheduling

Related Processes



Best Value
Found

Scheduling Unrelated Threads

Shared Queues

- A single or multi-level queue **shared** between all CPUs
- Advantage: automatic **load balancing**
- Disadvantages:
 - **Contention** for the queues (locking is needed)
 - Mainly applicable for **unrelated threads/processes**
 - *"All CPUs are equal, but some are more equal than others"*: does not account for **processor affinity**:
 - **Cache** becomes invalid when moving to a different CPU
 - Translation look aside buffers (**TLBs**) become invalid
- Windows will allocate the **highest priority threads** to the individual CPUs/cores

Scheduling Unrelated Threads

Private Queues

- Each process/thread is assigned to a queue **private** to an individual CPU (\Rightarrow two level scheduling)
- Advantages:
 - **CPU affinity** is automatically satisfied
 - **Contention** for shared queue is minimised
- Disadvantages: less **load balancing**
- **Push** and **pull migration** between CPUs is possible

Scheduling Related Threads

Working Together

- Threads belong to the same process and are **cooperating**, e.g. they **exchange messages** or **share information**
- The aim is to get threads **running**, as much as possible, at the **same time** across **multiple CPUs**
- Approaches include:
 - **Space** sharing
 - **Gang** scheduling

Scheduling Related Threads

Space Scheduling


- Approach:
 - N threads are allocated to N **dedicated CPUs**
 - N threads are kept waiting until N CPUs are available
 - **Non-preemptive**, i.e. blocking calls result in **idle CPUs** (less context switching overhead but results in CPU idle time)
 - Scheduling decisions are made using **scheduling algorithms** (e.g., FCFS, PQ)
- The number N can be **dynamically adjusted** to match processor capacity

Scheduling Related Threads

Gang scheduling

Imagine:

- Process A has thread A1 and A2, A1 and A2 cooperate
- Process B has thread B1 and B2, B1 and B2 cooperate
- The scheduler selects A1 and B1 to run first, then A2 and B2, and A1 and A2, and B1 and B2 run on different CPUs
- They try to send messages to the other threads, which are still in the ready state



Scheduling Related Threads

Gang scheduling

- Time slices are **synchronised** and the scheduler **groups threads** together to run simultaneously (as much as possible)
- A **preemptive** algorithm
- Blocking threads** result in idle CPUs

		CPU					
		0	1	2	3	4	5
Time slot	0	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
	1	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
	2	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
	3	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆
	4	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
	5	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
	6	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
	7	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆

Summary

Take Home Message

- User, kernel, and hybrid **thread implementations**
- **Multi-processor/core** scheduling is “a bit different” (load balancing, processor affinity, etc.)
 - **Related** and **unrelated** threads
 - **Shared** or **private** queues
 - **Space** scheduling or **gang** scheduling