

Memory Management

OPS Lecture 9, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

Recall

Last Lecture

- **Dynamic relocation** and protection \Rightarrow **base** (offset) and **limit** registers (logical/physical address)
- **Dynamic partitioning** \Rightarrow **internal** and **external fragmentation**
- **Memory management** using **linked lists** and bitmaps

Dynamic Partitioning

Example

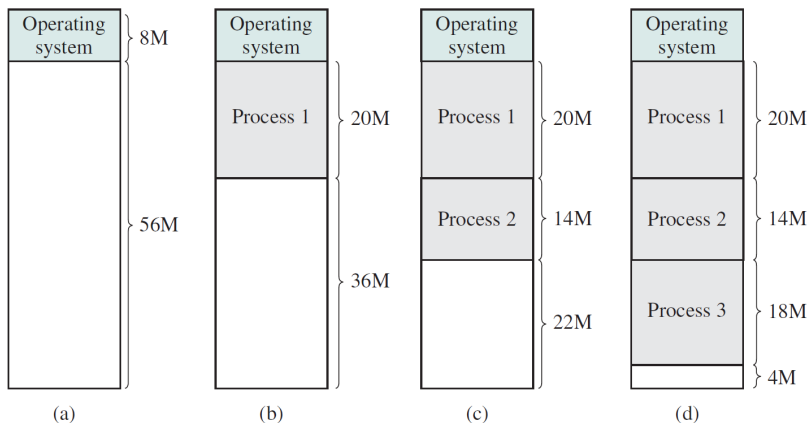


Figure: Dynamic partitioning, a problem occurs when process 2 grows (from Stallings)

Dynamic Partitioning

Swapping: Example

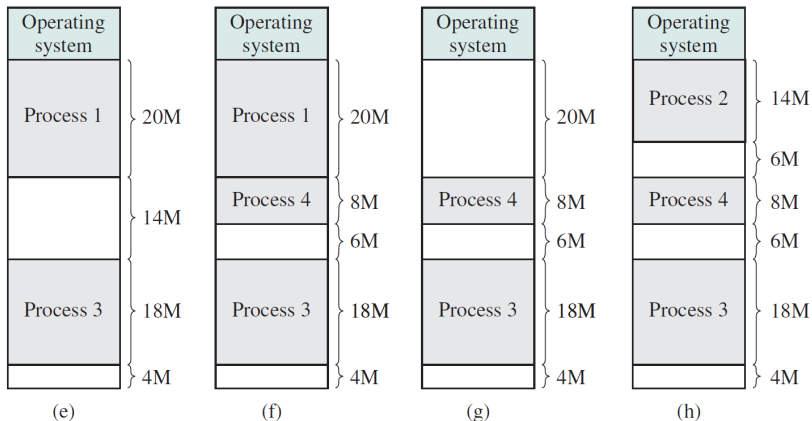


Figure: Dynamic partitioning, swapping results in external fragmentation (from Stallings)

Dynamic Partitioning

Memory Management

- How to keep track of **available memory**
 - Bitmaps
 - **Linked lists**
- The operating system is responsible for:
 - Applying strategies to (quickly) **allocate processes** to available memory (“holes”)
 - Managing **free space**




Figure: Memory management with linked lists

Dynamic Partitioning

Allocating Available Memory: First Fit

- The linked list is **initialised** to a **single link** of the entire memory size, flagged as “free”
- **First fit** starts scanning **from the start** of the linked list until a link is found which represents **free space of sufficient size**
 - If requested space is **exactly the size** of the free space, all the space is allocated (i.e., no **internal** fragmentation)
 - Else, the free link is **split into two**:
 - The first entry is set to the **size requested** and marked “**used**”
 - The second entry is set to **remaining size** and marked “**free**”

Dynamic Partitioning

Allocating Available Memory: Next Fit

- As a minor variation of first fit, the **next fit algorithm** maintains a record of where it got to:
 - The next time a block is requested the algorithm **restarts its scan from where it stopped last time**
 - The idea is to give an **even chance to all of memory to getting allocated**, rather than concentrating at the start
- However, simulations have shown that next fit actually gives **worse performance** than first fit!

Dynamic Partitioning

Allocating Available Memory: Best Fit

- **First fit** just looks for the **first available hole**
 - It doesn't take into account that there may be a **hole later** in the list that **exactly(-ish)** fits the requested size
 - First fit **may break up a big hole** when the right size hole exists later on
- The **best fit algorithm** always **searches the entire linked** list to find the **smallest hole big** enough to satisfy the memory request
 - However, it is **slower** than first fit because of searching
 - Surprisingly, it also results in **more wasted memory** because it tends to fill up memory with tiny (useless) holes

Dynamic Partitioning

Allocating Available Memory: Worst Fit

- **Tiny holes** are created when **best fit** breaks a hole of nearly the exact size into the required size and whatever is left over
- To get around the **problem of tiny holes**
 - How about always **taking the largest available hole** and breaking that up
 - The idea being that the **left over part will still be a large** and therefore **potentially useful** size
 - This is the **worst fit algorithm**
- Unfortunately, simulations have also shown that worst fit is **not very good either!**

Dynamic Partitioning

Allocating Available Memory: Summary

- **First fit:** allocate **first block** that is **large enough**
- **Next fit:** allocate **next block** that is large enough, i.e. **starting from the current location**
- **Best fit:** choose block that **matches** required size **closest** - $O(N)$ complexity
- **Worst fit:** choose the **largest possible block** - $O(N)$ complexity

Dynamic Partitioning

Allocating Available Memory: Quick Fit and Others

- As yet another variation, **multiple lists of different (commonly used) size blocks** can be maintained
 - For example a separate list for each of 4K, 8K, 12K, 16K, etc., holes
 - **Odd sizes** can either go into the **nearest size** or into a **special separate list**
- This scheme is called **quick fit**, because it is **much faster** to find the required size hole, however it still has problem of creating **many tiny holes**

Dynamic Partitioning

Allocating Available Memory: Quick Fit and Others

- As yet another variation, **multiple lists of different (commonly used) size blocks** can be maintained
 - For example a separate list for each of 4K, 8K, 12K, 16K, etc., holes
 - **Odd sizes** can either go into the **nearest size** or into a **special separate list**
- This scheme is called **quick fit**, because it is **much faster** to find the required size hole, however it still has problem of creating **many tiny holes**
- **Finding neighbours** for coalescing becomes more difficult/time consuming

Dynamic Partitioning

Managing Available Memory: Coalescing

- **Coalescing** (joining together) takes place when two **adjacent entries** in the linked list **become free**
 - There may be three adjacent free entries if an in-use block that is in-between two free blocks is freed
- Both **neighbours** are examined when a **block is freed**
 - If either (or both) are **also free**
 - Then the two (or three) **entries are combined** into one larger block by adding up the sizes
 - The earlier block in the linked list gives the **start point**
 - The **separate links are deleted** and a single link inserted




Figure: Memory management with linked lists

Dynamic Partitioning

Managing Available Memory: Compacting

- Even with coalescing happening automatically, **free blocks** may still **distributed across memory**
 - \Rightarrow **Compacting** can be used to join free and used memory (but is **time consuming**)
- **Compacting is more difficult and time consuming** to implement than coalescing (processes have to be moved)
 - Each process is **swapped out & free space coalesced**
 - Process swapped back in at lowest available location

Contiguous Allocation Schemes

Overview and Shortcomings

- Different contiguous memory allocation schemes have different advantages/disadvantages
 - **Mono-programming** is easy but does result in **low resource utilisation**
 - **Fixed partitioning** facilitates **multi-programming** but results in **internal fragmentation**
 - **Dynamic partitioning** facilitates **multi-programming**, reduces **internal fragmentation**, but results in **external fragmentation** (allocation methods, coalescing, and compacting help)
- Can we design a memory management scheme that **resolves the shortcomings** of contiguous memory schemes

Paging

Principles

- **Paging** uses the principles of **fixed partitioning** and **code re-location** to devise a new **non-contiguous management scheme**, however:
 - Memory is split into much **smaller blocks** and **one or more blocks** are allocated to a process
 - e.g., a 11kb process would take up 3 blocks of 4 kb
 - These blocks **do not have to be contiguous in main memory**, but the process still **perceives them to be contiguous**
- Benefits compared to contiguous schemes include:
 - **Internal fragmentation** is reduced to the **last “block”** only
 - There is **no external fragmentation**, since physical blocks are **stacked directly onto each other** in main memory

Paging

Principles (Cont'ed)

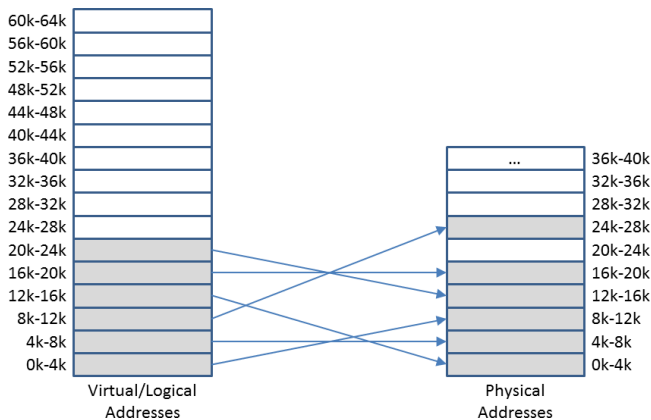


Figure: Paging in main memory with multiple processes

Paging

Principles (Cont'ed)




Figure: Concept of Paging (Stallings)

Paging

Principles: Definitions

- A **page** is a small block of **contiguous memory** in the **logical address space**, i.e. as seen by the process
- A **frame** is a **small contiguous block** in **physical memory**
- Pages and frames (usually) have the **same size**:
 - The size is usually a power of 2
 - Sizes range between 512 bytes and 1Gb

Paging

Relocation

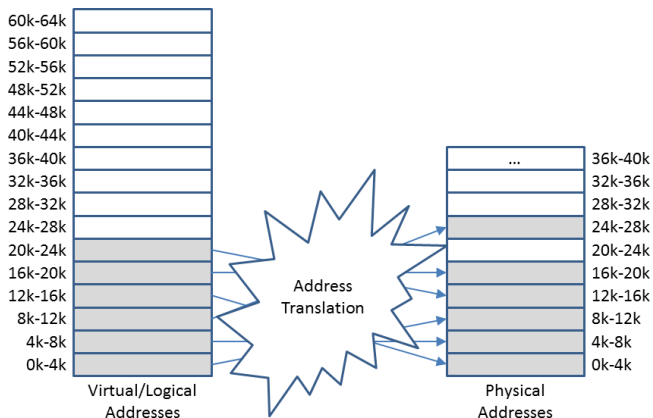


Figure: Address Translation

Paging

Relocation

- **Logical address** (page number, offset within page) needs to be **translated** into a **physical address** (frame number, offset within frame)
- **Multiple “base registers”** will be required:
 - Each logical page needs a **separate “base register”** that specifies the start of the associated frame
 - I.e, a **set of base registers** has to be maintained for each process
- The base registers are stored in the **page table**

Paging

Relocation: Address Translation

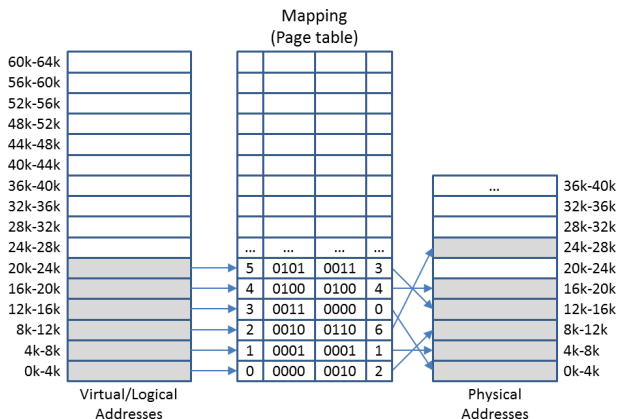


Figure: Address Translation

Paging

Relocation: Page Tables

- The page table can be seen as **a function**, that **maps the page number** of the logical address **onto the frame number** of the physical address
 - $\text{frameNumber} = f(\text{pageNumber})$
- The **page number** is used as **index to the page table** that lists the **number of the associated frame**, i.e. it contains the location of the frame in memory
- Every process has its **own page table** containing its own “base registers”
- The **operating system** maintains a **list of free frames**

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
 - The **right most m bits** that represent the **offset within the page (frame)**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
- The **offset** within the page and frame **remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**




Figure: Logical Address

Paging

Address Translation: Implementation

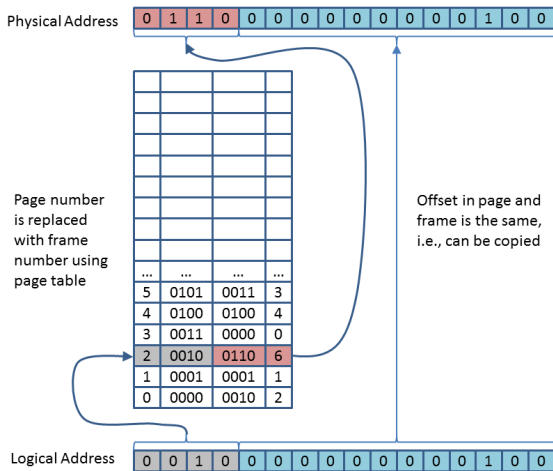


Figure: Address Translation

Paging

Relocation: Address Translation

- Steps in **address translation**:
 - 1 Extract the page number from logical address
 - 2 Use page number as an index to retrieve the frame number in the page table
 - 3 Add the “logical offset within the page” to the start of the physical frame
- **Hardware implementation** of address translation
 - 1 The CPU's memory management unit intercepts logical addresses
 - 2 Address translation using page table as above
 - 3 The resulting physical address is put on the memory bus

Recap

Take-Home Message

- Memory **allocation**, **coalescing** and **compacting** in dynamic partitioning
- **Paging**, **page tables**, and **address translation**

Paging

Address Translation: Examples

- Let us assume 4KB pages (2^{12}), leaving $2^4 = 16$ pages
- What is the physical address of 0, 8192, 20500 using the page table below

	Pages		Frames	
0	0000	0010	2	
1	0001	0001	1	
2	0010	0110	6	
3	0011	0000	0	
4	0100	0100	4	
5	0101	0011	3	
...	
...	
...	
9	1001	0101	5	
...	
11	1011	0111	7	
...	

Table: Page Table

Paging

Address Translation: Examples

- Virtual address 0 falls in page with index 0 which is mapped onto frame with index 2, starting at 2×4096 , i.e., the physical address is 8192
- Virtual address 8192 falls in page with index 2 which is mapped onto frame with index 6, starting at 6×4096 , i.e., the physical address is 24576
- Virtual address 20500 falls in page with index 5 ($20500 > 5 \times 4096$) which is mapped onto page with index 3, starting at 3×4096 , i.e., the physical address is $12288 + 20 = 12308$