

File Systems

OPS Lecture 13, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

- **Traditional hard drives** are still used for most of secondary storage, **solid state disks** are becoming more popular
 - They require less disk scheduling (reads are equally fast, writes are clustered)
 - They must balance wear and tear!
- Traditional hard drives have **physical limitations** (seek times, rotational latency) \Rightarrow **disk scheduling** can help to minimise the impact of this
- This **influences file system design**

Goals for Today

Overview

- ➊ User view of file systems
 - System calls
- ➋ Implementation view of file systems
 - Disk and partition layout
 - File tables
 - Free space management
 - ...

File Systems

Different Views

- A **user view** that defines a file system in terms of the **abstractions** that the operating system provides
- An **implementation view** that defines the file system in terms of its **low level implementation**

User View

Definition

- Important aspects of the user view include:
 - The **file abstraction** which hides away implementation details to the user (similar to processes and memory)
 - File **naming policies** (abstracts storage details), **extensions**, “user file **attributes**” (e.g. size, protection, owner, protection, dates)
 - There are also **system attributes** for files (e.g. non-human readable file descriptors (similar to a PID), archive flag, temporary flag, etc.)
 - **Directory structures** and organisation
 - **System calls** to interact with the file system
- The **user view** defines what the file system looks like to regular users (and programmers) and relates to **abstractions**

Implementation View

Definition

- Important aspects of the **implementation view** include:
 - **Partitions** and their implementation
 - The **file system implementation** (FAT-16, FAT-32, exFAT, UFS, NTFS, EXT2,3,4, ReFS, etc)
 - Logical to physical **mapping**
 - **Free space** management
 - File system consistency
- The **implementation view** relates to **low level details**, technical aspects, **data structures**, block sizes (influencing internal fragmentation), etc.

Files

Types

- Common **file types** include:
 - **Regular files** contain user data in **ASCII** or **binary** (well defined) **format**
 - **Directories** group files together (but are files on an implementation level)
 - **Character special files** are used to model **serial I/O devices** (e.g. keyboards, printers)
 - **Block special files** are used to model, e.g. hard drives
- Files are **sequential** or **random access**

System Calls

Types

- Two different categories of **system calls** exist:
 - **File manipulation**: open(), close(), read(), write(), ...
 - **Directory manipulation**: create(), delete(), readdir(), rename(), link(), unlink(), list(), update()
- System calls enable a **user application** to **ask the operating system** to carry out an **action** on its behalf (in kernel mode)
- In practice, a lot more happens underneath!

System Calls

Files

- `Open()` system call:
 - **Search directory** entries for the file location
 - Copy the directory entry to the **open file table** (specify access rights)
 - Increase the **open count**
 - `new Writer/Reader('`D:/.../.../.../file.txt`')`
- `close()` to free up the entry in the internal table

System Calls

Files (Cont'ed)

- System calls **after opening a file:**
 - `Write()`: write information to the file starting at the position of the current **write pointer**, update **write pointer**
 - `Read()`: read information starting at the current **read pointer**, store the information in the process memory space
- System calls **prior to opening a file:**
 - `Create()`: space is allocated, an **entry is made in a directory**, attributes are initialised by OS
 - `Delete()`: find the file, **remove directory entry** and release the file's space
- Others include `seek`, `append`, `getAttributes`, ...

File System Structures

Overview

- Different directory structures have been used over the years
 - **Single level**: all files in the same directory (reborn in consumer electronics)
 - **Two or multiple level directories** (hierarchical): tree structures
 - Absolute path name: from the root of the file system
 - Relative path name: the current working directory is used as the starting point
 - **Directed acyclic graph (DAG)**: allows files to be shared (i.e. **links** to files or sub-directories) but **cycles are forbidden**
 - **Generic graph structure** in which links and **cycles can exist**
- The use of **DAGs** and **generic graph structures** results in significant **complications** in the implementation

Directories

Possible Implementations




Figure: DAG Directory Implementation (Tanenbaum)

File System Structures

DAG and Graph Complications

- When searching the file system:
 - Cycles can result in **infinite loops**
 - Sub-trees can be **traversed multiple times**
- Files have **multiple absolute file names**
- **Deleting files** becomes a lot more complicated (i.e. links may no longer point to a file, **inaccessible cycles** may exist)
- A **garbage collection scheme** may be required to remove files that are no longer accessible from the file system tree (that are part of a **cycle only**)

Directories

System Calls

- Similar to files, **directories** are manipulated using **system calls**
 - `create/delete`: a new directory is created/deleted
 - `opendir, closedir`: add/free directory to/from internal tables
 - `readdir`, return the next entry in the directory file
 - Others: `rename, link, unlink, list, update`
- **Directories are special files** that **group files** together and of which the **structure is defined by the file system**
 - A bit is set to indicate that they are directories!

Directories

Structure

- Directories contain a list of **human readable file names** that are mapped on **unique identifiers** and **disk locations**
 - They provide a mapping of the logical file onto the physical location
- **Retrieving a file** comes down to using the path to search the directory file and retrieve the exact location of the file, e.g.:
 - E.g. first address and number of blocks
 - I-node number
- They can **store all file related attributes** (e.g. file name, disk address – Windows) or they can **contain a pointer** to the data structure that contains the details of the file (Unix)

Directories

Internal Structure

- Retrieving a file comes down to **searching a directory file** as fast as possible:
 - A **simple random order of directory** entries might be insufficient (search time is linear as a function of the number of entries)
 - Indexes or **hash tables** can be used
- Commonly accessed files/directories are **cached**

Directories

Possible Implementations

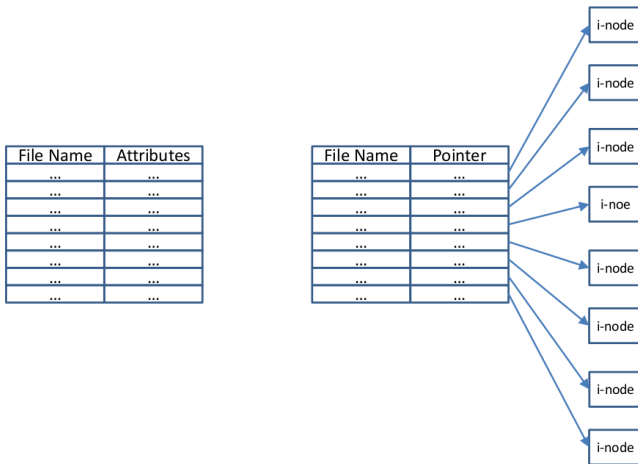


Figure: Directory Implementations

Implementation

Context

- Irrespective of the type of file system, a number of **additional considerations** have to be addressed, including
 - Disk **partitions**, **partition tables**, **boot sectors**, etc.
 - System wide and per process **file tables** (\Rightarrow process tables)
 - Free **space management** (\Rightarrow free memory)
- **Low level formatting** writes sectors to the disk, **high level formatting** imposes a file system on top of this (using **blocks** that can cover multiple **sectors**)

Hard Disk Structures

Partitions

- Disks are usually divided into **multiple partitions**
 - An independent file system may exist on each partition
- **Master is boot record** located at start of the entire drive:
 - Used to boot the computer (BIOS reads and executes MBR)
 - Contains **partition table** at its end with active partition
 - One partition is listed as **active** containing a boot block to load the operating system




Figure: Layout of a Disk

Partition Layouts

A Unix Partition

- The layout of a partition differs depending on the file system
- A UNIX partition contains:
 - The partition **boot block**:
 - Contains code to boot the operating system
 - Every partition has boot block – even if it does not contain OS
 - **Super block** contains the partition's details, e.g., partition size, number of blocks, I-node table size
 - **Free space management** contains, e.g., a bitmap or linked list that indicates the free blocks

Boot block	Super block	Free space MGT	I-nodes	Root dir	Files/directories
------------	-------------	----------------	---------	----------	-------------------

Figure: Layout of a Partition

Partition Layouts

A Unix Partition

- A UNIX partition contains (cont'ed):
 - **I-nodes** contains information on files, commonly maintained in I-nodes
 - **Root directory**
 - **Data:** files and directories




Figure: Layout of a Partition

Disk Space Management

Free Space Management

- Two methods are commonly used to keep track of free disk space: **bitmaps** and **linked lists**
 - Note that these approaches are very similar to the ones to keep track of free memory
- **Bitmaps** represent each block by a single bit in a map
 - The **size of the bitmap** grows with the size of the disk but is constant for a given disk
 - Bitmaps take **comparably less space than linked lists**, unless the disk is nearly full

Disk Space Management

Free Space Management

- Linked Lists:
 - Free blocks are used to hold the **numbers of the free blocks** (hence, they are no longer free)
 - Since the free list shrinks when the disk becomes full, this is not wasted space
 - **Blocks are linked together**, i.e., multiple blocks list the free blocks
 - The size of the list **grows with the size of the disk** and **shrinks with the size of the blocks**
- Linked lists can be modified by **keeping track of the number of consecutive free blocks** for each entry

Disk Space Management

Free Space Management




Figure: Free Block Management with Linked Lists (Tanenbaum)

File Tables

Implementation

- The OS maintains **two types of file tables**:
 - **System wide file table** contains **process independent information**
 - **Per process open file table** contains **process dependent information**
- **Process independent information** includes location on disk, file size, and “open count” (#processes that use the file)
- **Process dependent information** includes **read/write pointers**, a **reference** to the system wide file table

File Tables

Illustration

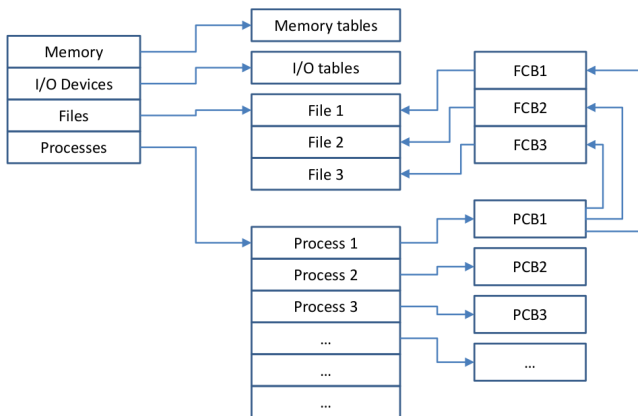


Figure: Illustration of File Tables

File Locking

Types

- Two different **lock types** exist:
 - A **shared lock** allows multiple processes to access a file at the same time
 - An **exclusive lock** allows at most one process to access the file at a time
- Locking enforcement can be:
 - **Mandatory**: no other process can access the file until the exclusive lock has been released, i.e., OS enforces locking (Windows)
 - Hold the lock only as long as necessary
 - Deadlock situations can occur
 - **Advisory locking**: other processes may have access to the file, i.e., OS does not enforce locking – programs themselves are responsible (UNIX)

File System Implementations

How are Files Stored

- File system **implementations**
 - ① Contiguous
 - ② Linked lists
 - ③ File Allocation Table (FAT)
 - ④ I-nodes
- File system **paradigms**, including journalling file systems, log structured file systems, virtual file system

Summary

Take-Home Message

- User vs. implementation view
- Implementation of files and directories
- System calls for file and directory management
- File tables, free space management, partitions, boot sectors, etc.