

File Systems

OPS Lecture 14, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

- ❶ **User view** of file systems
 - System calls
 - Structures, organisation, file types
- ❷ **Implementation view** of file systems
 - Disk and partition layout
 - File tables
 - Free space management
 - ...

- ➊ **User view** of file systems
 - System calls
 - Structures, organisation, file types
- ➋ **Implementation view** of file systems
 - Disk and partition layout
 - File tables
 - Free space management
 - ...
- ➌ There is a lot more happening than expected at first sight!

Goals for Today

Overview

- File system **implementations**
 - 1 Contiguous
 - 2 Linked lists
 - 3 File Allocation Table (FAT)
 - 4 I-nodes (linking files, the Unix V7 file system)
- The **log structured** file system paradigm

Contiguous Allocation

Concept

- **Contiguous file systems** are similar to **dynamic partitioning** in memory allocation:
 - Each file is stored in a single group of **adjacent blocks** on the hard disk
 - E.g. 1k blocks, 100k file, 100 contiguous blocks
- Allocation of free space can be done using **first fit**, **best fit**, **next fit**, etc.




Figure: External Fragmentation when Removing Files

Contiguous Allocation

Advantages

- **Simple to implement:** only location of the first block and the length of the file must be stored
- **Optimal read/write performance:** blocks are co-located/clustered in nearby/adjacent sectors, hence the seek time is minimised (remember the example in lecture on disks!)

Contiguous Allocation

Disadvantages

- Disadvantages of contiguous file systems include:
 - The **exact size** of a file (process) is not always known beforehand: what if the file size exceeds the initially allocated disk space
 - **Allocation algorithms** used to decide which free blocks to allocate to a given file (e.g., first fit, best fit, etc.)
 - Deleting a file results in **external fragmentation**: de-fragmentation must be carried out regularly (and is slower than for memory)
- Contiguous allocation is still in use: **CD-ROMS/DVDs**
 - External fragmentation is less of an issue here since they are write once only

Linked Lists

Concept

- To avoid external fragmentation, files are stored in **separate blocks** (similar to paging) that are **linked to one another**
- Only the **address of the first** block has to be stored to locate a file
- Each block contains a **data pointer** to the next block (which takes up space)




Figure: Linked List File Storage

Linked Lists

Advantages

- **Easy to maintain:** only the first block has to be maintained in the directory entry
- File sizes can **grow dynamically** (i.e. file size does not have to be known beforehand): new blocks/sectors can be added to the end of the file
- Similar to paging for memory, every possible block/sector of disk space can be used: i.e., there is **no external fragmentation!**
- **Sequential access is straightforward**, although **more seek operations/disk access** may be required

Linked Lists

Disadvantages

- **Random access is very slow**, to retrieve a block in the middle, one has to walk through the list from the start
- There is **internal fragmentation**, on average the last half of the block is left unused
 - Internal fragmentation will reduce for **smaller block sizes**
- May result in **random disk access**, which is very slow (remember the example in lecture on disks)
 - **Larger blocks** (containing multiple sectors) will be faster

Linked Lists

Disadvantages (Cont'ed)

- Space is lost within the blocks due to the pointer, the data in a **block is no longer a power of 2!**
- **Diminished reliability:** if one block is corrupt/lost, access to the rest of the file is lost

File Allocation Tables

Concept

- Store the linked-list pointers in a **separate index table**, called a **File Allocation Table (FAT)**




Figure: File Allocation Tables

File Allocation Tables

Advantages

- **Block size remains power of 2**, i.e., no space is lost due to the pointer
- **Index table** can be **kept in memory** allowing fast non-sequential/random access (one still has to walk through the table though)

File Allocation Tables

Disadvantages

- The **size of the file allocation** table grows with the number of blocks, and hence the size of the disk
- For a 200GB disk, with a 1K block size, 200 million entries are required. Assuming that each entry is 3 bytes, this requires **600Mb of main memory!**

I-nodes

Concept

- File attributes and block pointers are stored in separate data structures, called **I-nodes**
 - In contrast to FAT, I-nodes are **only loaded when the file is open** (stored in system wide open file table)
 - If every I-node consists of n bytes, and at most k files can be open at any point in time, at most $n \times k$ bytes of main memory are required
- I-nodes can use **direct block pointers** (usually 12), **indirect block pointers**, or a combination thereof (e.g., similar to **multi-level page tables**)

I-nodes

Concept





Figure: File Storage Using I-Nodes

File System Comparison

Contiguous vs. Linked vs. Indexed


contiguous

directory		
name	start	size
fred.doc	0	2
bill.doc	5	5
jack.doc	16	8



linked list

directory	
name	start
fred.doc	0
bill.doc	1
jack.doc	2



indexed

directory	
name	index
fred.doc	0
bill.doc	8
jack.doc	16




Figure: Contiguous vs. Linked List (or FAT) vs. I-nodes (or indexed)

I-nodes

Lookups

- **Opening a file** requires the disk blocks to be located
 - **Absolute file names** are located relative to the root directory
 - **Relative file names** are located based on the current working directory
- E.g. Try to locate `/usr/gdm/mbox`

/	/	/usr	/usr	/usr/gdm	/usr/gdm
inode	contents	inode	contents	inode	contents
1	2	6	132	26	406
size	1 .	size	6 .	size	26 .
mode	1 ..	mode	1 ..	mode	6 ..
times	4 bin	times	19 fred	times	64 research
2	7 dev	132	30 bill	406	92 teaching
	14 lib		51 jack		60 mbox
	9 etc		26 gdm		17 grants
	6 usr		45 cfi		
	8 tmp				

Figure: Locating a File

I-nodes

Lookups

Locate the root directory of the file system

- Its i-node sits on a fixed location at the disk (the directory itself can sit anywhere)

Locate the directory entries specified in the path:

- Locate the i-node number for the first component (directory) of the path that is provided
- Use the i-node number to index the i-node table and retrieve the directory file
- Look up the remaining path directories by repeating the two steps above

Once the file's directories have been located, locate the file's i-node and cache it into memory

I-nodes

Hard and Soft Links

- Two approaches exist to **share a file**, e.g. between directory B and C (owner):
 - **Hard links**: maintain two (or multiple) a reference to the same i-node in B and C
 - the i-node link reference counter will be set to 2
 - **Symbolic links**:
 - The owner maintains a reference to the i-node in, e.g., directory A
 - The “referencer” maintains a small file (that has its own i-node) that contains the location and name of the shared file in directory C
- What is the **best approach** \Rightarrow both have advantages and disadvantages

I-nodes

Hard Links

- Disadvantages of hard links:
 - Assume that the owner of the file **deletes** it:
 - If the i-node is also deleted, any hard link will, in the best case, **point to an invalid i-node**
 - If the i-node gets **deleted** and “**recycled**” to point to an other file, the hard links will **point to the wrong file!**
 - The only solution is to **delete the file**, and **leave the i-node intact** if the “**reference count**” is larger than 0 (the original owner of the file still gets “charged” for the space)
- Hard links are the **fastest way of linking files!**

I-nodes

Hard and Soft Links

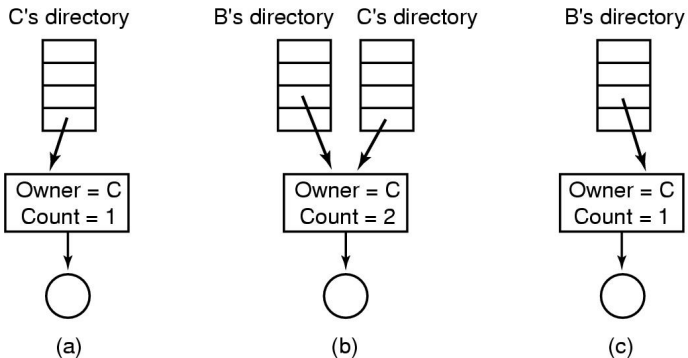


Figure: Hard Links (Tanenbaum)

I-nodes

Symbolic Links

- Disadvantages of symbolic links:
 - They result in an **extra file lookup** (once the link file has been found, the original file needs to be found as well)
 - They require an **extra i-node** for the link file
- Advantages of symbolic links:
 - They can **cross the boundaries of machines**, i.e. the linked file can be located on a different machine
 - There are **no problems with deleting** the original file \Rightarrow the file simply does not exist any more

The Unix V7 File System

Implementation

- **Tree structured** file system with links
- Directories contain **file names** and **i-node numbers**
- I-nodes contain **user attributes** and **system attributes** (e.g. count variable)
- One single, double, and triple **indirect blocks** can be used

Log Structured File System

Context

- Consider the **creation of a new file** on a Unix system:
 - ① Allocate, initialise and write the i-node for the file
 - i-nodes are usually located at the start of the disk
 - ② Update and write the directory entry for the file (directories are tables/files that map names onto i-nodes in unix)
 - ③ Write the data to the disk
- The corresponding blocks are not necessarily in **adjacent locations**!

Log Structured File System

Context

- **Disks are slow** compared to other components in a computer (e.g. CPU)
 - Is there any way in which we can make them faster?
 - I.e., can we develop a file system that copes better with the inherent delays of traditional disks
- This is what log structured file systems try to achieve: improve speed of a file system on a traditional hard disk by **minimising head movements** and **rotational delays**

Log Structured File System

Concept

- Log structured file systems **buffer read and write operations** (data, etc.) in memory, enabling us to **write “larger volumes”** in one go
- Once the buffer is full it is “flushed” to the disk and written as **one contiguous segment** at the end of “**a log**”
 - i-nodes and data are all written to the same segment
 - Finding i-nodes (traditionally located at the start of the partition) becomes more difficult
- An **i-node map** is maintained in memory to quickly find the address of i-nodes on the disk

Log Structured File System

Concept

- A **cleaner thread** is running in the background that retrieves data from the back of the log, removes the deleted files and places them in the buffer
- Once the buffer is full the **data retrieved at the end** will be **written to the front of the log**
- I.e., a hard drive is treated as a **circular buffer**




Figure: Log Structured File System

Recap

Take-Home Message

- File system **implementations**
- The **log structured** file system paradigm

Next Lecture

File System Paradigms

- Journaling file systems
- Virtual file system
- Consistency checking/fixing