

Linux Case Study

OPS Lecture 18, G53OPS/G52OSC

Geert De Maere

(Jason Atkin – OSC)

Geert.DeMaere@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2015

Linux

Structure of the Linux Kernel

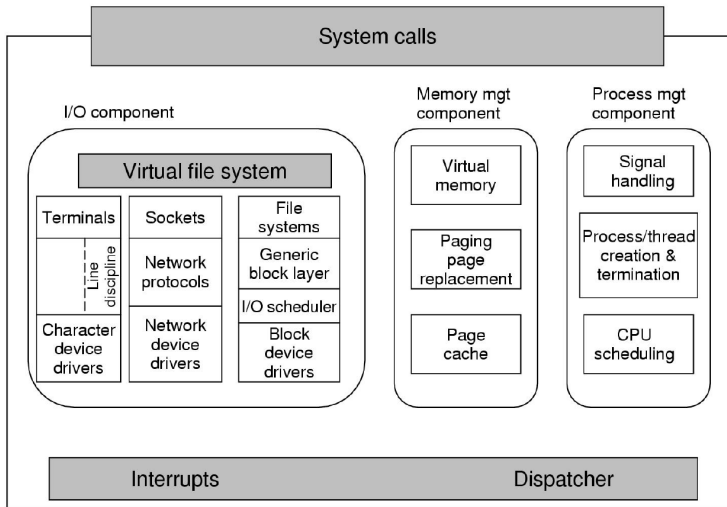


Figure: Structure of the Linux Kernel (Tanenbaum)

Memory Management

Swapping and Paging in Linux

- Linux uses a **pure demand paging** approach (i.e. no pre-paging or working sets) and pages are reclaimed by the **paging daemon**
 - Loading a process **sets up the page table** but does not load the pages
- “**Swapping**” is implemented on a **page granularity** (\Leftrightarrow entire process)
 - I.e. paging and virtual memory are used
 - The process does not have to be entirely in memory

Memory Management

Page Replacement in Linux

- Linux's **page replacement algorithm** is a modified version of the **second chance clock algorithm** based on **age**:
 - Each pass of the clock reduces the **age value**
 - Each **reference** to a page **increases the age value**
- The age value is used as a basis for the **least recently used** algorithm




Figure: Memory Image (Tanenbaum)

Memory Management

Swapping and Paging in Linux

- Pages can be mapped:
 - Onto files on the disk, e.g. for **text segments** and **memory mapped files**
 - Onto a **paging file** or **paging partition**
- **Partitions are faster** since it avoids file system overheads, pages are usually written contiguously
- A **next fit algorithm** is used for swap partitions, aiming to write pages **contiguously**

File System

History of the Linux file system

- **Minix file system**: the maximum file size was 64MB and file names were limited to 14 characters
- The “**extended file system**” (extfs): file names were 255 characters and the maximum file size was 2 GB
- The “**ext2**” file system: larger files, larger file names, better performance
- The “**ext3-4**” file system: journaling etc.

File System

History of the Linux file System

- Anything capable of handling input or output of data is regarded as a file in Linux
 - **Character special files** model serial devices (e.g. keyboards, mice, printers)
 - **Block special files** model raw block devices (e.g. hard drives without file systems, used for swap partitions)
- **Directories** are **implemented as files** and lookups take place in a similar way to Unix

File System

History of the Linux file system

- **Files** and **directories** are manipulated through **system calls**, e.g. `open` and `close`
- **File descriptors** are non-negative integers used as an index for the **file descriptor table** (0, 1, and 2 are used for standard input, output, and error)

System Call	Description
<code>fd = creat(name, mode)</code>	Create a new file
<code>fd = open(file, how, ...)</code>	Open file
<code>s = close(fd)</code>	Close file
<code>n = read(fd, buffer, nbytes)</code>	Read data into buffer
<code>n = write(fd,buffer, nbytes)</code>	Write data into buffer
<code>position = lseek(fd, offset, whence)</code>	Move file pointer

Table: Examples of system calls for file manipulation

File System

History of the Linux File Sytem

System Call	Description
<code>s = mkdir(path, mode)</code>	Create a new directory
<code>s = rmdir(path)</code>	Remove directory
<code>s = link(oldpath, newpath)</code>	Create link
<code>n = unlink(path)</code>	Unlink a file
<code>n = chdir(path)</code>	Change CWD
<code>dir = opendir(path)</code>	Open directory
<code>s = closedir(dir)</code>	Close directory
<code>dirent = readdir(dir)</code>	Read one directory entry

Table: Examples of system calls for directory manipulation

File System

The Virtual File System

- Linux and Unix use a **virtual file system** (VFS) that hides the implementation specific details behind a generic interface
- **Key objects** in the VFS implementation include:
 - **Super block object** containing file system information for every mounted file system (e.g. number of i-nodes and access to i-nodes)
 - **File objects** for in memory representations of files created by an `open ()` call , and containing e.g a file pointer, file permissions, etc.
 - **I-node object** containing information on individual files
 - **Dentry objects** representing directory entries and which are cached
- Objects contain **pointers to function tables** that list addresses for the actual **implementations**

File System

The Extended 2 File System




Figure: Ext2 Partition Layout (Tanenbaum)

- The **superblock** contains file system information (e.g. the number of i-nodes, disk blocks)
- The **group descriptor** contains bitmap locations, the number of free blocks, i-nodes and directories
- A **data block bitmap** and **i-node bitmap**, used to keep track of free disk blocks and i-nodes (Unix uses lists)
- A **table of i-nodes** containing file and disk block information
- **Data blocks** containing file and directory blocks

File System

The Extended 2 File System

- An ext2 partition is split into several **block groups** to:
 - **Reduce fragmentation** by storing i-nodes and files, and parent directories and files in the same block group if possible
 - Reduce **seek times** and improve performance
- All block groups have the same size and are stored sequentially (which allows direct indexing)

File System

Directory Entries

- Every **directory** contains the following fields:
 - i-node number
 - Entry size in bytes
 - Type field, i.e. file, directory, special file, etc.
 - File name length in bytes
- Directories are **searched linearly** (i.e. they are unsorted) and a cache is maintained for recently accessed items




Figure: Ext2 Directory Entry (Tanenbaum)

- **File lookups** are similar to the Unix file system
- The **i-node structure** is similar to the Unix i-nodes
 - 12 block addresses are contained in the i-node
 - Single, double and triple indirect blocks are used

File System

The Ext3 File System

- When making changes to an Ext2 file system, files are ideally **written immediately** to prevent inconsistency:
 - This generates **significant head movement**
 - Ext2 File system is more suitable for **flash disks** (no journal)
- Ext3 builds upon the Ext2 file system by adding:
 - **Tree based structures** for directory files to facilitate indexing (HTrees)
 - **Journaling** capabilities
 - **Meta structures** remain stored in fixed locations

Summary

Take-Home Message

- Demand paging, swapping, page replacement
- File system management, including Ext2 and Ext3